

The officially “not recommended” resource on the Black Magic Probe

Embedded Debugging with the Black Magic Probe



Copyright 2022 © Thiadmer Riemersma, CompuPhase.

This work is licensed under the Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International License (BY-NC-ND). To view a copy of this license, visit <https://creativecommons.org/licenses/by-nc-nd/4.0/> or send a letter to Creative Commons, PO Box 1866, Mountain View, CA 94042, USA.

This book is available in PDF format on the CompuPhase web site and on GitHub, along with utilities and coding examples. See [Further Information](#) on [page 119](#) for a link.

The cover image is by Arek Socha.

ISBN 9789090349879

Typeset with T_EX in the “DejaVu” typeface family.

Contents

Introduction	1
Hardware and Software	1
Why bother, why choose the difficult route?	2
About this Book	3
License	4
The Debugging Pipeline	5
GDB Architecture	5
The Serial Wire Debug Protocol in a Nutshell	7
Embedded Debugging: Points for Attention	10
Requirements for Front-ends	11
Hardware Overview	13
Accessories	15
Setting up the Black Magic Probe	18
Microsoft Windows (USB set-up)	18
Linux (USB set-up)	20
Wi-Fi set-up for ctxLink	22
Connecting the Target	24
Checking the Set-up	24
Running Commands on Start-up	26
Design for Debugging	27
Debugging Code	29
Prerequisite Steps	30
Loading a File and Downloading it to the Target	30
Starting to Run Code	34
Getting help and information	35
Listing Source Code	35
Downloading code into the microcontroller	36
Stepping and Running	36
Breakpoints and watchpoints	38
Examining Variables and Memory	40
The Call Stack	42
Inspecting Machine Code	43
Debug Probe Commands	43
The BlackMagic Debugger Front-end	47
Edit-Compile-Debug Cycle	54
Debugging Optimized Code	55

Run-Time Tracing	56
Levels of Tracing	56
Secondary UART	57
Semihosting	57
SWO Tracing	61
Tracing with Command List on Breakpoints	71
The Common Trace Format	72
Binary Packet Format	74
A Synopsis of TSDL	74
Generating Trace Support Files	83
Integrating Tracing in your Source Code	84
Mixing Common Trace Format with Plain Tracing	85
Applications for Run-Time Tracing	86
Code Assertions	86
Tracing Function Entry and Exit	89
Code Profiling	92
Sampling on ARM Cortex	92
Calltree Analysis	94
Firmware Programming	97
Using GDB	97
Using the BlackMagic Flash Programmer	97
Updating Black Magic Probe Firmware	103
Troubleshooting	105
Check whether the system detects the probe	105
Check whether the probe detects the target	106
Target scan hangs	109
GDB crashes on “attach”	110
Failure to erase Flash memory	111
Spying on the communication	111
TRACESWO Capture	112
TTL-Level UART	113
GDB on Microsoft Windows	114
Microcontroller Driver Support	115
Linking TRACESWO to UART-RxD	118
Further Information	119
Hardware	119
Software	119
Articles, Books, Specifications	121
Index	122

Introduction

The “Black Magic Probe” is a combined hardware & software project. At the hardware level, it implements JTAG and SWD interfaces for ARM Cortex A-series and M-series microcontrollers. At the software level, it provides a “gdb-server” implementation and Flash programmer support for ranges of microcontrollers of various brands. The hardware was designed by 1BitSquared in collaboration with Black Sphere Technologies. The embedded software of the Black Magic Probe is an open-source project (the hardware is less open: schematics are available for legacy versions of the hardware, but not for the current version).

At the time of writing, the Black Magic Probe hardware is version 2.3, and the firmware is at version 1.8. The firmware in the Black Magic Probe is in ongoing development, so the firmware in newly purchased probes may be behind the current release, see also [Updating Black Magic Probe Firmware on page 103](#)). Derivatives of both hardware and firmware exist, with sometimes different capabilities or limitations. This book focuses on the *native* hardware, and firmware version 1.6 or later —though notes on ctxLink (a derivative of the Black Magic Probe with enhanced features) appear where applicable.



Hardware and Software

Separate from the MCU core, the ARM Cortex series have a Debug Access Port (DAP) that gives you access to the debugging features of the microcontroller. On older architectures, the debugging interface used the JTAG port and protocol, but for the ARM Cortex series, a new protocol that required less physical pins was designed: the ARM *Serial Wire Debug protocol* (SWD). This protocol gives you access to features like single-stepping, hardware breakpoints and watchpoints, dumping memory regions and programming Flash memory. Like was the case with the JTAG interface, the SWD interface is meant to be driven by a hardware interface, a *debug probe*.

The Black Magic Probe is such a debug probe. The “black magic” that it adds to alternative debug probes is that it embeds a software interface for GDB, the debugger for GNU GCC compiler suite —a widely used compiler for microcontroller projects. It is the closest that a debug probe can come to plug-&-play operation.

Next to the Black Magic Probe, you need GDB, and more specifically, the GDB from the toolchain that you use to build your embedded code. For the ARM Cortex-A and Cortex-M microcontrollers, this typically means the GDB from the arm-none-eabi toolchain.¹

While you do not need a debugger front-end, it is beneficial to get one. When you are running on Linux, you may get by with GDB’s integrated Text User Interface —it’s rudimentary, though. See [Requirements for Front-ends \(page 11\)](#) for tips to select a front-end.

Why bother, why choose the difficult route?

Advice that I have repeatedly seen on blogs and answers on stackoverflow (and others), is to make the software modular, debug each module on a desktop PC or laptop, and to then assemble the embedded application from these fully tested and debugged modules. The implied message is that embedded software is fundamentally the same as desktop software, but you have the cream of the crop in development tools on desktop systems.

Allow me to draw a parallel from a different field: From the earliest days of medicine, the focus has been on studying the physiology of men. It was assumed that the female body responds to medication and drugs in the same way as that of men. Up to the 1960s, clinical trials for a new drug were done on sometimes thousands of men, and zero women. Women, after all, would supposedly only bring “confounding issues” to the trials, due to their alleged emotional instability and fluctuating hormone levels. It leads to absurdities like a clinical trial in 2015 for Addyi, a drug to treat female sexual dysfunction. The trial involved 23 men and 2 women —a drug exclusively for women tested almost exclusively on men. All the while, the presumption that the female body is that of a man (though with confounding issues) is entirely unfounded. Sex bias in medicine is a symptom of complacency and indifference.

Embedded devices are a varied lot, but as a general rule, they are not just a PC with confounding issues. Software that runs fine on a desktop system may fail on the target microcontroller. Not every microcontroller handles unaligned memory access alike, for example. Embedded devices commonly

¹ With a caveat for GDB versions 11.x, up to 12.1. These releases have a bug that makes GDB fail to connect to the target microcontroller. See [page 110](#) for details.

have (integrated) peripherals that desktop systems lack, and on an embedded device those peripherals will be driven with the SPI or I2C protocols, rather than USB.

The recommendation to develop and debug embedded software on a desktop, on the dogma that it should then run alike on the embedded device, is similarly based on an invalid assumption and an ill-advised desire to stick with the familiar tools and environment. It will actually work on specific cases, such as a generic data structures library, but it is a bad strategy overall.

In my consulting work, I get occasionally to listen to a “war story” by a fellow developer, about failures, glitches and missed interrupts. But in a simulator on a PC it ran flawlessly, so... Often, the issue was circumvented rather than solved. For example, to “fix” a case of an occasionally missed interrupt, the developer set an interrupt on both rising and falling edges of a pulse because it hadn’t happened yet that the MCU missed two interrupts in succession. Sometimes the hardware was redesigned to use an MCU of a different brand or architecture (but yet, without evidence that the original MCU was at fault). Frequently, the frustration about the wasted time and resources hadn’t subsided yet —one developer claimed to have “even switched to Torx screws” so much he had come to detest Philips.²

I was not there to debug their systems, so we will never know the truth. The point is, developing code intended to run on an embedded device and testing it exclusively on a desktop system, is as absurd as developing a drug exclusively for women and testing it on men. And while these companies found workarounds, the real point is the wasted time and resources, both of which cost money.

About this Book

This guide is not a book on GDB. That book is *The Art of Debugging with GDB, DDD and Eclipse* by Norman Matloff and Peter Salzman,³ and which is highly recommended. This guide does not delve into the hardware and software design of the Black Magic Probe, either. The software of the Black Magic Probe is open-source, hosted on GitHub, and information on the hardware is best obtained from 1BitSquared (the manufacturer).

Instead, this guide aims at describing how to use the Black Magic Probe to debug embedded software running on an ARM Cortex microcontroller. It starts

² It didn’t help me laughingly pointing out that the well known cross-slotted screw head is Phillips —double “l” and entirely unrelated to the semiconductor brand Philips (now NXP).

³ Matloff, Norman and Peter Jay Salzman; *The Art of Debugging with GDB, DDD, and Eclipse*; No Starch Press, 2008; ISBN 978-1593271749.

with an overview of the debugging pipeline, from the target microcontroller to the visualization of the embedded code on your workstation. Debugging embedded code usually implies remote debugging (with the code that is being debugged running on a different system than the debugger), but also cross-platform debugging. A broad understanding of these is helpful when making practical use of the Black Magic Probe.

The next chapters focus on setting up the hardware and software for the Black Magic Probe, and then a selection of GDB commands, with a special focus on those that are particularly useful for debugging embedded code.

Run-time tracing is an essential debugging technique for embedded systems, due to the real-time requirements that these systems often have. Coverage is split in three chapters: the first on the hardware and software support in the Black Magic Probe, the second on generic techniques to perform tracing efficiently, and the third on particular applications of run-time tracing.

The Black Magic Probe can also be used for production programming of devices, through the same mechanism that GDB uses to download code to the target for purposes of debugging. This is the topic of another chapter, using both GDB and a separate utility.

The final (short) chapters are on updating the firmware of the Black Magic Probe itself and adding support for new microcontrollers to the GUI utilities that accompany this guide.

License

This guide is written by Thiadmer Riemersma and copyright © 2020–2022, CompuPhase. It is licensed under the Creative Commons BY-NC-ND 4.0 International License (Attribution-NonCommercial-NoDerivatives).

The associated software is copyright © 2019–2022 CompuPhase and licensed under the Apache License version 2.

The Debugging Pipeline

Developing embedded software on small microcontrollers presents some additional challenges in comparison with desktop software. The software is typically developed on a workstation and then transferred to the target system. Accordingly, cross-compiling and remote debugging are the norm. Remote debugging implies the use of a hardware box or interface to connect the workstation to the microcontroller's debug port & protocol. On the ARM Cortex processors, the most common debug and Flash programming protocols are JTAG and SWD (Serial Wire Debug).

In the idiom of remote debugging, the target is the device being debugged, and the host is the workstation that the debugger runs on. The interface between host and target is the probe. A debug probe typically connects to the workstation's USB, RS232 or Ethernet port.

GDB Architecture

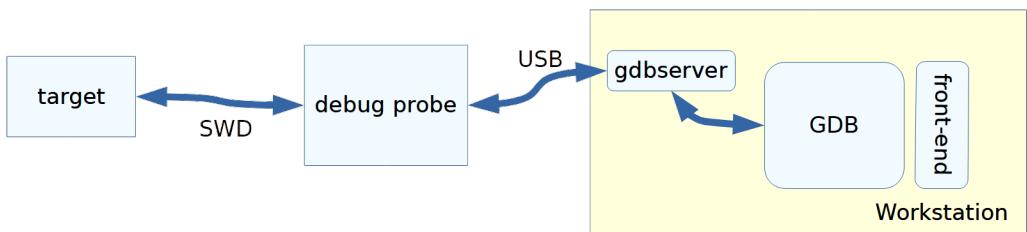
GDB is the GNU Debugger for programs built with GCC. It is also a debugger framework, with third-party front-ends and machine/protocol-specific back-ends.

GDB's user interface is, by today's standard, rather rudimentary, but GDB provides a "machine interface" to "front-ends", so that these front-ends can provide a (graphical) user interface with mouse support, source browser, variable watch windows, and so forth, while leaving symbol parsing and execution stepping to GDB. Most developers who use GDB actually run it hidden behind a front-end like Eclipse, KDbg, DDD, or the like. As a side note, a text-based front-end is built-in: TUI, and while it is an improvement over no front-end at all, TUI is not as stable as the alternatives (it is also broken on Microsoft Windows, and there is no plan to fix it).

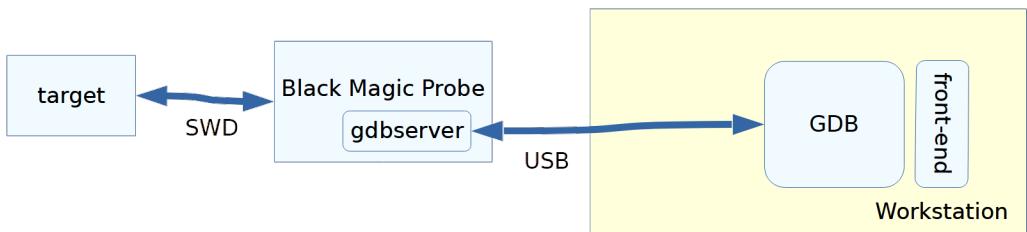
To debug a different system than the one where the debugger runs on, GDB provides the *Remote Serial Protocol* (RSP). This is a simple text-based protocol with which GDB on the workstation communicates with a debugger "stub" on the target system. This stub acts as a server that GDB connects to, over an RS232 or Ethernet connection, and it is referred to as a `gdbserver`.

Directly implementing a `gdbserver` on the *target* is impractical for microcontrollers such as the ARM Cortex M series. These microcontrollers provide hardware support for setting breakpoints and stepping through code, but make it available on a *separate* interface with dedicated pins for the task. On the ARM Cortex, this is *Serial Wire Debug interface* (SWD). To drive the serial wire protocol, a debug probe is needed: a hardware interface that drives the clock and data lines according to the SWD protocol. Common debug probes

are Segger J-Link and Keil ULINK-ME. The `gdbserver` functions as an interface to translate between GDB-RSP and the protocol of the hardware interface.



As is apparent, the debug data goes through a few hoops before the developer sees the code and data on the computer display in “GDB.” The OpenOCD project is an example of this set-up.¹ The main `openocd` program implements `gdbserver`, and it opens a Telnet port for the communication link to GDB and a USB, RS232 or Ethernet connection to the debug probe.



The Black Magic Probe embeds `gdbserver`. One advantage of this design is that its `gdbserver` has in-depth knowledge of the capabilities of the debug probe as well as what the debug probe has discovered about the target. The only configuration that needs to be done in GDB is the (virtual) serial port of the Black Magic Probe (the USB interface of the Black Magic Probe is recognized as a serial port on the workstation).

The ctxLink debug probe functions identically to the Black Magic Probe when connected to the USB —in fact, it even uses the same VID:PID codes. However, ctxLink also offers connection over a Wi-Fi link. In relation to the diagram above, this changes very little: in essence, you only need to change the caption of the “USB” link to “Wi-Fi” (disregarding that there is also a wireless switch or access point thrown in the mix). But the implication is that while the range of a USB-connection is limited, ctxLink makes the debug probe accessible over the local network and (after configuring the router) over the internet. Thereby, ctxLink enables debugging over a technically unlimited distance.

¹ In a “hosted” set-up, the Black Magic Probe also uses this configuration: the main firmware of the Black Magic Probe (with the embedded `gdbserver`) runs as a desktop application on the workstation, and the Black Magic Probe hardware is reduced to function as a dumb probe. See section [Check whether the probe detects the target](#) on page 106 for more information about this option.

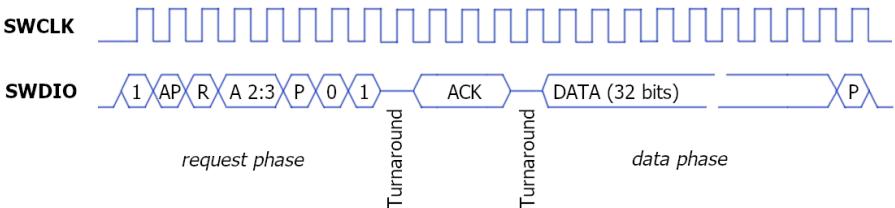
Bypassing GDB

While the *Remote Serial Protocol* (RSP) is specifically designed for GDB (to communicate with `gdbserver` in a debug probe), it is well-defined and well-documented, and therefore other tools can use it without requiring GDB. In fact, the [BMTrace](#) (page 69) and [BMFlash](#) (page 97) utilities do exactly this. These utilities use a fairly limited subset of the capabilities of `gdbserver`.

Troll is a source-level debugger for ARM Cortex architecture using GDB's RSP, and it is thereby compatible with the Black Magic Probe. The Troll debugger is still an experimental project; see [Further Information](#) on page 119 for a link to the project.

The Serial Wire Debug Protocol in a Nutshell

The Serial Wire Debug protocol (SWD) is designed as an alternative to the JTAG protocol, for microcontrollers with a low pin count. It is part of the ARM Debug Interface specification version 5, abbreviated as ADI5. At the physical layer, it needs two lines at the minimum (plus ground), as opposed to five for JTAG. These two lines are the clock (SWCLK, driven by the debug probe) and a bidirectional data line (SWDIO). Tracing output goes over a third (optional) line: TRACESWO, but using an unrelated protocol (independent of SWCLK) —see section [TRACESWO Protocol](#) on page 8.



The SWCLK signal is driven by the debug probe, regardless of the direction of the transfer. Each transaction starts with a request, that the probe sends to the target. The target replies by sending an acknowledgement back. After that, a *data phase* follows, which may be in either direction, depending on the request.

The target microcontroller polls the SWDIO pin on a rising edge of SWCLK, and also drives the pin on a rising edge. When idle, the SWCLK and SWDIO pins are driven low by the debug probe.

As is apparent, the direction of the SWDIO line switches between input and output at least once during a transaction, on both sides. The SWD protocol calls this the *turnaround*, and there is an extra clock cycle for each turnaround in the transaction. The pictured example is for a *write* transaction; in a *read*

transaction, there is no turnaround after the ACK —however, if another transaction follows head-to-tail, a turnabout is added after the data phase.

The request phase is a sequence of 8 bits. First comes a start bit (always 1). The AP bit is 0 if this transfer is for the Debug Port (DP), and 1 if it is for an Access Point (AP) such as MEM-AP, which provides access to core memory and peripheral registers. The R bit is 1 for a read request and a 0 for a write request. There are two address bits, to access the debug registers. The P bit is a parity bit, it is set such that the sum of the bits in the request byte is even. Following the parity bit are a stop bit and a park bit, which are 0 and 1 respectively.

The ACK is a three bit sequence with the value 1 (on success), sent with the low bit first. The data is likewise transmitted low bit first. After the 32-bits of data are transmitted, another parity bit follows (calculated such that the sequence of 33 bits has even parity).

With two address bits in a transfer request, you can only address four registers. To access code or data memory, the access port of the AHB provides the TAR register. In this register, you set a memory address so that you can read from or write to that memory location on a subsequent transfer. A peculiarity of the SWD protocol is that a read transfer returns the value from the previous transaction. Hence, to read the current value of a register or memory location, you need to perform the read operation twice, and discard the first result.

Before a microcontroller's SWD port is serviceable, an initialization sequence must be performed, part of which is to switch the protocol from JTAG to SWD. Some ARM Cortex microcontrollers do not support JTAG, but the protocol requires that the JTAG-to-SWD switch is still performed.

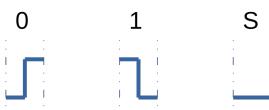
TRACESWO Protocol

The TRACESWO protocol is independent of the SWD protocol. You can trace without debugging, as well as debug without tracing. The data is transmitted over a single line using one of two serial formats: asynchronous encoding and Manchester encoding. The ARM documentation occasionally refers to these encodings as NRZ and RZ (Non-Return-to-Zero and Return-to-Zero) respectively. The TRACESWO protocol is handled by the *Instrumentation Trace Macrocell* (ITM) of the ARM Cortex core.

The asynchronous encoding is in essence TTL-level UART, with a start bit, eight data bits, one stop bit and no parity. As is common with UART protocols, the target and the debug probe must use the same bit rate within a narrow margin. Since the UART clock is typically derived from the microcontroller clock, at high bit rates it becomes harder to find a bit rate shared by both the target and the debug probe within the required margin.

Manchester encoding, on the other hand, has the property that the clock frequency can be established from the data stream. This makes it a self-adapting protocol, tolerant to jitter and insusceptible to clock drift. These properties make Manchester encoding the option of choice for microcontrollers that lack hardware support for SWO tracing (such as the ARM Cortex-M0 and Cortex M0+ architectures) because it is easier to implement it with bit-banging. A drawback of Manchester encoding is that the encoding takes two clocks per bit, which means that the maximum bit rate is typically half as high as for asynchronous encoding.

The physical Manchester protocol on the TRACESWO pin transmits sequences of 1 to 8 bytes, where each sequence is prefixed with a start bit (a 1-bit) and suffixed with a “space.”



The pin is low on idle; a 0-bit has a rising edge halfway the bit period, a 1-bit has a falling edge halfway the bit period, and a space is a low level for the full bit period.

Obviously, since a 1-bit starts high, if the pin is low at the start of the bit period, there is also a rising edge at the start of the 1-bit. This occurs when the previous bit is also a 1-bit, or when the previous state was idle or space. Similarly, there is a falling edge at the start of a 0-bit if the pin is high at the start of the 0-bit, which occurs when the previous bit was also a 0-bit. A space resets the decoder state back to idle.

Although Manchester is a *bit* transmission protocol, the ITM always transmits a multiple of 8 bits of data (least-significant bit first). After a start bit and up to 64 data bits (8-bytes) have been transmitted, a space follows and after that (if there is more data to transmit) a new start bit plus another sequence of data. This short interruption after every 64-bits is to resynchronize the bit stream. The start bit is needed to determine the clock frequency of the protocol (the start bit is transmitted from idle state, so there is a rising edge at the start of the bit and a falling edge half way), and the space at the end of a sequence is needed to properly decode the next start bit (it needs to come after a known state).

At a higher level, the TRACESWO protocol transmits *packets* consisting of an 8-bit packet header followed by a 32-bit payload (transmitted low byte first). The protocol uses trailing-zero compression on the payload, which means that if only one or two bytes are transmitted, these form the low bytes of the 32-bit word and the high bytes of that word are assumed zero.

The header byte contains the channel number in the highest five bits. The

low three bits indicate the number of payload bytes that follow; the value can be 1, 2 or 3, where 3 means that *four* payload bytes follow.



Events generated by the *Data Watchpoint & Trace* unit (DWT), that the ITM passes through, use a packet header as well. The three low bits in the header are set to combinations that are invalid for a trace message. Thus, trace monitoring applications can test the three low bits to check whether to process or reject a packet.

Embedded Debugging: Points for Attention

On desktop computers and single-board computers, programs run in RAM. A debugger sets a breakpoint at a location by storing a special software interrupt instruction at that location (after first saving the instruction that was originally at that location). When the instruction pointer reaches the location, the software interrupt instruction causes the corresponding exception to be raised, which is intercepted by the debugger, which then halts the target program. The debugger also quickly puts the original instruction back into RAM, so that when you resume running the target, it will execute the original instruction.

Current microcontrollers often have limited SRAM, but a larger amount of Flash memory. The program for microcontroller projects therefore typically runs from Flash memory. For the purposes of running code, you may regard Flash memory as ROM; technically, it is re-writable, but re-writing is slow and needs to be done in full sectors. The upshot is: a debugger cannot set a breakpoint by swapping instructions in memory because the memory (for practical purposes) is read-only.

The solution for the debugger is to team up with the microcontroller and tell the microcontroller to raise an exception if the instruction pointer reaches a particular address. This is called a hardware breakpoint (the former breakpoints are occasionally called software breakpoints). However, ARM Cortex microcontrollers provide only a few hardware breakpoints, typical values are below:

Core	Breakpoints	Watchpoints
Cortex M0 / M0+	4	2
Cortex M3 / M4	6	4
Cortex A / R	6	2

A common architecture for an embedded application is one where the system responds to events (from sensors, switches, or a databus) in a timely manner.

The criterion “timely” regularly means: as quickly as possible, which then means that it is common to handle the event (and its response) in an interrupt. With crucial activity happening in various interrupt service routines, a puzzle that frequently pops up is that a global variable (or a shared memory buffer) takes on an unexpected value. A *watchpoint* can then tell you where in the code that variable got set. A watchpoint is a breakpoint that triggers on data changes. As with breakpoints, you will want hardware watchpoints, so that setting a watchpoint won’t interfere with the execution timing of the code.

Code that is stopped and stepped-through may not follow the same logic flow as code that executes in normal speed, because events or interrupts are missed or arrive in a different context (and those interrupts may set global variables or set semaphores). This change of behaviour may lead to bugs that “disappear” as soon as you try to debug them. The approach to tackle this situation is by tracing the execution path. Tracing can take multiple forms, from “printf-style” debugging to hardware support that records the entire execution flow of a session for post-mortem analysis.

A tracing technique that is unique to GDB is to add a command list to a (hardware) breakpoint, where the last command in the list immediately continues execution after recording that the breakpoint was passed. This way, you can evaluate which points in the code were visited and which were not, move the breakpoints to closer to the area where the bug is suspected and run another session —all without needing to edit and rebuild the code.

Requirements for Front-ends

GDB has powerful and flexible commands, but its console interface falls short of what is needed. Code is hard to follow if you only see a single line at a time. While you can routinely type the `list` command on the “`(gdb)`” prompt, it is clumsy, and it distracts you from focusing on locating any flaws in your code. A front-end that provides a full-screen user interface is therefore highly desirable.

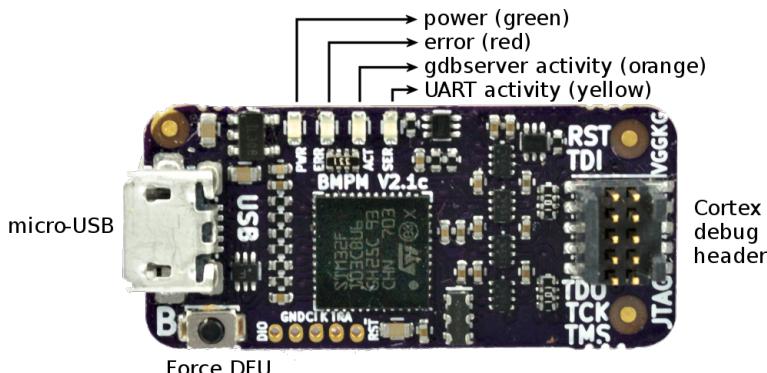
The front-end should not do away with the console, though. Some of the more advanced commands of GDB are not easily represented with icons and menu selections. This is especially true for remote debugging, and even more so for remotely debugging embedded systems. Without the ability to set or read the debug probe’s configuration, via the `monitor` command, your set-up depends on the defaults in the probe, which may not be appropriate for the target. Without the ability to set hardware breakpoints, you may not be able to debug code that runs from Flash memory; and as mentioned, running from Flash memory is the norm on small microcontrollers.

In a misguided attempt to increase “user-friendliness,” KDbg, Nemiver and the Eclipse IDE hide the GDB console (Eclipse has a console tab in its “debug

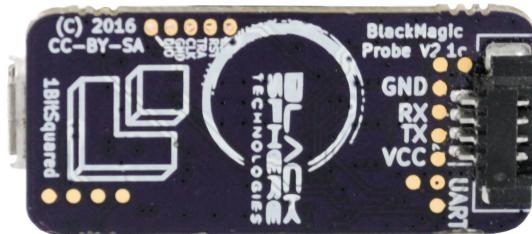
mode,” but it is not the GDB console). Fortunately, this still leaves several front-ends to choose from on Linux: DDD & gdbgui work well, and GDB’s internal TUI is adequate. The TUI is not available on Windows builds of GDB, and DDD has not been ported to Windows. However, gdbgui runs in a browser, there is a “Cortex Debug” extension for the Visual Studio Code editor, and two (commercial) alternative front-ends for Microsoft Windows are WinGDB and VisualGDB (both function as plug-ins to Microsoft’s Visual Studio). Finally, a few GDB front-ends specifically designed for the Black Magic Probe exist. One of these was developed along with this book, and it is covered extensively in section [The BlackMagic Debugger Front-end](#) on [page 47](#). For an alternative, see [Further Information](#) on [page 119](#) and specifically the front-end “turbo.”

Hardware Overview

There are two versions of the *native* Black Magic Probe hardware in current use. Version 2.1, the Black Magic Probe “mini,” was available from 2016 to 2022. It is a 33×15 mm PCB, with a micro-USB connector for linking it to a workstation and a 2×5 -pins 1.27 mm pitch “debug” header for connection to the target microcontroller. See section [Connecting the Target on page 23](#) for details on the Cortex Debug header.



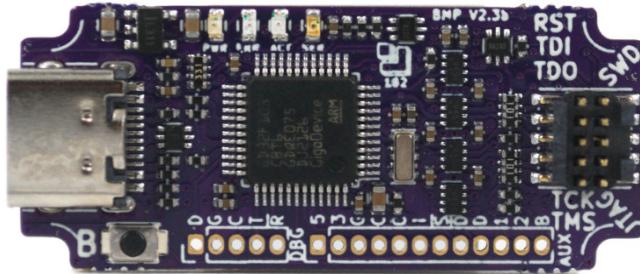
Next to the two connectors, the Black Magic Probe has an on-board switch (that you will only use to upgrade the firmware to the Black Magic Probe, see [Updating Black Magic Probe Firmware on page 103](#)) and four LEDs that signal power and activity status.



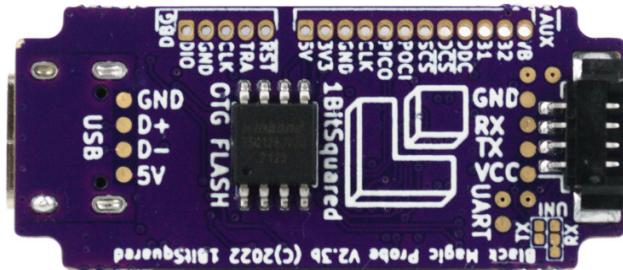
On the reverse site, the Black Magic Probe has a third connector, for a secondary TTL-level UART. This is a 4-pins 1.25 mm pitch “PicoBlade” connector. The function of the four pins is annotated in the silk-screen text on the back. A 145 mm cable with four coloured wires and a suitable PicoBlade connector is provided with the Black Magic Probe.

Hardware version 2.3 is the current release (as of 2022) of the Black Magic Probe, after the last batch of version 2.1 sold out. At 39×17 mm, it is slightly larger than its predecessor. The USB connector has been upgraded to USB-C. Other changes to the hardware design, such as the extra Flash memory on the bottom side for “on-the-go” programming, may become more important

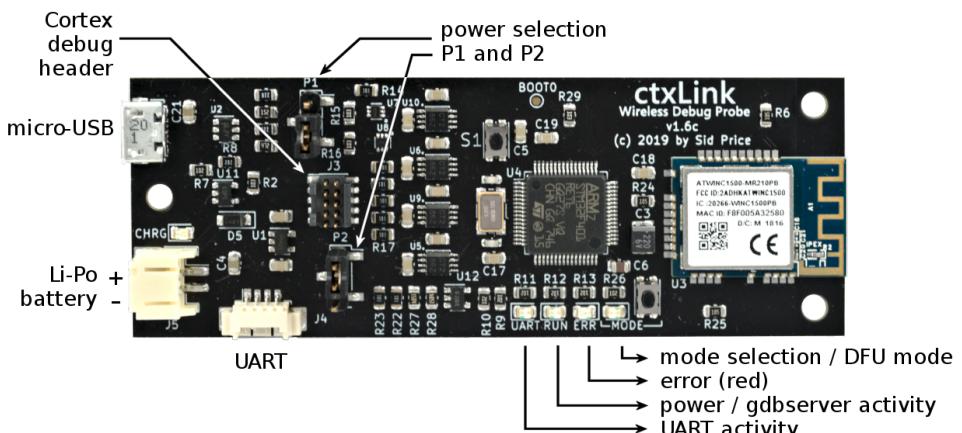
in future releases of the firmware —but they are currently unused. The LEDs and the on-board switch are the same as on the earlier version, and positioned in the same positions.



Like version 2.3, there is a PicoBlade connector for a TTL-level UART.



The ctxLink probe is larger than the Black Magic Probe, at 89×33 mm. All components and connectors are on the top side — the PCB uses only surface-mount connectors, thus the bottom side is completely flush. There are three 3.2 mm mounting holes.



Like the Black Magic Probe:, the ctxLink probe has a USB port, a Cortex De-

bug header, and a TTL UART connector (a 4-pins 1.25 mm pitch “PicoBlade” connector). It also has four LEDs; three of them have the same functions as on the Black Magic Probe, the fourth shows the connection mode. In addition to the shared features, the ctxLink probe has a connector for a rechargeable battery (a JST PH-series, 2-pin with a pitch of 2 mm) and a Wi-Fi module.

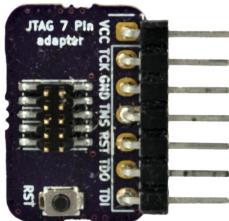
The ctxLink probe can be powered from multiple sources. It uses two jumpers, P1 and P2, for selecting the power source (see the image above for the locations of P1 and P2). These jumpers should be appropriately set before connecting the ctxLink to power.

Power selection	P1 (source)	P2 (voltage)
USB-connection, USB-adapter, or Li-Po battery (3.7V)	 jumper on 2 & 3	 jumper on 1 & 2
Powered from Target (5V)	 jumper on 1 & 2	 jumper on 1 & 2
Powered from Target (3.3V)	 no jumper	 jumper on 2 & 3

Accessories

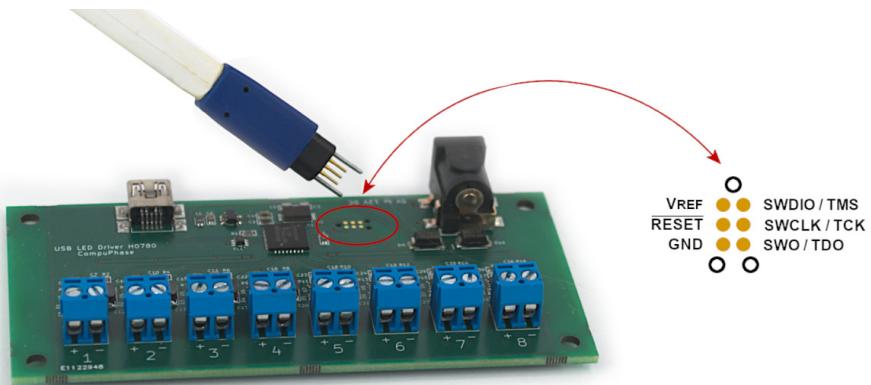
The Black Magic Probe has a 2×5 -pins 1.27 mm pitch debug header for connection to the target, and if the target board has the same connector, it can be connected with the provided ribbon cable. For the other cases, an adapter board or “break-out” board is needed.

A common adapter board is one where that converts between the 10-pin Cortex Debug header and the 20-pin JTAG header. An example of a break-out board is pictured below; it makes the signals of the Cortex Debug available on a single-row 7-pin header (of the ten pins of the Cortex Debug header, three are ground and one is not-connected, so seven pins cover all signals).

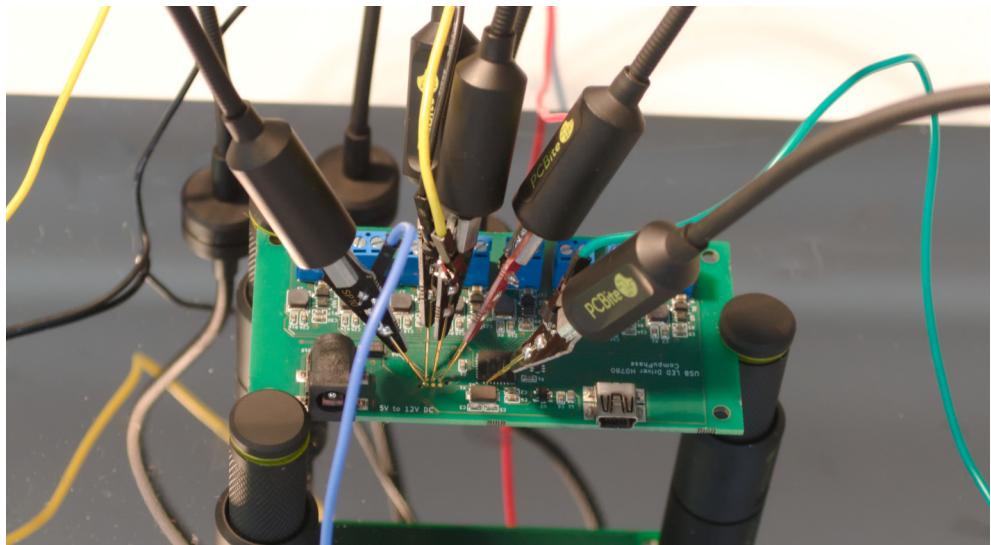


Our favourite debug connector is the decal for the tag-connect cable. This cable has a plug with six pogo-pins, plus three fixed pins that serve to align the plug. The benefit of the tag-connect cable is that it requires less space on the target board than for most other connectors, and that the matching “connector” on the target board is simply a decal. For the target board, the added cost for the programming & debugging connector is therefore zero.

The tag-connect lacks the TDI pin, and hence the tag-connect cable is not suitable for JTAG scanning purposes.



The freeconnect project is an open-source design for a set of pogo-pin connectors that are compatible with the tag-connect cable. See chapter [Further Information](#) on page 119 for a link to the project.



We have found a set of needle probes an indispensable accessory for debugging, especially after a mishap. Like downloading code that inadvertently disables the SWD port. To restore access to the target microcontroller, you will then need to put it up in bootloader mode — as described in section [Design for Debugging](#) on page 27. With a few needle probes on the test pads, or directly on microcontroller pins, you can do so conveniently. We have good experience with the PCBite probes by Sensepeek, again see chapter [Further Information](#) on page 119 for a link to the project.

Information on [page 119](#) for a link.

The Black Magic Probe comes without an enclosure, but if you have access to a 3D printer, it is recommended to print one. An enclosure gives electrical insulation (the Black Magic Probe has a series of exposed test pads at the bottom), as well as mechanical protection. Especially the header for the Cortex Debug connector is somewhat fragile. A few printable designs of enclosures are freely available, see chapter [Further Information](#) on [page 119](#). It feels fitting to print these enclosures in black, but you are of course free to choose any colour.



Likewise, designs for 3D printed enclosures for the ctxLink probe are available on Sid Price's GitHub page. When using ctxLink with a rechargeable battery, an enclosure is recommended, because it protects the battery as well. The ready-to-print STL files are for a Lithium Polymer (Li-Po) "503562"-style battery, which stands for 5.0 mm thick, 35 mm wide and 62 mm long. When using a different battery size, you may need to adjust the design of the enclosure; the design files for AutoDesks Fusion 360 are provided.

The Li-Po battery itself is also a useful accessory for the ctxLink probe, especially for those situations where it is cumbersome to pull a (long) cable to the remote target. The ctxLink probe has a 2-pin JST PH-series connector for the battery (2 mm pitch). For the polarity, see the picture at [page 14](#). The ctxLink probe charges the battery with a constant current of 500mA. Since charging current of Li-Po batteries should not exceed 1C, where C is the capacity in Ampere-hours, the deduction is that a 500mAh capacity is the minimum to be suitable for ctxLink. For a prolonged lifetime of the battery, it is recommended to charge at 0.5C. Therefore, a 1000mAh battery (or higher) is recommended.

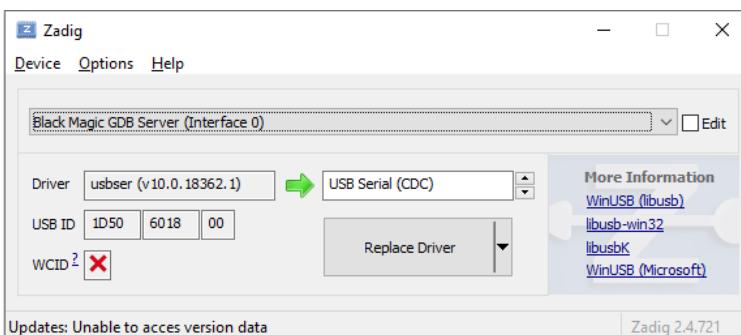
Depending on the project, a galvanic isolation adapter for the USB port may be advisable. The Black Magic Probe provides ESD-protection on the pins of the Cortex Debug header and the UART, but these pins are also rated for an absolute maximum voltage of 6V. If you are working on boards that carries voltages well above that 6V limit (e.g. a LED-driver for eight white LEDs in series may exceed 30V), there is the risk that this voltage breaks out to the low-voltage logic of the board. A multimeter probe that slips and shorts out two pins, may be enough to burn a chip on the target board. If you are unlucky, it may also burn an attached debug probe, but with an isolator you will *not* burn the USB port of your laptop or workstation.

Setting up the Black Magic Probe

Details for adding the Black Magic Probe to your workstation as a USB device, depend on the operating system that you are using, and the data link that you use. This chapter therefore starts with three sections: [Microsoft Windows \(USB set-up\)](#), [Linux \(USB set-up\)](#), and [Wi-Fi set-up for ctxLink](#). You can skip the sections not relevant to you.

Microsoft Windows (USB set-up)

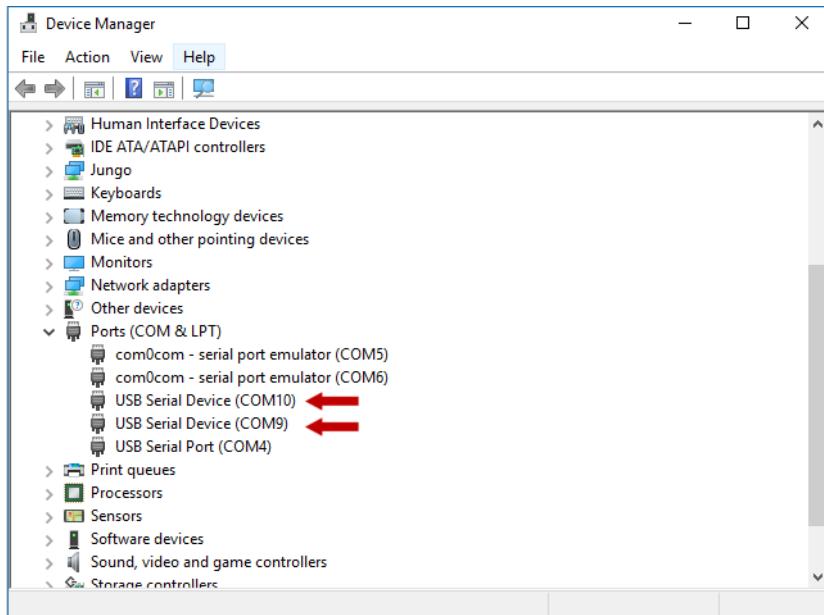
On connecting the Black Magic Probe to a USB port on a workstation, four devices are added. The principal ones are two (virtual) serial ports (COM ports). One of these is for gdbserver and the other is the generic TTL-level UART. The other two are vendor-specific interfaces for firmware update (via the DFU protocol) and trace capture.



On Windows 10, no drivers are needed (a class driver is built-in and automatically set up). Earlier versions of Microsoft Windows require that you install an “INF” file that references the CDC class driver that Microsoft Windows has already installed (“usbser.sys”). A suitable INF file can be found on the project site for the Black Magic Probe, as well as with this book.

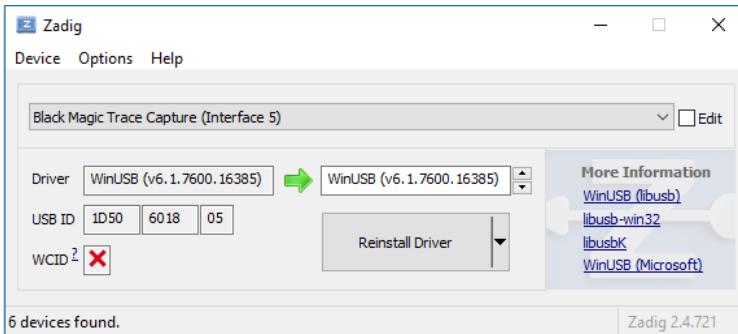
Alternatively, you can set up the CDC driver for the Black Magic Probe with the free utility “Zadig” by Akeo Consulting (see also [Further Information](#) on [page 119](#)). When using Zadig, you need to set up both interfaces 0 (“Black Magic GDB Server”) and 2 (“Black Magic UART Port”) to “USB Serial (CDC)”. You may need to first select List All Devices in the Options menu to see the interfaces of the Black Magic Probe.

Once the CDC driver is configured, two COM ports are assigned to the Black Magic Probe. You can find out which ports in the Device Manager, where they are listed under the item “Ports (COM & LPT).” Alternatively, you can run the BMScan utility on the command line (this is one of the utilities that comes with this book).



Note that in Windows 10, as we are using the built-in CDC driver, the name for the Black Magic Probe interfaces is the generic “USB Serial Device” (see the red arrows in the picture above).

For trace capture and for firmware update, the two generic interfaces of the Black Magic Probe must be registered as either a WinUSB device or a libusbK device.¹ The most convenient way to do so is by running the aforementioned “Zadig” utility (see [Further Information on page 119](#)).



You need to register both interfaces 4 (“Black Magic Firmware Upgrade”) and 5 (“Black Magic Trace Capture”) separately. Both are on USB ID 1D50/6018.

¹ More on the choice between WinUSB and libusbK on the next page.

You may need to first select List All Devices from the Options menu, to make the Black Magic Probe interfaces appear in the drop-down list of the Zadig utility.

For firmware update, you should also register the DFU interface (in DFU mode, USB ID 1D50:6017) as a WinUSB or libusbK device. This interface is hidden until the Black Magic Probe switches to DFU mode. To force the Black Magic Probe in DFU mode, keep the push-button (next to the USB connector) pressed while connecting it to the USB port of the workstation. The red, orange and yellow LEDs will blink in a pattern as a visual indication that the Black Magic Probe is in DFU mode. When you launch the Zadig utility at this point, the interface will be present.

Note that the Black Magic Probe has different USB IDs (VID:PID pairs) in DFU mode versus normal mode (“run mode”). In DFU mode, the ID is 1D50:6017; in run mode it is 1D50:6018.

The choice between WinUSB and libusbK depends on the PC-hosted software that you wish to use for trace capture. The firmware upgrade tool dfu-util (see [Updating Black Magic Probe Firmware on page 103](#)) supports both WinUSB and libusbK. The debugger front-end and trace viewer that accompany this book also support both WinUSB and libusbK, and in this case WinUSB is preferred (because it is pre-installed). The Windows port of the Orbuculum tool-set (see [Monitoring Trace Data on page 68](#)), however, is based on libusb and requires the libusbK driver.

Linux (USB set-up)

After connecting the Black Magic Probe to a USB port, two virtual serial ports appear. One of these is for gdbserver and the other for the generic TTL-level UART. Since the Black Magic Probe implements the CDC class, and Linux has drivers for CDC class devices built-in, no drivers need to be set up.

The device paths for the serial ports are /dev/ttyACM* where the “*” stands for a sequence number. For example, if the Black Magic Probe is the only virtual serial port connected to the workstation, the assigned device names will be /dev/ttyACM0 and /dev/ttyACM1.

You can find out which ttyACM devices are assigned to the Black Magic Probe by giving the dmesg command (in a console terminal) shortly after connecting the Black Magic Probe (see also the arrows in the picture below). Alternatively, you can run the BMScan utility from inside a terminal (BMScan is a companion tool to this book).

```

thiadmer@thinkcentre: ~
File Actions Edit View Help
thiadmer@thinkcentre: ~
ration="mknod" profile="/usr/bin/evince-thumbnailer" name="/home/thiadmer/.cache/thumbnails/normal/a2a3e5a3a665600c5d9b235cc5032418.png" pid=5021 comm="evince -thumbnail" requested_mask="c" denied_mask="c" fsuid=1000 ouid=1000
[ 9003.719555] audit: type=1400 audit(1560938604.609:82): apparmor="DENIED" operation="mknod" profile="/usr/bin/evince-thumbnailer" name="/home/thiadmer/.cache/thumbnails/normal/9974470d2a9e9916825daa685a8a3f63.png" pid=5025 comm="evince -thumbnail" requested_mask="c" denied_mask="c" fsuid=1000 ouid=1000
[ 9003.822613] audit: type=1400 audit(1560938604.713:83): apparmor="DENIED" operation="mknod" profile="/usr/bin/evince-thumbnailer" name="/home/thiadmer/.cache/thumbnails/normal/f207ee2ecbaae58143e04198a3576608.png" pid=5029 comm="evince -thumbnail" requested_mask="c" denied_mask="c" fsuid=1000 ouid=1000
[15089.612036] usb 5-1: new full-speed USB device number 3 using uhci_hcd
[15089.810061] usb 5-1: New USB device found, idVendor=1d50, idProduct=6018, bcdDevice= 1.00
[15089.810066] usb 5-1: New USB device strings: Mfr=1, Product=2, SerialNumber=3
[15089.810069] usb 5-1: Product: Black Magic Probe
[15089.810073] usb 5-1: Manufacturer: Black Sphere Technologies
[15089.810076] usb 5-1: SerialNumber: 7BB180B4
[15089.816169] cdc_acm 5-1:1.0: ttyACM0: USB ACM device ←
[15089.819138] cdc_acm 5-1:1.2: ttyACM1: USB ACM device ←
thiadmer@thinkcentre: ~$ 

```

To be able to access the serial ports, a non-root user must in most cases be included in the dialout group. To add the current user to the group, use:

```
sudo usermod -a -G dialout $USER
```

After this command, you need to log out and log back in, for the new group assignment to be picked up. Reportedly, some distributions use the plugdev group rather than the dialout group.

No driver needs to be installed for the firmware update and trace capture interfaces, but if you wish to use those features as a non-root user (so without needing sudo), a file with udev rules must be installed. For firmware update, it may not be a burden to use sudo, as you will update the Black Magic Probe's firmware only occasionally, but trace capture is a valuable debugging tool for everyday use.

When you copy the file `55-blackmagicprobe.rules` (printed below) into the directory `/etc/udev/rules.d`, it allows any user to access the trace capture interface of the Black Magic Probe.

```
# Standard mode
ACTION=="add", SUBSYSTEM=="usb_device", SYSFS{idVendor}=="1d50", SYSFS{idProduct}=="6018", MODE=="0666"
ACTION=="add", SUBSYSTEM=="usb", ATTR{idVendor}=="1d50", ATTR{idProduct}=="6018", MODE=="0666"

# DFU mode
ACTION=="add", SUBSYSTEM=="usb_device", SYSFS{idVendor}=="1d50", SYSFS{idProduct}=="6017", MODE=="0666"
ACTION=="add", SUBSYSTEM=="usb", ATTR{idVendor}=="1d50", ATTR{idProduct}=="6017", MODE=="0666"
```

The provided udev rules file does not configure stable device names for the ttyACM devices for the Black Magic Probe. If so desired, add the following lines to the rules file (55-blackmagicprobe.rules):

```
SUBSYSTEM=="tty", ATTRS{interface}=="Black Magic GDB Server", SYMLINK+="ttyBMPGDB"  
SUBSYSTEM=="tty", ATTRS{interface}=="Black Magic UART Port", SYMLINK+="ttyBMPUart"
```

After adding the udev rules, you must reload the rules (or alternatively: refresh the session by logging out and logging in again).

```
sudo udevadm control --reload-rules  
sudo udevadm trigger
```

Wi-Fi set-up for ctxLink

When ctxLink is connected to a workstation via USB, the set-up is the same as for the Black Magic Probe. To use the Wi-Fi interface for debugging, the first issue to decide on is how to power the ctxLink. The options are to use a net adapter with a USB-micro connector, a 3.7V Li-Po battery, or to power ctxLink from the target. See [page 15](#) for setting the jumpers for the power selection of ctxLink.

The “mode” LED periodically performs a blink sequence, which indicates both the Wi-Fi status and the battery status. See the picture at [page 14](#) for the “mode” LED and the associated button.

Pulses per sequence	Status
none (LED off)	Wi-Fi not active, power good
1	Battery low
2	Connected to a Wi-Fi access point
3	WPS configuration in progress
4	HTTP Provisioning in progress

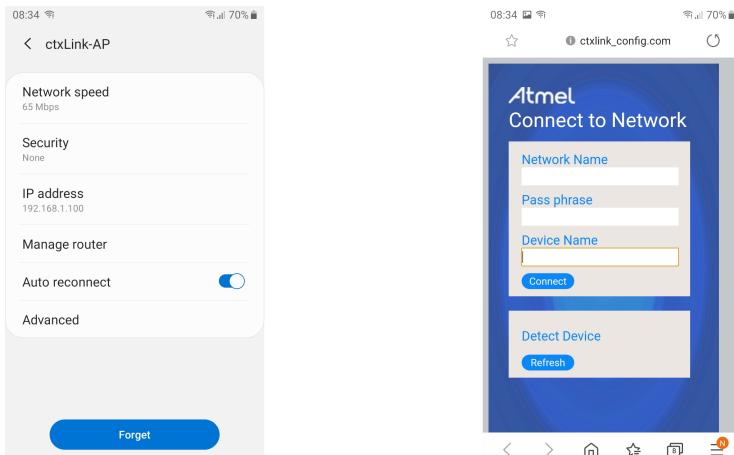
The “mode” button enables you to activate either of the Wi-Fi configuration modes, or to cancel a Wi-Fi configuration, by pressing it for a particular duration: 4 seconds for WPS configuration, 6 seconds for HTTP Provisioning, 7 seconds or more to cancel either configuration mode.

To add the ctxLink to a wireless LAN via WPS (Wi-Fi Protected Set-up), first start the WPS function at the access point (e.g. the wireless router). This is typically done by pushing a button at the router. Then press the “mode” button on the ctxLink for 4 seconds (more accurately: between 3 and 5 seconds). On release of the button, the mode LED will blink in sequences of 3 pulses until the set-up completes. Note that an access point typically shuts WPS off after 2 minutes, so you have to start WPS configuration on the ctxLink fairly quickly after starting it on the access point.

If WPS is not an option, the alternative is to use HTTP Provisioning. With this method, you temporarily set up the ctxLink as an access point, after which you connect to that access point with a laptop or smartphone. You will then be presented with a form that allows you to enter the SSID and pass-phrase of the Wi-Fi access point that ctxLink must connect to.

The first step is to press the “mode” button for 6 seconds (to be precise: between 5 and 7 seconds). Then, use a laptop or smartphone to connect to the new access point with the name “ctxLink-AP.” Possibly you will be asked to confirm that you want to log in. On the form that appears next, fill in (or select) the network name (“SSID”) of the access point and the pass-phrase (those are the SSID and pass-phrase of the Wi-Fi router that you want ctxLink to connect to), and click on the “connect” button. After a short while, the “mode” LED of the ctxLink should start blinking in sequences of two pulses, as an indication of a successful connection.

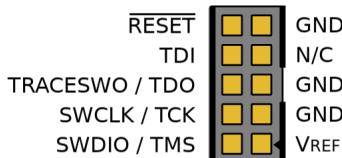
On some systems, the form for the network name and pass-phrase does not automatically pop up. To open it explicitly, you can tap on the ctxLink-AP link, which gives you the screen at the left (in the picture below). In this screen, choose the option “Manage router,” to arrive at the form, as shown at the right. Alternatively, you can open a browser and open 192.168.1.1.



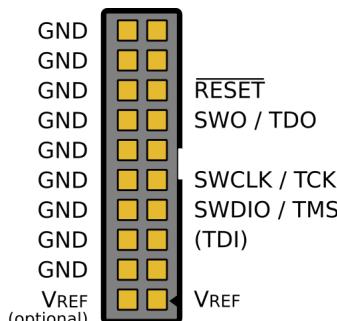
Once connected to the access point, the ctxLink acquires an IP address via DHCP. One way to retrieve this address is to look it up in the list of the DHCP server (typically the wireless router). The ctxLink announces itself as ctxLink_0001 to the DHCP server. In case an access point does not show the device names, the MAC address of the ctxLink is printed on the Wi-Fi module. Alternatively, you may be able to use the BMScan utility to scan the network for ctxLink devices; see [Checking the Set-up on page 24](#) for more information.

Connecting the Target

The Black Magic Probe has a 2×5 -pins 1.27 mm pitch IDC header. This is the Cortex Debug header for JTAG and SWD. If your target board has the same connector, the two can be readily connected with the provided ribbon cable. Otherwise, you will need an adapter board or a break-out board.



When using a break-out board, you may get away with wiring less than the seven distinctive pins/signals from the Cortex Debug header. Of the pins on the debug connector, SWCLK, SWDIO and GND are essential. These must always be connected to the target. The RESET pin is strongly recommended, especially for downloading firmware to Flash memory. The VREF pin should in most cases be connected as well, because the Black Magic Probe uses the target's voltage level at this pin to shift the level on the signal lines to this same voltage. The alternative is to drive VREF to 3.3V from the Black Magic Probe (see the [monitor tpwr](#) command on [page 44](#)). Finally, the TRACESWO pin is for debug tracing, which requires support code in your firmware, and the TDI line is used for JTAG scan, not for SWD debugging.



When the target board has a 20-pin JTAG header, you can also use a break-out board and wire the pins individually. For debugging, the TDI pin is not used, and therefore it does not need to be connected in most cases.

Checking the Set-up

When the Black Magic Probe is connected to a USB port, the green and orange LEDs (labelled “PWR” and “ACT” respectively) should be on. The ACT LED may be bright or dim, depending on the firmware version and the operating system.

If you have not checked which serial port the Black Magic Probe uses for its gdbserver, run BMScan on the command line.

```
d:\Tools> bmscan  
Black Magic Probe found:  
gdbserver port: COM9  
TTL UART port: COM10  
SWO interface: {9A83C3B4-0B99-499E-B010-901D6C2826B8}
```

The above command works for a ctxLink debug probe connected to USB too. For a ctxLink that is set up for Wi-Fi, you can scan the local network for the debug probe with the following command:

```
d:\Tools> bmscan ip  
ctxLink found:  
IP address: 192.168.0.214
```

To check whether the drivers were installed correctly, launch GDB from the command line. You should be using the GDB that was build for the architecture that matches the microcontroller (typically arm-none-eabi). On the “(gdb)” prompt, type (where you replace “*port*” with the COM port for gdb-server):

```
(gdb) target extended-remote port
```

In Microsoft Windows, when the port is above 9, the string “\\.\\” must be prefixed to the port name. So, COM port 10 is specified as “\\.\\com10”. On Linux, the device path for the port must be used, like in “/dev/ttyACM0”.

There is no need to configure the baud rate or other connection parameters; what the operating system presents as a serial port is a USB connection running at 12 Mbits/s, irrelevant of what baud rate it is configured to.

After setting the remote port in GDB, the orange LED (“ACT”) may increase in brightness (this depends on the firmware version of the Black Magic Probe). This LED responds to the DTR signal set by GDB; this was a physical line on the RS232 port, but now just a command on a virtual serial port.

The next step is to scan for the target microcontroller. There are two ways to do this: swdp_scan for microcontrollers supporting SWD and jtag_scan for devices supporting only JTAG.

```
(gdb) monitor swdp_scan  
Target voltage: 3.3V  
Available Targets:  
No. Att Driver  
1 LPC11xx
```

The output shows the driver name for the microcontroller. Note that multiple devices may be returned, for both the SWD scan (using the SW-DP protocol) and the JTAG scan (JTAG devices may be daisy-chained).

The command also shows that the target is not yet “attached” to `gdbserver` (otherwise, there would be a “*” in the “Att” column of the target list). Attaching the target is done with the `attach` command.

```
(gdb) attach 1
Attaching to Remote target
```

At this point, the Black Magic Probe is attached to GDB and you can proceed to download firmware and/or to start debugging it, which is the topic of the next chapter starting at [page 29](#).

Running Commands on Start-up

The above commands have to be repeated on each debugging session. On start-up, GDB reads a file called `.gdbinit` and executes all commands in it. This file is read from the “home” directory on Linux, and from the path set in the `HOME` environment variable in Microsoft Windows (this environment variable is not set by default, so you may need to create it, see also section [GDB on Microsoft Windows](#) on [page 114](#)).

Following the examples in this chapter, a suitable `.gdbinit` file could be:

```
target extended-remote com9
monitor swdp_scan
attach 1
```

If the Black Magic Probe is not yet connected when starting GDB, or if the operating system decided to assign the Black Magic Probe to a different serial port, the above start-up code will fail. GDB aborts parsing the `.gdbinit` file on the first error, so the remainder of the file is not executed either. My recommendation is, therefore, to only add user-defined commands in `.gdbinit`.

```
define bmconnect
  if $argc < 1 || $argc > 2
    help bmconnect
  else
    target extended-remote $arg0
    if $argc == 2
      monitor $arg1 enable
    end
    monitor swdp_scan
    attach 1
  end
end
```

```
document bmconnect
    Attach to the Black Magic Probe at the given serial port/device.
        bmconnect PORT [tpwr]
            Specify PORT as COMx in Microsoft Windows or as /dev/ttyACMx in Linux.
            If the second parameter is set as "tpwr", the power-sense pin is driven
            to 3.3V.
end
```

The above definition gives you a shorthand for conveniently connecting to the Black Magic Probe with a single command.

```
(gdb) bmconnect com9
```

Other settings can be added to the .gdbinit too. If you have per-project settings, these can be in a secondary .gdbinit file in the current directory. GDB will load the “current directory” .gdbinit file when adding the following command in the “home” .gdbinit file:

```
(gdb) set auto-load local-gdbinit
```

You can also load a GDB “command file” explicitly with the “source” command, as below (and you can include such statements in the .gdbinit to load auxiliary command files):

```
(gdb) source ../share/orbcode/gdbtrace.init
```

Design for Debugging

Like almost any other debug probe, the Black Magic Probe can be used for Flash memory programming as well as for debugging the code that runs from Flash memory. For the development cycle, this is very convenient: you build the code and then load it into the target and into the debugger in a single flow.

However, it is common for microcontrollers that several functions are shared on each single pin. If the code redefines one of the pins for SWD to some other function, by design or by accident, the debugging interface will stop functioning. If the code redefines the pins quickly after a reset, the Black Magic Probe may not have a chance to regain control of the SWD interface, even after a reset. The result is that not only the code cannot be debugged any more, but also that no new code can be flashed into the microcontroller.

Depending on your microcontroller, a way to circumvent this is to enable the option `connect_srst` in the Black Magic Probe (see [page 44](#) for the command description). The Debug Access Port of the ARM Cortex is designed such that it may stay active while the remainder of the microcontroller is in reset, so that a debug probe can attach to it. This is precisely what the `connect_srst` option does: it pulls the reset pin on the connector low while performing a

SWDP scan, as well as during the `attach` command. Whether or not the ARM Cortex debug port is enabled during reset, depends on the microcontroller, however. For example, NXP's low-end microcontrollers with a Cortex M0(+) core use the `RESET` pin to switch between JTAG and SWD (disabling SWD while `RESET` is low).

As an alternative, you can often use system-specific pins to force a microcontroller into boot mode. The LPC series of microcontrollers from NXP have a `BOOT` pin that forces the microcontroller into bootloader mode when it is pulled low on reset (or on power cycle). The STM32 series from STMicroelectronics have two boot pins for the same purpose —though `BOOT0` must be pulled high, rather than low. Bootloader mode is designed for Flash programming over a serial port or USB, but the side effect is that it blocks the firmware from running. As a result, the pins for SWD have not been redefined and you can now start GDB and attach to the target (after which you can upload new firmware). The recommendation for PCBs with an LPC or STM32 microcontroller is therefore to branch out the “boot” pin(s) to a jumper or a test pad, so that you can recover from an accidental pin redefinition.

If the pin redefinition of the SWD pins is by design, because you need these pins for other purposes, this will thwart your ability to debug the code. If possible, arrange the design such that the SWD pins are used for a non-essential function. Then, you can implement the firmware such that it redefines the SWD pins only when not running under control of a debugger. While debugging, you will miss the functionality that would otherwise be driven by SWD pins, but you can debug the rest.

Two methods are available for the firmware to detect whether it is running under a debugger. The first is to test that the low bit of the *Debug Halting Control & Status Register* (`DHCSR`) is set. This works on a Cortex M3/M4/M7 microcontroller, however; on the Cortex M0/M0+ microcontroller architecture, this register is not accessible from firmware (it is accessible from the JTAG/SWD interface).

```
if ((CoreDebug->DHCSR & 1) == 0) {  
    /* not running under a debugger, free to redefine pins */  
}
```

The alternative is to have a weak pull-up on the `SWCLK` pin and probe it (as a general-purpose I/O pin) on start-up. The Black Magic Probe pulls the clock line low (provided that it senses a voltage on the `VREF` pin). This does require some pin juggling, though: you first have to configure the `SWCLK` pin as an “input” I/O pin (with a pull-up) to be able to read it, and depending on the value read, either quickly change it back to `SWCLK` pin, or set it to its intended configuration. Also, if the pin is connected to other circuitry that drives the pin low, this trick won't work.

Debugging Code

Debugging code for embedded systems has its own challenges, in part due to the way that microcontroller projects differ from typical desktop applications. Some commands of GDB are skipped over in almost every book because they are not relevant for desktop debugging. This chapter focuses on the commands that are pertinent to the Black Magic Probe and ARM Cortex targets. It is therefore more an addendum to books/manuals on debugging with GDB, than a replacement of them.

As mentioned in [The Debugging Pipeline \(page 5\)](#), you will probably prefer a front-end to do any non-trivial debugging. Below is a screen-capture of gdbgui connected to the Black Magic Probe, and ready to debug "blinky."

The screenshot shows the gdbgui interface running in a browser window. The title bar says "gdbgui - gdb in a browser". The address bar shows "127.0.0.1:5000". The main area displays the C source code for the "blinky" program:

```
32 int main(void)
33 {
34     uint32_t led_ioport, led_iobit;
35
36     if (SysTick_Config(SystemCoreClock / 1000)) { /* Setup SysTick Timer */
37         while (1); /* Capture error */
38     }
39
40     led_ioport = IOPORT(PIN_LED);
41     led_iobit = IOBIT(PIN_LED);
42     LPC_GPIO->DIR[led_ioport] |= led_iobit; /* LED = output */
43
44     for ( ;; ) {
45         LPC_GPIO->SET[led_ioport] = led_iobit; /* turn LED on */
46         mdelay(500);
47         LPC_GPIO->CLR[led_ioport] = led_iobit; /* turn LED off */
48         mdelay(500);
49     }
50 }
```

To the right of the code editor is a sidebar with several sections:

- threads**: Shows a single thread named "Remote target, id 1" which is currently "stopped".
- local variables**: Shows variables `curTicks`, `dlyTicks`, and `dlyTicks@entry`.
- expressions**: A collapsed section.
- Tree**: A collapsed section.
- memory**: A collapsed section.

At the bottom of the sidebar, there is a memory dump for address 8, which shows "no memory to display".

The bottom half of the screen shows the GDB console output:

```
No. Att Driver
1    LPC11xx
attach 1
Attaching to program: d:\\products\\usbkey\\source\\obj\\blinky.elf, Remote target
0x0000033a in mdelay (dlyTicks=dlyTicks@entry=500) at blinky.c:27
27   while ((msTicks - curTicks) < dlyTicks)

(gdb) enter gdb command. To interrupt inferior, send SIGINT.
```

The gdbgui front-end is a fairly thin graphical layer over GDB: you have to type most commands in the console. However, the limited abstraction from GDB is actually an advantage. Front-ends typically aim at desktop debugging, and so the set of commands specific to embedded code are not wrapped in

dialogs and pop-up menus.

Yet, while we recommend the use of a front-end with GDB, the commands and examples in this chapter use the GDB console. While a front-end may provide a more convenient way to perform some task, each will have its own interface for it. The GDB console is a common denominator for all GDB-based debuggers.

Prerequisite Steps

On every launch of GDB, it has to connect to the Black Magic Probe, scan for the attached target and attach to it. Unless you are using the [BMDebug](#) front-end that handles these steps automatically, they have to be given through the console.

```
(gdb) target extended-remote COM9
Remote debugging using COM9
(gdb) monitor swdp_scan
Target voltage: 3.3V
Available Targets:
No. Att Driver
 1      LPC11xx
(gdb) attach 1
Attaching to Remote target
0x0000033a in ?? ()
```

These commands can be wrapped in a user-defined command in a `.gdbinit` file, see [Running Commands on Start-up](#) on page 26. In that case, you would type only a single command:

```
(gdb) bmconnect COM9
Target voltage: 3.3V
Available Targets:
No. Att Driver
 1      LPC11xx
0x0000033a in ?? ()
```

Loading a File and Downloading it to the Target

The first step in running code in a debugger, is to generate debug symbols while building it. The GNU GCC compiler (and linker) use the command line option `-g` for that purpose. The default format for the debug symbols is typically DWARF, and this would also be the preferred format.

You can specify the target executable file on the command line when launching GDB, but alternatively, you set it with the `file` command. The filename

may be a relative or full path, with a / as the directory separator (this is of notice to users of Microsoft Windows, where directories are usually separated with a "\").

```
(gdb) file blinky.elf
A program is being debugged already.
Are you sure you want to change the file? (y or n) y
Reading symbols from blinky.elf...done.
(gdb) load
Loading section .text, size 0x7da lma 0x0
Start address 0xd8, load size 2008
Transfer rate: 6 KB/sec, 669 bytes/write.
```

Note that the GDB load command downloads only the executable code to the target. The ELF file contains debug symbols, which makes the executable file much larger than when the code is compiled without debugging information. However, the size of the code that is downloaded to the target remains the same; the debug symbols are not transferred.

Flash Memory Remap

For the LPC microcontroller series, an additional step is recommended before the load command. NXP designed the microcontrollers such that the bootloader always runs on reset (or power-up). The bootloader then samples the boot pin, verifies whether there is valid code in the first Flash sector, and jumps to it if it checks out. The conflict is: the ARM Cortex starts running at the reset vector stored at address 0, which must initially point to ROM (where the bootloader resides) and then to Flash memory (where the user code sits). The LPC microcontrollers have the feature to remap address range 0...511 to either Flash, RAM or ROM via either the SYSMEMREMAP or the MEMMAP register. According to the NXP manuals, after a reset, the register is initialized such that address 0 maps to Flash memory. However, that is not what happens: the SYSMEMREMAP (or MEMMAP) register is initially 0 (remap to bootloader ROM) and the bootloader then modifies it to map to Flash before jumping to the user code in Flash. However, when the microcontroller is halted by the debug probe, SYSMEMREMAP is still 0. Then, if you download new code in the microcontroller, the bottom 512 bytes will be sent to ROM, and be lost.

The fix is to force mapping the SYSMEMREMAP register to the appropriate value from GDB (as is apparent, SYSMEMREMAP is a memory-mapped register). The example below is for the LPC8xx, LPC11xx, LPC12xx and LPC13xx series.

```
set mem inaccessible-by-default off
set {int}0x40048000 = 2
```

For convenience, the above can be wrapped in a user-defined command in the .gdbinit file, see [Running Commands on Start-up](#) on page 26:

```

define mmap-flash
    set mem inaccessible-by-default off
    set {int}0x40048000 = 2
end

document mmap-flash
    Set the SYSMEMREMAP register for NXP LPC devices to map address 0 to
    Flash.
end

```

You would then give the command `mmap-flash` before using the `load` command. The address of the `SYSMEMREMAP` register (and the value to set it to) is different in other series in the LPC microcontroller range.

NXP series	Register	Address	Flash map
LPC800, LPC1100, LPC11U00, LPC1200, LPC1300	SYSCON.SYSMEMREMAP	0x40048000	2
LPC1500	SYSCON.SYSMEMREMAP	0x40074000	2
LPC1700	SCB.MEMMAP	0x400FC040	1
LPC2100, LPC2200, LPC2300, LPC2400	SCB.MEMMAP	0xE01FC040	1
LPC4300	M4MEMMAP	0x40043100	0

The “`mmap-flash`” snippet in `.gdbinit` therefore needs to be adapted for the particular microcontroller as well. A more complete version of the `mmap-flash` user-defined command is in the `.gdbinit` file that comes with this book.

Reset Code Protection

On the STM32Fxx family of microcontrollers, the `load` command may give the following error:

```
(gdb) load
Error erasing flash with vFlashErase packet
```

This implies that **readout protection** (RDP) is set in the option bytes. No new code can be downloaded unless the option bytes are erased first —which in turn wipes the entire Flash memory. To check whether code read protection is set, use the `monitor` command.

```
(gdb) monitor option
usage: monitor option erase
usage: monitor option <addr> <value>
0xFFFF800: 0x5aa5
0xFFFF802: 0x00ff
0xFFFF804: 0x00ff
0xFFFF806: 0x00ff
0xFFFF808: 0x00ff
0xFFFF80A: 0x00ff
0xFFFF80C: 0x00ff
0xFFFF80E: 0x00ff
```

If the first option word is anything other than 0x5aa5, the code is read protected. As a side note, option bytes are written in pairs: value and complement. The option is only valid if the complement matches. For unprotected code, the value for the option is 0x55, and its complement is 0xaa.

If the value is in the first option word is 0x33cc, RDP Level 2 is in effect. RDP Level 2 cannot be undone. Any other value (so neither 0x55aa, nor 0x33cc) indicates RDP Level 1.¹ RDP Level 1 can be reverted to Level 0 (i.e. “unprotected”) by erasing the option bytes. As stated, the side effect of clearing RDP Level 1 is that the Flash memory is fully erased (which is, of course, by intent).

To erase the option bytes, again use the `monitor` command, but now with the “`erase`” option.

```
(gdb) monitor option erase
0xFFFFF800: 0x0000
0xFFFFF802: 0x0000
0xFFFFF804: 0x0000
0xFFFFF806: 0x0000
0xFFFFF808: 0x0000
0xFFFFF80A: 0x0000
0xFFFFF80C: 0x0000
0xFFFFF80E: 0x0000
```

After erasing the option bytes, the microcontroller must be power-cycled to reload them (the output of the `option erase` command does not reflect the true values of the option bytes; after reset, you will see that the first option word is actually set to 0x5aa5 instead of 0x0000). Note that GDB will lose the connection to the target on a power-cycle, so you must rescan and re-attach to the target again.

To set code protection on a STM32Fxx microcontroller, by the way, use the command below, followed by a power-cycle.

```
(gdb) monitor option 0xfffff800 0x00ff
```

When code protection is enabled on the LPC microcontroller series, Flash memory must also be fully erased before new firmware can be downloaded. These microcontrollers do not use option bytes, however. Instead, you must erase the Flash memory either by a `monitor` command, or by using a tool that talks directly to the Black Magic Probe (there is no GDB command for erasing Flash memory). For target drivers that support it, use the following monitor command:

```
(gdb) monitor erase_mass
```

¹ However, RDP Level 2 disables the SW-DP port, so you will not be able to list the option bytes.

Only a subset of the target drivers of the Black Magic Probe support this command. See also [Using the BlackMagic Flash Programmer](#) on page 97 as an alternative tool for downloading firmware via the Black Magic Probe. The [BM-Flash](#) utility has an option to erase the entire Flash memory, even for target drivers that do not support the monitor `erase_mass` command.

However, setting code protection on the LPC series disables the SWD interface after a reset. After that, it is no longer possible to remove code protection using the Black Magic Probe. Instead, your options are to erase Flash memory either via the serial bootloader (ISP), or from within your firmware (that is, you've added a piece of self-destruct code to the firmware, which is triggered by a special command or special status on power-up).

Verify Firmware Integrity

To verify that the code in the microcontroller is the same as the code loaded in GDB, you can use the `compare-sections` command. This command also lets you verify that downloading code was successful.

```
(gdb) compare-sections  
Section .text, range 0x0 -- 0x7d8: matched.
```

There is a caveat with the LPC series of microcontrollers from NXP: these microcontrollers require a checksum in the vector table at the start of the Flash code. The checksum can only be calculated at or after the link stage, but the GNU linker is oblivious of this requirement. Instead, firmware programmers calculate and set the checksum while downloading, and the Black Magic Probe is no exception. The upshot is that `compare-sections` will now always return a mismatch on the first section, since its contents were changed on the flight while downloading it.

To fix `compare-sections`, the checksum must be set in the vector table in the ELF file after the link phase. The Black Magic Probe will calculate it again during downloads, despite that it is already correctly set—but that is harmless. After downloading, the code in the microcontroller will be identical to the code in the ELF file.

```
elf-postlink lpc11xxx blinky.elf
```

The program `elf-postlink` is one of the utilities that come with this book.

Starting to Run Code

The `run` command starts to run the loaded code from the beginning. If you have not set any breakpoints, the code runs until it is interrupted through `Ctrl+C`. The `start` command sets a temporary breakpoint at function `main` and then runs; the program will therefore stop at `main`.

```
(gdb) start
The program being debugged has been started already.
Start it from the beginning? (y or n) y

Temporary breakpoint 1 at 0x348: file blinky.c, line 33.
Starting program: c:\Source\blinky\blinky.elf
Note: automatically using hardware breakpoints for read-only addresses.

Temporary breakpoint 1, main () at blinky.c:33
33      {
```

Note the mention of the automatic use of hardware breakpoints. With the help of the Black Magic Probe, GDB indeed inserts a hardware breakpoint on the break command.

Getting help and information

Oft overlooked, but the `help` and `info` commands are among the most useful for a command-line tool like GDB.

<code>help</code>	Show a list of topics you can get help on.
<code>help topic</code>	Show help on a topic, usually including a list of relevant commands. For the list of topics, type "help" without any parameter.
<code>help command</code>	Show the syntax and parameters of the command, plus a brief description of its purpose or function. For a list of commands, type "help all".
<code>info</code>	Show a long list of topics you can get information about.
<code>info topic</code>	Show the information GDB has on the topic. Topics can range from variables and arguments, stack, targets and the sources that built it, breakpoints, watchpoints, etc.

The difference between the `help` and `info` commands is that `help` gives information on how to do things in GDB, and `info` informs you about the status of GDB or the program loaded into it. For example, "`help break`" gives you a page that describes how to set a breakpoint, whereas "`info break`" lists the breakpoints that are set, plus the properties of each of these.

Listing Source Code

The commands listed below are a subset of the full GDB command & parameter set for listing the source code of a target. These are the most common commands.

<code>list line</code>	Show the source code around the given line number in the current source file.
------------------------	---

<code>list file:line</code>	Show the source code in the file with the name in the first parameter, and around the line number in the second parameter.
<code>list function</code>	Show the source code starting at the given function.
<code>list</code>	Show the next lines (below the current position). You can optionally add a + as a parameter ("list +").
<code>list -</code>	Show the preceding lines (above the current position).
<code>info sources</code>	List the names of the source files for the target executable.
<code>info line *address</code>	Print the line number and source file associated with the address. The address parameter must start with "0x" if it is in hexadecimal.

Downloading code into the microcontroller

These commands were covered in section [Loading a File and Downloading it to the Target](#) on page 30.

<code>file path</code>	Set the path of the ELF file to debug (if it was not already given on the GDB command line).
<code>load</code> <code>load filename</code>	Download the code sections of the ELF file into the microcontroller. In the common case, no parameter is given with the <code>load</code> command, in which case the file set with the <code>file</code> command is downloaded. However, you may specify the file to download explicitly.
<code>compare-sections</code>	Check whether the code in the microcontroller matches the ELF file being debugged; see section Verify Firmware Integrity on page 34.
<code>flash-erase</code>	Erases the Flash memory regions on the target.

Stepping and Running

These are the basic commands needed for debugging. Several of these commands were already informally introduced in earlier sections.

<code>start</code>	Start or re-start the program and break at function main. If no function "main" exists, it is the same as the <code>run</code> command.
<code>run</code> <code>continue</code> (can be abbreviated to c)	Start or re-start the program (from the beginning). Continue running (from the current execution point). A count may follow the command, but it is only relevant if code stopped due to a breakpoint. If present, the breakpoint is ignored the next "count" times it is hit. This is particularly useful in when the breakpoint is inside a loop: the command "continue 10" will run 10 more iterations before stopping at the breakpoint again.

step <i>(can be abbreviated to s)</i>	Step a single source line, step into functions if the current execution point is at line with a function call. A count may follow the command. If present, the command repeats the step “count” times.
next <i>(can be abbreviated to n)</i>	Step a single source line, step over functions (if there is a function call at the current execution point). A count may follow the command. If present, the command repeats the step “count” times.
until <i>(can be abbreviated to u)</i>	Run until a source line is reached that is below the current line (this command is intended for stepping out of loops). Alternatively, you can set a line number after the until command, and then it runs until that line is reached.
finish <i>(can be abbreviated to fin)</i>	Continues execution until it steps out of the current function, then stops at the location from where the function was called.

The `step` command will not step *into* functions *without* symbolic information, such as a function from the standard library. Instead, `step` will step *over* the function call, and behave identical as `next` in this case. You can also instruct GDB to always step over particular functions, with the `skip` command (it *skips* stepping *into* the function). The purpose of `skip` is easiest explained with an example:

```
▶ transmit_data(array, setup_connection());
```

When a function has a call to another function in its parameter list, the `step` command will step *into* the nested function first. In this example, it will step *into* `setup_connection()` and only go *into* `transmit_data()` afterwards. Suppose we want to step *into* `transmit_data()`, but that needing to step through `setup_connection()` first is a chore. This is where you want to mark `setup_connection()` to be skipped.

The `skip` command is very flexible; the two most common variants are below.

<code>skip function name</code>	Skip the stated function (or skip the current function if no name is given).
<code>skip file name</code>	Skip all functions in the stated file (or skip all functions in the current file, if no name is given).

When stepping through optimized code, the current line may jump back and forth on occasion, because the compiler has re-arranged the generated machine code. See section [Debugging Optimized Code](#) on page 55 for details.

Altering execution flow

In case that you need to break out of an endless loop, or a case where you know that continuing running the remainder of the function will do no good, you can alter the execution flow.

<code>jump line</code> <code>jump file:line</code>	Start running at the given location. You can use this command to break out of an endless loop, or to jump back a few lines to re-examine the control flow.
<code>return</code> <code>return expression</code>	Skips to the return address of the current function (without executing the code between the current location and the return point). The program stays in stopped state.

In brief: `jump` is like `continue`, but starting from a different location; and `return` is like `finish`, but without executing the code. See also to set command on [page 40](#), because you can often also change control flow (or exit a loop) by changing a variable.

Breakpoints and watchpoints

When creating a breakpoint or watchpoint, it gets assigned an ID. This is simply a unique number to identify the breakpoint or watchpoint. Several of the commands listed below take the breakpoint ID as a parameter.

<code>break line</code> (can be abbreviated to b)	Set a breakpoint at the line number in the current source file.
<code>break file:line</code> <code>break function</code> <code>tbreak ...</code>	Set a breakpoint at the line number in the specified source file. Set a breakpoint at the start of the named function. Sets a one-time breakpoint, which auto-deletes itself as soon as it is reached. The <code>tbreak</code> command takes the same parameter options as the <code>break</code> command.
<code>watch expr</code>	Set a watchpoint, which causes a break as soon as the expression changes. In practice, the expression is typically the name of a variable, so that GDB halts execution of the program as soon as the variable changes.
<code>rwatch expr</code>	A watchpoint that triggers when the variable (or memory location that the expression points to) is <i>read</i> . This requires hardware breakpoints.
<code>awatch expr</code>	A watchpoint that triggers on <i>access</i> —either read or write. It requires hardware breakpoints.
<code>info break</code>	Show the list of breakpoints and watchpoints, together with the sequential index numbers (i.e. the breakpoint “IDs”) that each breakpoint got assigned.
<code>delete</code> <code>delete id ...</code>	When given without parameters, this command deletes all breakpoints. Otherwise, if one or more breakpoint IDs follow the command (separated by spaces), the command deletes the breakpoints with those IDs.
<code>clear</code>	Without parameters, this command deletes the breakpoint that is at the current code execution point. The primary use is to delete the breakpoint that was just reached.

clear <i>line</i> clear <i>file:line</i> clear <i>function</i> disable <i>id</i> ...	Delete a breakpoint on the given line or function. It allows the same options as the break command.
enable <i>id</i> ...	Disables the breakpoints with the given IDs. There may be one or more IDs on the command list (separated by spaces).
cond <i>id expr</i>	Enables the breakpoints with the given IDs. There may be one or more IDs on the command list (separated by spaces). You may also use enable once to enable the breakpoints, but disable them when they are reached. Attaches a condition to the breakpoint with the given ID. The condition is what you would write between the parentheses of an "if" statement in the C language. For example: <pre>cond 3 count == 5</pre> causes breakpoint 3 to only halt execution when variable count equals 5 (assuming, of course, that variable count is in scope). When the expression is absent on this command, the condition is removed from the breakpoint (but the breakpoint stays valid).
command <i>id</i> ... end	Sets a command list on the given breakpoint. These commands are executed when the breakpoint is reached. It can be used, for example, to automatically print out the stack trace on arriving at the breakpoint. See section Tracing with Command List on Breakpoints on page 71 .

For embedded development, enabling and disabling breakpoints (and watchpoints) is all the more useful, because hardware breakpoints & watchpoints are a scarce resource. Most Cortex-M microcontrollers offer 6 hardware breakpoints and 2 hardware watchpoints. What counts for the Black Magic Probe, is not the number of breakpoints that have been set, but the number that is active. When you need more breakpoints than the microcontroller offers, you keep them defined, but disable the ones that are not immediately relevant for the next step in debugging the code.

For a hardware watchpoint, the data type of the expression cannot be wider than the word size of the microcontroller. For example, the word size is 32-bit on an ARM Cortex-M microcontroller, and the expression to watch can therefore be up to four bytes wide.

Setting a breakpoint plus a condition on a breakpoint may be combined in a single step. To do so, put the keyword "if" followed by condition expression at the end of the break command. For example:

```
break blinky.c:168 if count == 5
```

The Cortex-M microcontrollers can also break on specific exceptions or interrupts. An exception trap is set with the `monitor vector_catch` command, see [page 45](#). When the exception is caught, the microcontroller will halt on the first instruction of the exception/interrupt handler.

Examining Variables and Memory

In addition to the commands below, most front-ends show a variable's value when hovering the mouse cursor over it. Front-ends typically also allow setting "variable watches" (which is the equivalent to the `display` command), and they may also automatically list all local variables and their values (the equivalent of the `info locals` command). The `gdbgui` front-end even allows you to add a graph for numeric variables, to give you a visualization of the value of the variable over time.

<code>print var</code> <code>print /fmt various</code> <code>print var@count</code> (can be abbreviated to p)	Show the contents of the variable. GDB can parse C-language expressions to show array elements or dereferenced pointers, like in: <code>print var[6]</code> show the value of an array element <code>print *ptr</code> dereference the pointer and show the value The format to print the variable in (e.g. decimal, hexadecimal, or other) is a single letter; see the list on page 41 . The "var@count" syntax interprets var as the start of an array and prints count elements.
<code>info args</code>	Show the names and values of the function arguments.
<code>info locals</code>	Show the names and values of all local variables.
<code>ptype var</code>	Show the type information of the variable.
<code>display var</code> <code>display /fmt var</code> (can be abbreviated to disp)	Watch the variable. Show the variable's value each time that the execution is halted. The format to print the variable in (e.g. decimal, hexadecimal, or other) is a single letter; see the list on page 41 .
<code>undisplay num</code> (can be abbreviated to undisp)	Remove the watch with the given sequence number.
<code>x address</code> <code>x /options address</code>	Display the memory at the given address. The options start with a slash, followed by zero or more digits, and then followed by one or two letters. The digits are the count of elements, the first letter is the format and the second letter the size of each element in bytes.
<code>set var=value</code> <code>set addr=value</code>	Set the variable to the value, or store the value at the address. You can use C-style type-casts on the address to specify the size of the memory field.

The GDB print command records each value that it prints in its “value history.” Each entry in that value history is labelled. The first label is \$1 and the number is incremented for each successive print command. You can use these labels on subsequent print commands.

While the print command is typically used to show variables, it is able to evaluate C-style expressions. As such, you can use the GDB print command as a built-in calculator.

```
(gdb) p sampleDelay  
$1 = 50  
(gdb) p $1 / (double)ticksPerSecond  
$2 = 0.05000000000000003
```

The x command displays the memory at any given address. The address can be an expression that evaluates to an address—which includes variables. If no options are given, GDB uses its defaults, and at start-up the default display format is: a single 32-bit value displayed in hexadecimal. The number of elements (bytes or words) to display can be set after a slash, for example “x /4 0x1234” displays four elements starting at the given address. This element count then also becomes the new default for the x command.

The print, display and x commands each allow a format specification behind the slash (or, for the x command, behind the count). The format code consists of a single letter. In case of the x command, a second letter may be added to indicate the size of each item.

o	octal
x	hexadecimal
d	decimal
u	unsigned decimal
t	binary
f	floating point
a	address
c	character
s	string (zero-terminated)
i	instruction
b	size modifier: byte (8-bit value)
h	size modifier: halfword (16-bit value)
w	size modifier: word (32-bit value)
g	size modifier: “giant” word (64-bit value)

As is the case for the count of elements in the x command, the given display format becomes the default for any subsequent x command. Unless the count of elements is specified together with the display format, the element count is reset to 1.

Peripheral registers are memory-mapped in the ARM architecture, but by default, GDB won’t show data outside the range for program and data memory.

That is, the address of a peripheral register is considered an invalid memory address. To view peripheral registers, first issue the following command:

```
set mem inaccessible-by-default off
```

After this command, GDB considers any address outside the memory map as RAM. The [BMDebug](#) front-end (see [page 47](#)) automatically runs this command, and other GDB front-ends may do so too. Alternatively, you can include the command in `.gdbinit` —the `.gdbinit` file was covered in section [Running Commands on Start-up on page 26](#).

The Call Stack

A stack frame stores the local variables, arguments and the return address for each sub-routine (or function). The scope of local variables and arguments is restricted to the sub-routine that they are declared in.² Stack frames form a list, that a debugger can walk up and down. Moving up the stack frame allows you to look at the local variables in the routine that the current routine (that contains the execution point) was called from.

<code>backtrace num</code> (can be abbreviated to <code>bt</code>)	Show a list with the call-stack that lead to the current execution point. The call stack is optionally limited to the given number of levels.
<code>up</code>	Move to the frame one higher in the call-stack, which is the frame that contains the call to the current frame. You can go up multiple levels by adding the count, as a parameter.
<code>down</code>	Move back to a lower frame. You can go down multiple levels by adding the count, as a parameter.
<code>frame idx</code> (can be abbreviated to <code>f</code>)	Move to the given frame index (the <code>backtrace</code> command prints these index numbers). The <code>frame</code> command <i>without</i> parameter prints the active frame index.

GDB numbers the stack frames sequentially, starting from zero for the sub-routine that the current execution point is in. With the command “`frame 0`”, you will return to the frame that corresponds with the execution point.

After changing to a different stack frame (with the `up`, `down` or `frame` commands), commands like `info locals` will reference to the local variables of that frame. This may help you in determining what conditions caused the call to the function that contains the execution point.

² This is a simplification —more accurately, local variables have a scope that runs from their declaration to the end of the compound block that the declaration appears in.

Inspecting Machine Code

GDB is primarily used as a source-level debugger, but at times, you may want to look at what happens at the CPU level.

<code>disassemble</code> <code>disassemble start,end</code> <code>disassemble /s</code> (can be abbreviated to <code>disas</code>)	When used without start & end arguments, it shows the assembly code of the function that the execution point is in. The alternative is to specify an address range (start, end) to disassemble. The /s option mixes the assembly code with the source code, which often makes it easier to follow the assembly code.
<code>set disassemble-next-line on</code> <code>set disassemble-next-line off</code>	When GDB halts execution, it shows the source code line that it stopped on (if that source code line is available). When the option <code>disassemble-next-line</code> is switched on, GDB will in addition show the disassembly for the instruction at the execution point.
<code>stepi</code>	Like the step command, see page 37 , but stepping a single instruction (instead of a source line).
<code>nexti</code>	Like the next command, see page 37 , but stepping a single instruction (instead of a source line).
<code>info registers</code>	Print the names and values of the registers of the microcontroller.

GDB allows you to refer to a register anywhere where it expects a variable, by prefixing it with a \$. In other words, you can add a watch on register r0 by giving the command:

```
(gdb) display $r0
```

Debug Probe Commands

GDB has a pass-through command to configure or query a gdbserver implementation: `monitor` (`monitor` can be abbreviated to `mon`). Whatever follows the keyword `monitor` is passed to the `gdbserver`, in our case the embedded `gdbserver` in the Black Magic Probe.

The supported `monitor`-commands are listed below. Note that some of these commands are only available on particular microcontroller series; if this is the case, the applicable microcontroller series is noted.

<code>help</code>	Show a summary of the commands that the debug probe supports (basically this list, but restricted to commands relevant to the detected target).
<code>version</code>	Show the current version of the firmware and the hardware.
<code>jtag_scan</code>	Scan the devices on the JTAG chain.

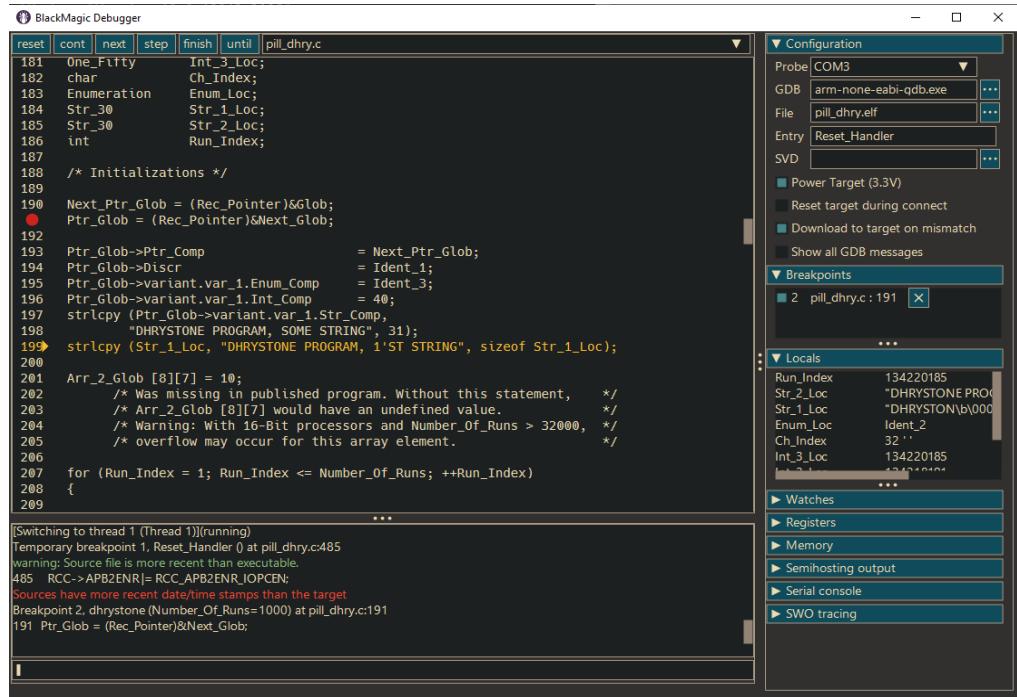
<code>swdp_scan</code>	Scan for Serial Wire Debug devices (using the SW-DP protocol). The command prints the I/O voltage and the list of targets. See also the <code>tpwr</code> command (below) for the I/O voltage and the <code>targets</code> command for the device list.
<code>auto_scan</code>	Scan either JTAG or SWD protocols <small>firmware 1.8</small> Performs a <code>jtag_scan</code> first, followed by an <code>swdp_scan</code> if the JTAG scan does not detect devices.
<code>traceswo</code> <code>traceswo rate</code>	Enable the SWO capture pin to for trace capture. The rate parameter is the bitrate of the SWO trace protocol. It is used only for asynchronous encoding and on firmware versions 1.7 and later it defaults to 115.2 kbps. Note that the native Black Magic Probe only supports Manchester encoding. The ctxLink probe supports only asynchronous encoding. For capturing SWO output using the BlackMagic Debugger front-end, see the Trace Views on page 51 . Decode SWO output in the Black Magic Probe. <small>firmware 1.7</small> The trace output is transmitted over the virtual UART, so that it can be viewed on a serial terminal. The SWO channels to decode can be appended as a space-separated number list to the command. If absent, all channels are active.
<code>targets</code>	Show the detected targets. This is the same list as the one returned by the <code>jtag_scan</code> and <code>swdp_scan</code> commands. For each detected microcontroller, it displays the driver (the driver is often specific to a microcontroller family).
<code>tpwr enable</code> <code>tpwr disable</code>	Enables or disables driving the VCC pin on the 2×5 pin header to 3.3V. See page 23 for the pin-out of the connector. When the Black Magic Probe drives the VCC pin, it can power the target (maximum current: 100mA). The VCC pin must always be driven, either by the target or by the Black Magic Probe, because the voltage at this pin is also used by level shifters on the logic pins on the connector. The default is that the VCC pin must be driven by the target. A special case is to not wire the VCC pin between the Black Magic Probe and the target. The VCC pin must now also be driven by the Black Magic Probe, and the level shifters are therefore set to 3.3V TTL levels.
<code>connect_srst</code>	Enables or disables the option to keep the target microcontroller in reset while scanning and attaching to it. See the discussion at page 27 .
<code>hard_srst</code>	Resets the target by briefly pulling the RESET pin low on the 2×5 pin header (see page 23 for the connector).

<code>morse</code>	When the Black Magic Probe encounters an error that it cannot handle otherwise, it will start to blink the red LED (labelled “err”) in a Morse code pattern. In case your Morse code decoding skill is a little rusty, you can instead use this <code>morse</code> command to return the error message in plain text on the GDB console. But in fact, the only such error message is “target lost.”																
<code>frequency value</code>	Sets maximum SWCLK frequency. The value is in Hz, but may use a “k” or “M” suffix, to multiply the value by one thousand or one million respectively; for example, 1M stands for 1 MHz. A lower frequency for the SWD protocol may be needed depending on the target microcontroller, or on the wiring between the debug probe and the target. When no parameter is given, the command returns the active frequency (however, in firmware 1.8, it is returned in hexadecimal).																
<code>halt.timeout delay</code>	Set time to wait for device to halt. The time to wait for the Cortex-M core to halt, so that the debug probe can attach to it. This value is in milliseconds. The default is 2000 ms.																
<code>heapinfo hb hl sb sl</code>	Set semihosting heapinfo values. Four values in hexadecimal format must follow this command, for the values of the heap base, the heap limit, the stack base and the stack limit respectively. The startup code in the newlib C library (and possibly others) uses a semihosting call to get the heap and stack space, if running under a debugger. This command lets you override the defaults for the heap and the stack.																
<code>vector_catch enable vec</code> <code>vector_catch disable vec</code>	Break on specific exceptions. The first parameter must be enable or disable. The second parameter must be the exception for which the “catch” must be enabled or disabled. It is one of: <table> <tr><td>hard</td><td>Hard fault</td></tr> <tr><td>int</td><td>Interrupt/exception service errors</td></tr> <tr><td>bus</td><td>Bus fault</td></tr> <tr><td>stat</td><td>Fault state error</td></tr> <tr><td>chk</td><td>Divide by zero, misaligned memory access, etc.</td></tr> <tr><td>nocp</td><td>Missing coprocessor (on coprocessor instruction)</td></tr> <tr><td>mm</td><td>Memory Manager fault</td></tr> <tr><td>reset</td><td>Core reset</td></tr> </table> Cortex-M0 and M0+ microcontrollers only support “reset” and “hard” fault. A hard reset cannot be caught, though.	hard	Hard fault	int	Interrupt/exception service errors	bus	Bus fault	stat	Fault state error	chk	Divide by zero, misaligned memory access, etc.	nocp	Missing coprocessor (on coprocessor instruction)	mm	Memory Manager fault	reset	Core reset
hard	Hard fault																
int	Interrupt/exception service errors																
bus	Bus fault																
stat	Fault state error																
chk	Divide by zero, misaligned memory access, etc.																
nocp	Missing coprocessor (on coprocessor instruction)																
mm	Memory Manager fault																
reset	Core reset																

<code>erase.mass</code>	Erase entire flash memory.	<i>LPC17xx, LPC4300 Cortex-M4, EFM32 Gecko, nRF51xxx series, SAMD, STM32Fxx, STM32L4xx</i>
<code>erase.bank1</code>	Erase entire flash memory in bank 1.	<i>STM32L4xx</i>
<code>erase.bank2</code>	Erase entire flash memory in bank 2.	<i>STM32L4xx</i>
<code>reset</code>	Reset target.	<i>LPC4300 Cortex-M4</i>
<code>mkboot bank</code>	Make flash bank bootable. The parameter is the bank number, 0 or 1.	<i>LPC4300 Cortex-M4</i>
<code>serial</code>	Print the MCU serial number.	<i>EFM32, Gecko, SAMD</i>
<code>unsafe</code>	Allow programming the security byte. The parameter must be enable or disable.	<i>Kinetis</i>
<code>read</code>	Read target device parameters. The parameter is one of: help Show brief help on the command hwid The hardware identification number fwid The pre-loaded firmware ID deviceid The unique device ID deviceaddr The device address	<i>nRF51xxx series</i>
<code>gpnvm.get</code>	Get value of the GPNVM register.	<i>SAM3N, SAM3S, SAM3U, SAM3X, SAM4S</i>
<code>gpnvm.set</code>	Set bit in the GPNVM register. The first parameter is the bit number. The second parameter is the value for the bit (0 or 1).	<i>SAM3N, SAM3S, SAM3U, SAM3X, SAM4S</i>
<code>lock_flash</code>	Lock Flash memory against accidental change.	<i>SAMD</i>
<code>unlock_flash</code>	Unlock Flash memory.	<i>SAMD</i>
<code>user_row</code>	Print the user row from Flash.	<i>SAMD</i>
<code>mbist</code>	Run the “Memory Built-In Self Test” (MBIST).	<i>SAMD</i>
<code>option erase</code> <code>option address value</code>	Set option bytes. The first syntax is “option erase” to erase the option bytes. If read protection set in option bytes, erasing it implicitly erases the entire Flash memory. The second syntax is “option address value” which stores a value at the given address.	<i>STM32Fxx, STM32L0x, STM32L1x, STM32L4xx</i>
<code>eeprom</code>	Set values in EEPROM (non-volatile memory). The first parameter is one of: byte 8-bit value halfword 16-bit value word 32-bit value The second parameter is the address in the EEPROM. The third parameter is the value (with the size as specified in the first parameter).	<i>STM32L0x, STM32L1x</i>

The BlackMagic Debugger Front-end

The BMDebug utility is a front-end for GDB that is designed for the Black Magic Probe and ctxLink. It handles the [Prerequisite Steps](#) described on [page 30](#) on start-up: it locates the debug probe and attaches to it, optionally provides power to the target, verifies whether the code in the microcontroller matches the file loaded in GDB and optionally downloads it on a mismatch. Another distinguishing feature of BMDebug is that it combines traditional debugging with run-time tracing.



Starting up

After loading an ELF file, BMDebug stops at function `main` in that code. You may set an alternative function as the entry point of the executable. If the entry point function (typically “`main`”) cannot be found, BMDebug keeps the microcontroller in halted state, so that you can set a breakpoint at some code of interest before giving the `run` command (or pressing the “`cont`” button).

Unlike the [BMFlash](#) utility (see [page 97](#)), the BMDebug front-end is not able to calculate the header checksum for the LPC microcontroller family *before* uploading it. This is because BMDebug is based on GDB (it is a “front-end”), whereas [BMFlash](#) is independent of GDB. As a consequence, GDB (and thereby BMDebug) will always see a CRC mismatch between the (LPC-specific) ELF

file loaded in the debugger and the one downloaded in the target, and re-download it at every run. To avoid this, include a call to `elf-postlink` on the ELF file as part of the build process (e.g. the Makefile). See the discussion of the `elf-postlink` utility in section [Verify Firmware Integrity \(page 34\)](#) for more information on the checksum for LPC microcontrollers. The other option is to disable automatic download of the ELF file in BMDebug (option “Download to target on mismatch” in the “Configuration” section in the sidebar), and instead use the `load` command to explicitly download new code. See also section [Edit-Compile-Debug Cycle on page 54](#).

GDB Console and Command Line

BMDebug is a “thin” front end: it has controls and shortcuts for the basic operations of a debugger, but more advanced commands (like adding a condition to a breakpoint) need to be typed as a command. The output of those commands typically appears in the GDB console.

BMDebug has the GDB console and the command line (for input to GDB) in the bottom-left section. The GDB console shows the output of GDB. Some messages from GDB are filtered out by default. You can set the option “Show all GDB messages” in the “Configuration” section in the sidebar to see all output.

The command line keeps a history of commands that are typed in. The `Ctrl+↑` and `Ctrl+↓` key combinations scroll through earlier commands on the command line, and `Ctrl+R` key pair searches in the command history for matching the text.

Another feature is autocompletion of commands or parameters, on the `TAB` key. This is especially convenient when the parameter of a command is a file or a function: just type in the first few letters of the function or file name and press `TAB`. Pressing `TAB` multiple times cycles through all candidates.

Source View

The source view shows the execution point with a rightwards pointing triangle in the left margin. The execution point is the line that will be executed next when continuing execution.

The “cursor line” in the source view is highlighted. You can freely move the cursor line, using the standard keys for cursor movement (using Arrow Up/Down, Page Up/Down, `Ctrl+Home` and `Ctrl+End`). Every time the target microcontroller stops, BMDebug sets the cursor line to the execution point. Alternatively, you can also run to the cursor line with the button `Until` (or use function key `F7`).

When stepping through code, the source view automatically switches to the source file that the execution point is in. You can select any source file from

the drop-down list in the button bar above the source view. Alternatively, you can use the `list` command in the console line (see section [Listing Source Code on page 35](#)). For switching to another source file, the file extension may be omitted. For example, the following command will load the file `blinky.c` or `blinky.cpp` (whichever is available).

```
(gdb) list blinky
```

You may also type a function name or a line number as the parameter to the `list` command. This will make the source view jump to that line or to the start of the given function. The `Ctrl+G` key combination is a shorthand for the `list` command, and if you type only the first letters of a file or function, pressing TAB will autocomplete the name.

An additional command is provided to search for text in the source file that is displayed (this is not a GDB command, but one specific to BMDebug).

<code>find text</code>	Finds the first occurrence of the text starting from the cursor line. The search wraps from the bottom of the text to the top. The text search is case-insensitive. The key combination <code>Ctrl+F</code> inserts the <code>find</code> command on the edit line.
<code>find</code>	Repeats the last search. Function key <code>F3</code> is a shorthand for this action.

Running code

The button bar above the source code view has the essential functions for running and stepping through code. The names of most buttons reflect the GDB command that it executes: the “step” button executes a `step` command, and the “finish” button lets GDB execute a `finish` command.

The exception is the “reset” button, which reloads and restarts the target firmware, and then runs up to `main`.

All buttons have a function key associated with them. For example, `F10` does a next command (step over) and `F11` does a `step` command (step into). A tooltip on each button shows the equivalent function key.

Breakpoints

You can set a breakpoint either by clicking in the left margin in the *source view*, or with function key `F9`, or with a `break` command in the console.

When clicking in the source view, clicking a second time on an existing breakpoint *disables* the breakpoint (rather than removing it). To remove the breakpoint, you need to click on it a third time (while staying on the line with the mouse cursor). The breakpoints can also be toggled between enabled and disabled in the Breakpoints view in the sidebar.

When debugging code in Flash ROM, you can set as many breakpoints as you like, but only a limited number can be enabled at any time (most Cortex-M microcontrollers provide 6 hardware breakpoints).

The `break` command (see [Breakpoints and watchpoints on page 38](#)) can also be used on the console line. The command line allows you to set temporary breakpoints and watchpoints as well.

Viewing Variables and Registers

Hovering over a variable name in the source view shows the current value of that variable in a tooltip. Note that the tooltip only appears when the target is in a stopped state.

The “Locals” view in the right sidebar shows all local variables that are currently in scope. GDB uses heuristics (based on the variable type) to choose whether to display integer variables in decimal, hexadecimal or other. In BMDebug, you can select a different display format after right-click of the mouse on the value.

The “Watches” view in the sidebar shows the current value of all expressions that have been added to it. The expression can be as simple as the name of a variable, but it may include redirections or arithmetic operations. When adding a watch, all variables that are mentioned in the expression are evaluated in the active scope. The expression of the watch retains this scope. When stepping into a sub-routine or function, the Watches view keeps showing the watches in the scope that the watch was declared in.

A watch can be added by typing the expression in the edit field in the Watches view and clicking on the  button. You can also use the `display` command in the console line (see section [Examining Variables and Memory on page 40](#)). The BMDebug front-end handles the `display` and `undisplay` commands internally.

Standard registers of the microcontroller can be views in the “Registers” view in the right sidebar. For peripheral registers, BMDebug supports “System View Description” files (SVD files). SVD files contain the definitions of the core and peripheral registers of the microcontroller. When an appropriate SVD file is loaded, hovering over a register name in the source view shows the value of the register; likewise, you can add a watch to a peripheral register. That said, this feature depends on the source code and the SVD file to agree on the peripheral and register names. In practice, this means that SVD files combine neatly with CMSIS as the hardware abstraction layer, because the System View Description format is a subproject of CMSIS. When using a different hardware abstraction layer, like libopencm3, SVD files may not be of much use.

Microcontroller manufacturers typically provide SVD files for their microcontrollers, and the CMSIS project comes with the SVDConv utility to generate a C/C++ header file from an SVD file. A collection of SVD files for various brands and series of microcontrollers is available on GitHub, see [Further Information on page 119](#) for the link.

Viewing Assembly Code

BMDebug can show disassembled machine code, interleaved with the source code. It uses its own disassembler (rather than the one in GDB), so that the assembly code can be annotated with peripheral and register names from SVD files (as covered above).

<code>assembly</code>	Switches assembly mode on or off. When used without parameters, the command toggles the mode.
<code>disassemble</code>	The standard GDB disassemble command is redefined to be the equivalent to the assembly command.

When in assembly mode, function keys F10 and F11 step by machine instruction, rather than by source line. Specifically, F10 performs a `nexti` command in assembly mode, and a `next` command when in source mode. Likewise, F11 executes a `stepi` or a `step` command, depending on the mode.

Viewing Memory

Viewing memory at some address that is not related to a symbol in the program, is quite common on microcontrollers. Embedded peripherals are often memory-mapped and a microcontroller may define special memory regions for buffers or queues. GDB has the “x” command that fits this purpose (see [page 40](#)). The BMDebug front-end improves on it by displaying the memory dump in a separate view, and by updating this view at each halting point. It functions like a *watch* on a memory range: bytes or words that have changed since the last refresh are coloured red.

BMDebug supports the same options on the `x` command as GDB, but its defaults are different. Where GDB defaults to displaying a single 32-bit word, BMDebug defaults to displaying sixteen 8-bit bytes.

Trace Views

Three trace views are provided: one for semihosting output, one for a serial monitor, and one for SWO tracing. See chapter [Run-Time Tracing on page 56](#) for more information on tracing.

The view for semihosting is always active, and it requires no configuration (except that the target firmware must be built to send output via the semihosting interface).

The serial monitor and SWO tracing view must be configured through commands on the console line. These commands are specific to the Black Magic Probe and the BMDebug front-end; they are not passed to GDB. Both the serial monitor and the SWO tracing view support [The Common Trace Format](#) (see [page 72](#)), for tracing with reduced overhead.

On the topic of SWO tracing: note that while the BMDebug front-end supports both Manchester encoding and asynchronous encoding. However, the hardware implementation of the debug probe determines which of the two you can use. At the time of writing, the native Black Magic Probe supports only Manchester encoding (version 2.3 of the hardware also supports asynchronous encoding, but firmware support is still pending). The ctxLink probe supports asynchronous encoding (but not Manchester).

<code>serial port bitrate</code> <code>serial bitrate</code>	Open the serial port at the given bitrate (Baud), to monitor received data. If the port name is omitted, the command uses the secondary TTL-level UART of the Black Magic Probe. The protocol settings that the serial port are set to, are: 8 data bits, 1 stop bit, no parity.
<code>serial disable</code>	Disable the serial monitor, closes the serial port.
<code>serial enable</code>	Open the serial monitor with the most recent settings (for port and bitrate).
<code>serial clear</code>	Clear the viewport of the serial monitor (deletes all received text).
<code>serial filename</code>	Set the metadata file for decoding the Common Trace Format (see page 72). The metadata file is in TSDL format. When a metadata file is set, incoming serial data is interpreted as Common Trace Format packets.
<code>serial plain</code>	Disable the Common Trace Format decoding and unload a previously loaded TSDL metadata file.
<code>serial info</code>	Show the current configuration.
<code>trace clock bitrate</code> <code>trace passive</code>	Enable tracing in Manchester encoding. If the clock of the target microcontroller and bit rate are set, the BMDebug front-end configures the target for SWO tracing. The clock and bitrate parameters may have a MHz or kHz suffix. For example, the clock may be specified as either 12mhz or 12000000. The bitrate parameter may also use the “kbps” unit. If “passive” is set as the command parameter, SWO tracing is turned on in the Black Magic Probe, but the target is not configured. Use this option if the firmware of the target configures SWO tracing itself (in code). The parameter “passive” may also be written “pasv”.

<code>trace async <i>clock</i> <i>bitrate</i></code> <code>trace async passive <i>bitrate</i></code>	Enable tracing in Asynchronous encoding with the given clock of the target microcontroller and bit rate. The clock and bitrate parameters are the same as with the preceding command. If “passive” or “pasv” is set as the command parameter, SWO tracing is turned on in the Black Magic Probe, but the target is not configured (see also the preceding command). In the case of asynchronous encoding, the bit rate must still be set for passive mode.
<code>trace disable</code>	Disable SWO tracing.
<code>trace enable</code>	Enable SWO tracing using previously configured settings.
<code>trace clear</code>	Clear the viewport.
<code>trace 8-bit</code> <code>trace 16-bit</code> <code>trace 32-bit</code> <code>trace auto</code>	Set the width of the data in an SWO tracing packet (in relation to trailing-zero compression). This value must match the value that the target uses. The ubiquitous implementation is 8-bit data width (which is the default setting). When the parameter is auto, the debugger derives the data width from the incoming data. See page 9 for more information.
<code>trace <i>filename</i></code>	Set the metadata file for decoding the Common Trace Format (see page 72). When no file is explicitly set, the BMDebug front-end looks for a file with the same base name as the ELF file and a “.tsdl” extension, and it searches in the same directory as the ELF file, as well as in the directories where the source files are.
<code>trace plain</code>	Disable the Common Trace Format decoding and unload a previously loaded TSDL metadata file.
<code>trace channel <i>index</i> enable</code> <code>trace chan <i>index</i> enable</code> <code>trace ch <i>index</i> enable</code>	Enable the display of the given channel (range 0..31).
<code>trace channel <i>index</i> disable</code> <code>trace chan <i>index</i> disable</code> <code>trace ch <i>index</i> disable</code>	Disable the display of the given channel.
<code>trace channel <i>index</i> name</code> <code>trace chan <i>index</i> name</code> <code>trace ch <i>index</i> name</code>	Set a name for the channel marker in the view (the default name is the channel number). Note that when using the Common Trace Format, the channel names are initially set to the “stream” names in the trace metadata.
<code>trace channel <i>index</i> #colour</code> <code>trace chan <i>index</i> #colour</code> <code>trace ch <i>index</i> #colour</code>	Set the background colour of the channel marker. The colour must be in “HTML format” with three pairs of hexadecimal digits following the “#”, in the order R/G/B.
<code>trace info</code>	Show the current configuration and all active channels.

The BMDebug front-end saves target-specific settings, such as the settings for SWO tracing, in a file with the same name as the target ELF file, but with the

added file extension “.bmcfg.” The settings of this file are reloaded when you load the ELF file again in BMDebug. Therefore, to enable SWO tracing and restore all settings and channel configurations from a previous session, the following command is sufficient:

```
(gdb) list trace enable  
Active configuration: Manchester encoding, passive, data width = 8-bit
```

Help and info

BMDebug adds a few topics to the `help` and `info` commands —see [Getting help and information](#) on page 35 for these commands. When typing `help` without parameters, these topics are listed under the sub-head “Front-end topics.”

Edit-Compile-Debug Cycle

While stepping through code or analysing trace output, you may spot something that needs to be fixed. However, you do not need to leave the debugger to edit and re-compile the code. It is recommended that you switch to your editor or IDE and rebuild it, and then reload it in GDB. This way, breakpoints and other settings are preserved. The code still restarts at main, though.

With the BMDebug front-end, the recommended way to reload the ELF file is to use the button “reset” at the top left of the source view, or the key combination `Ctrl+F2`. This button not only reloads the file in GDB, it also downloads the file into the target (provided that the “Download to target on mismatch” option is ticked in the “Configuration” section in the sidebar).

The `gdbgui` front-end keeps all source files cached until the “reload file” button is clicked. Likewise, the BMDebug front-end loads all source files right after GDB loads the debugging symbols for the ELF file and keeps them in memory. As a result, if you edit a source file, those changes will not appear in BMDebug until the ELF file is reloaded (through the “reset” button or `F2`). The rationale for this operation is that it keeps the source code, as presented in BMDebug in line with the debugging information in the ELF file. The upshot is that you can edit the source code for a program without hesitation while continuing to debug it. A pitfall with `gdbgui`, though, is that if you re-run the program (which reloads the symbolic information), but forget to reload each source file (with the “reload file” button), the source and the executable are still out of sync.

Note that when the “Download to target on mismatch” option is disabled in the configuration, the `reset` command or button in BMDebug restarts debugging and reloads the source files, but does not download any changed ELF file to the target. You will need to use the `load` command, or temporarily force reloading with the command:

```
(gdb) reset load
```

Another reset option that you may need in special occasions, such as when the code has accidentally redefined the SWCLK or SWDIO pins, is:

```
(gdb) reset hard
```

This option does a full reset of GDB and either resets or power-cycles the target (depending on whether the “Power Target” option is set in the configuration).

Debugging Optimized Code

When stepping through the code, the current line may on occasion jump over a few lines and then jump back up later. This is especially the case with optimized code. The reason is that GDB steps sequentially through the machine code, and at each point where it stops, it looks up the line number in the source file that matches the address where it stopped. The GCC compiler may have rearranged the code that it generated, in order to get a more optimal result. While it is common advice to compile with optimizations disabled, GDB is actually very capable to debug optimized code —if you can live with an occasional surprising order of execution.

Another optimization that the GCC compiler may perform, is to inline small functions. You may not immediately notice this, because GDB is smart enough to simulate a call to the inlined function when stepping through the code. That is, you can step into an inlined function, even though there isn’t a call in the machine code. What you cannot do, however, is place a breakpoint on the inlined function: the function does not exist as a separate block of instructions. Instead, you must place the breakpoint at the point (or points) where the inlined function is called.

Run-Time Tracing

The standard “stop & stare” style of debugging, where you step through code one line at a time, may not be suitable for an embedded system. When the code hits a breakpoint or is in “step”-mode, the microcontroller stops, and this may be *too little* or *too much* (even both at the same time). The microcontroller may not run in isolation: if it drives a linear actuator, that actuator will continue to run while the MCU is in stopped state, until it reaches a safety end stop —unless that end stop is handled by an interrupt routine on the same MCU, in which case the actuator will run until it damages itself. Stopping the microcontroller does *too little* in this case: it does not stop the linear actuator, but it also does *too much*: it blocks the ISR that handles the safety end stop from running.

The alternative debugging technique for such circumstances is run-time tracing. The goal of tracing is to be non-intrusive: it gives you insight in what the code does *without* interfering with it. Run-time tracing is similar to logging, the differences between the two are mostly due to their distinctive purposes (logging is used by system administrators to review activity of the system; tracing is used by developers to spot software faults). Run-time tracing is also akin to post-mortem analysis, in the sense that you are analysing the code flow (and the logic behind that code flow) after the fact.

This chapter has an overview of the various methods for tracing that the Black Magic Probe offers. Each of these has its own advantages and disadvantages. The next chapter then delves into an efficient binary format and protocol for run-time tracing.

Levels of Tracing

The ARM CoreSight architecture has hardware support for both low-level tracing and high-level tracing. Specifically, the Cortex microcontrollers provide for three trace sources:

- ◊ Instruction trace, which creates a log of every instruction executed by the microcontroller. It is generated by the *Embedded Trace Macrocell* (ETM).
- ◊ Data trace, to monitor changes of variables or memory. It is generated by the *Data Watchpoint & Trace* unit (DWT).
- ◊ Software trace, or “debug message,” which sends out *printf* or *transmit* statements that are embedded in the source code of the firmware. Software trace is also called instrumented trace, because it requires the firmware to be “instrumented” with trace instructions.

The tracing techniques in this chapter mostly fall in the last category: software trace. The exception, in a way, is [Tracing with Command List on Breakpoints](#) (see [page 71](#)) because it does not require instrumenting the source code.

The main drawback of code instrumentation is that it makes the firmware code bigger and run slower. Unless you also build a method to disable tracing dynamically in the production code (the code that you distribute), you will want to remove the trace instrumentation from the production build. It is therefore common that the code instrumentation is implemented with conditionally compiled macros.

Secondary UART

The Black Magic Probe combines the `gdbserver` interface with a TTL-level UART interface (on the same USB connection). If the target board has the TxD and RxD lines of a UART branched out of the microcontroller, and the target does not need the UART for other purposes, you can use that port to output trace messages and capture those on a general purpose serial terminal or monitor,

Sending trace messages over a UART is a boilerplate technique, because it works everywhere: all microcontrollers offer one or more UART peripherals, and (virtual) serial ports on workstations are commonplace too. Other than its ubiquity, a benefit of the UART is that it requires only a single pin —configuring RxD is superfluous for tracing purposes. Of course, this is only valid in the case that you use tracing as your only means of debugging; otherwise, the UART pins are *in addition to* the pins reserved for the JTAG or SWD interface.¹

The RS232 transmission rates are, for today's standards, rather slow. Therefore, there is the risk that tracing slows down the code flow too much, defeating the entire purpose of run-time tracing.

Semihosting

Semihosting uses the debug protocol and interface, so that it does not require extra pins if you already have the JTAG or SWD pins branched out. This

¹ This refers to the number of pins on the microcontroller. With regard to the wiring, the ground wire must be connected in addition to the TxD and (optionally) RxD pins. When using the secondary UART of the Black Magic Probe, the device's power should normally be connected to the VCC pin of the UART connector of the Black Magic Probe —see [page 113](#).

is especially convenient if you are using an ST-Link clone instead of the original Black Magic Probe hardware, because the ST-Link clones have neither a secondary UART for tracing, nor the TRACESWO pin branched out (see [page 61](#)).

Due to additional overhead by the debug probe, semihosting has lower performance than using a UART. Semihosting also requires support from the debug probe and the debugger running on the remote host, but both the Black Magic Probe and GDB provide the necessary support. The source code must furthermore be instrumented with calls to `trace`, `printf` or similar.

At a low level, semihosting works by inserting a software breakpoint (or sometimes a software exception) in the code, followed by a special token value. When the microcontroller reaches that instruction, it halts and signals the debug probe. The debug probe first looks at the address of the break instruction, sees the token, and enters semihosting state. It then analyses two registers, `r0` and `r1`, which carry a command code and a pointer to a parameter block. The debug probe forwards the commands to the debugger (GDB in our case), which runs it and may transmit results back.

The ARM semihosting protocol is extensive and flexible. In principle, it allows the embedded target to relegate console and file I/O to the host. For tracing, only a single command code is relevant (`SYS_WRITE`). The snippet below is a function for transmitting a trace message using semihosting, implemented in GCC.

```
void trace(const char *message)
{
    uint32_t command = 5;      /*SYS_WRITE*/
    uint32_t packet[3] = { 2 /*stderr*/, (uint32_t)message, strlen(message) };
    __asm__ (
        "mov r0, %0\n"
        "mov r1, %1\n"
        "bkpt #0xAB\n"
        :
        : "r" (command), "r" (packet)
        : "r0", "r1", "memory"
    );
}
```

The command code 5 is defined for writing to a file, and file handle 2 (the first word in the packet array) is the predefined handle for “standard error” console output. When calling `trace("Hello world")` from your code (and running it from GDB), this text will be printed on the GDB console.

The reason for writing to file handle 2 (`stderr`) instead of handle 1 (`stdout`) is that when using GDB without a front-end, `stderr` can be redirected to a file or separate terminal (instead of being mixed with GDB console output).

Note, however, that GDB prints errors messages to `stderr` also, so GDB output and trace messages can still wind up interlaced. A front-end may write semihosting output to a separate view or window (regardless of whether it is sent to `stderr` or `stdout`), however in this case, output from the Black Magic Probe itself may also wind up in that view. The [BMDebug](#) front-end shows semihosting output in the “Semihosting output” view, see [page 51](#).

The above snippet is for the ARMv6-M and the ARMv7-M architectures (ARM Cortex M0, M0+ M1, M3, M4 and M7 series). On other architectures, you may need the `SVC` instruction rather than `BKPT`.

Depending on the standard libraries that you use, you may not need to implement a trace function yourself, but simply use `printf()` via semihosting. In particular, the library `librdimon` (part of `newlib`) implements semihosting calls. If you use `newlib`, it is sufficient to add the following option to the linker command line:

```
--specs=rdimon.specs
```

A drawback of semihosting is that if *no* debugger is attached, the software breakpoint triggers a *HardFault* exception —and typically stops the entire device in its tracks. Trace calls via semihosting are therefore typically wrapped inside macros whose definition is conditional on the build: debug versus release, and you must be careful to never run a debug build outside a debugger.

An alternative is to determine at run-time whether a debugger is attached, and adjust the `trace()` function to return straight away if otherwise. On a Cortex M3/M4/M7 microcontroller, this is as easy as testing the lowest bit of the *Debug Halting Control & Status Register* (DHCSR):

```
if (CoreDebug->DHCSR & 1) {
    /* debugger attached */
} else {
    /* not running under a debugger */
}
```

On the Cortex M0/M0+ microcontroller architecture, the CoreDebug registers are only accessible from the JTAG/SWD interface, however, not from the code that runs on the microcontroller. Instead, you can implement a *HardFault* handler to check the cause of the exception and return to the caller if it turns out to be a semihosting call. This way, the `trace()` function still drops on the `BKPT` instruction and still causes a *HardFault* exception (in absence of a debugger), but the *HardFault* handler ignores it and moves the program counter to the instruction behind it.

The *HardFault* handler approach for run-time debugger detection works on all Cortex architectures, it is not restricted to Cortex M0/M0+. On projects build with CMSIS and `libopencm3`, a user-defined exception handler automatically replaces the default implementation, provided that it has the correct

name. This is `HardFault_Handler()` for CMSIS, and `hard_fault_handler()` for `libopencm3`.

```
__attribute__((naked))
void HardFault_Handler(void)
{
    __asm__ (
        "mov r0, #4\n"          /* check bit 2 in LR */
        "mov r1, lr\n"
        "tst r0, r1\n"
        "beq msp_stack\n"      /* load either MSP or PSP in r0 */
        "mrs r0, PSP\n"
        "b get_fault\n"
    "msp_stack:\n"
        "mrs r0, MSP\n"
    "get_fault:\n"
        "ldr r1, [r0,#24]\n" /* read program counter from the stack */
        "ldrh r2, [r1]\n"      /* read the instruction that caused the fault*/
        "ldr r3, =0xbeab\n"   /* test for BKPT 0xAB (or 0xBEAB) */
        "cmp r2, r3\n"
        "beq ignore\n"         /* BKPT 0xAB found, ignore */
        "b .\n"                /* other reason for HardFault, infinite loop */
    "ignore:\n"
        "add r1, #2\n"         /* skip behind BKPT 0xAB */
        "str r1, [r0,#24]\n" /* store this value on the stack */
        "bx lr"
    );
}
```

The way the *HardFault* handler works is slightly convoluted, because the ARM Cortex microcontroller has two stack pointers, for the “main stack” and the “process stack.” When the exception occurred, the microcontroller has pushed a set of registers on the stack, including the program counter, but the first thing the *HardFault* handler must do is to check *which* stack. Once it has the appropriate stack pointer, by testing bit 2 in the LR register, it gets the value of the program counter. The program counter is the address of the instruction that caused the exception, so the handler reads from that address and tests for opcode 0xBE with parameter 0xAB. On a match, it is a semihosting breakpoint and it increments the program counter value on the stack before returning; effectively returning to the instruction that follows the breakpoint. Otherwise, it drops into an infinite loop, just like the default implementation for the *HardFault* handler.

SWO Tracing

The ARM Cortex M3, M4, M7 and A architectures provide a separate pin for tracing system and application events at a high data rate. This is the TRACESWO pin on the Cortex Debug header (see [page 23](#)). The ARM Cortex M0 and M0+ architectures lack support for SWO tracing, but see section [SWO Tracing on the Cortex M0/M0+](#) on [page 66](#) for a workaround.

The SWO Trace protocol allows messages to be transmitted on 32 channels (or *stimulus* ports, per the ARM documentation). This allows you to separate output for different modules in the firmware or to implement different levels of trace detail, because each channel can be individually enabled or disabled. By convention, the last channel (channel 31) is reserved for use by an RTOS. Sending a trace message on a channel that is disabled takes negligible time, and therefore it may be an option to leave the trace calls in the production code.

With CMSIS, a typical implementation of a `trace()` function is as below. Note, however, that the CMSIS function `ITM_SendChar()` is hard-coded to use channel 0.

```
void trace(const char *msg)
{
    while (*msg != '\0')
        ITM_SendChar(*msg++);
}
```

Apart from being limited to channel 0, the above function is also inefficient. With tracing disabled, the function still runs over all characters in the message and calls a function. Moreover, as explained in section [TRACESWO Protocol](#) ([page 8](#)), this protocol transmits packets of 1 to 4 bytes, and it prefixes each packet with a header byte. With the CMSIS implementation of `ITM_SendChar()`, each packet has a payload of only a single byte. As a result, the effective transfer speed of SWO tracing has just been halved (sending one byte in reality sends two: a header byte and a payload byte).

A more flexible and efficient function is below. It starts by checking whether tracing is enabled, both globally and on the chosen channel, so that it doesn't even run through the message string if nothing would be output anyway. If that test drops through, it collects up to 4 characters from the message into a 32-bit word, before storing it in the FIFO of the *Instrumentation Trace Macro-cell* (ITM). The FIFO is accessed via the register PORT, which is in fact an array of 32 registers. Before storing every next packet in the FIFO for the trace subsystem, the function waits in a while loop until the FIFO has space to hold the new packet.

```

void trace(int channel, const char *msg)
{
    if (((ITM->TCR & ITM_TCR_ITMENA) != 0UL && /* ITM tracing enabled */
        (ITM->TER & (1 << channel)) != 0UL) /* ITM channel enabled */
    {
        /* collect and transmit characters in packets of 4 bytes */
        uint32_t value = 0, shift = 0;
        while (*msg != '\0') {
            value |= (uint32_t)*msg++ << shift;
            shift += 8;
            if (shift >= 32) {
                while (ITM->PORT[channel].u32 == 0UL)
                    {} /* null statement */
                ITM->PORT[channel].u32 = value;
                value = shift = 0;
            }
        }
        /* transmit last collected characters */
        if (shift > 0) {
            while (ITM->PORT[channel].u32 == 0UL)
                {}
            ITM->PORT[channel].u32 = value;
        }
    }
}

```

The PORT register allows 8-bit, 16-bit and 32-bit accesses, and this relates to the trailing-zero compression used by the SWO Trace protocol (again, see section [TRACESWO Protocol on page 8](#)). In fact, the implementation in the above snippet could be optimized a little further still: when transmitting the last collected bytes, it now always sends a 32-bit payload —due to the assignment to PORT[].u32.

For trace viewers, zero compression adds the complexity that on reception of a packet with a 1-byte or 2-byte payload, there is no automatic way to know whether it should possibly be expanded to a 32-bit value. Text messages do not contain zero bytes, so that is our escape here, but the above becomes relevant in chapter [The Common Trace Format \(page 72\)](#), which uses a binary stream.

SWO Tracing must first be configured in the microcontroller, which can be done either in the firmware (i.e. source code), or via the debug probe. Joseph Yiu, author of *The Definitive Guide to ARM Cortex-M3 Processors*, argues that configuration should be done by the debugging tool, as to avoid that the firmware and the debugging tool overwrite each-other's settings. On the other hand, some microcontrollers require additional device-specific configu-

ration that is not standardized by ARM. Configuring the tracing in the source code (at least partially) may therefore be unavoidable.

The Orbuculum project allows both approaches. The trace capture tools of this project do not perform any configuration, but the project comes with .gdbinit files with settings and definitions to perform the configuration from within GDB. The Orbuculum trace tools do not require GDB in itself, but even if you perform the trace configuration in code, you still need GDB to enable the trace option on the Black Magic Probe.

The command to enable tracing in the Black Magic Probe is below. Once set, it remains enabled (there is no way to disable the capture of SWO tracing in the Black Magic Probe, except for unplugging and re-plugging it).

```
(gdb) monitor traceswo
```

The SWO Trace protocol uses one of two serial formats: asynchronous encoding and Manchester encoding. The ARM documentation occasionally refers to these encodings as NRZ and RZ (Non-Return-to-Zero and Return-to-Zero). A property of Manchester encoding is that the clock speed can be determined from the data stream, so the bit rate does not need to be specified on the traceswo command. However, the Black Magic Probe lacks a hardware decoder for the Manchester bit stream, and therefore (since it handles the decoding in software) the supported bit rates are limited to roughly 200 kb/s.

The asynchronous protocol generally allows for higher bit rates. The clock speed cannot be recovered from the data stream though, so for asynchronous encoding, the bit rate must be set on the traceswo command.

```
(gdb) monitor traceswo 2250000
```

The target must be able to configure the same bit rate, within an error margin of 3%. Also note that the debug probe may have additional limits on the supported bit rates. For example, on the Black Magic Probe (and clones that use the STM32F10x microcontroller), the bit rate must be 4.5 Mb/s divided by an integer value, and with a maximum of 2.25 Mb/s.²

The choice between the two protocols may be dictated by the debug probe. The native Black Magic Probe supports only Manchester encoding,³ whereas the ctxLink probe (and a few other derivatives of the Black Magic Probe) instead support asynchronous encoding exclusively.

² Running at 72 MHz, the USART of the STM32F10x is limited to 4.5 Mb/s. However, the USB peripheral of the STM32F10x overflows at a continuous data stream of 4.5 Mb/s, which is why the “traceswo” bit rate is limited to half that rate.

³ Hardware version 2.3 of the Black Magic Probe is adapted to support asynchronous encoding as an option, but firmware support is “pending” at the time of writing.

The initialization that is generic for all ARM Cortex microcontrollers starts below. It involves a number of subcomponents of the CoreSight architecture, notably the *Instrumentation Trace Macrocell* (ITM) and the *Trace Port Interface Unit* (TPIU, also called TPI), but registers in the *Core Debug and Data Watchpoint & Trace* (DWT) modules may come into play as well.

```
void trace_init(int protocol, uint32_t bitrate, uint32_t channelmask)
{
    uint32_t clockfreq = (protocol == 1) ? 2 * bitrate : bitrate;
    CoreDebug->DEMCR = CoreDebug_DEMCR_TRCENA_Msk;

    TPI->CSPSR = 1;           /* protocol width = 1 bit */
    TPI->SPPR = protocol;    /* 1 = Manchester, 2 = Asynchronous */
    TPI->ACPR = CPU_CLOCK_FREQ / clockfreq - 1;
    TPI->FFCR = 0;           /* turn off formatter, discard ETM output */

    ITM->LAR = 0xC5ACCE55;   /* unlock access to ITM registers */
    ITM->TCR = ITM_TCR_SWOENA_Msk | ITM_TCR_ITMENA_Msk;
    ITM->TPR = 0;             /* privileged access is off */
    ITM->TER = channelmask;  /* enable stimulus channel(s) */
}
```

Parameter “protocol” must be 1 for Manchester encoding, or 2 for asynchronous encoding. Parameter “channelmask” is a bit mask where a “1” bit enables the respective channel. Note that for Manchester encoding, the clock frequency is twice the bit rate, because there may be transitions halfway the bit period.

An extra device-specific initialization step often needs to precede the generic initialization. Examples for a few microcontroller series are below. Note that some microcontrollers do not need any device-specific initialization (for example, the LPC175x and LPC176x series).

STM32F10x series

```
void trace_init_STM32F10x(void)
{
    RCC->APB2ENR |= RCC_APB2ENR_AFIOEN; /* enable AFIO access */
    AFIO->MAPR |= AFIO_MAPR_SWJ_CFG_1; /* disable JTAG to release TRACESWO*/
    DBGMCU->CR |= DBGMCU_CR_TRACE_IOEN; /* enable IO trace pins */
}
```

If AFIO_MAPR_SWJ_CFG_1 is not defined in the device header file of your development suite, note that it is “ $(2 \ll 24)$.”

STM32F4xx series⁴

```
void trace_init_STM32F4xx(void)
{
    RCC->AHB1ENR |= RCC_AHB1ENR_GPIOBEN; /* enable GPIOB clock */
    GPIOB->MODER = (GPIOB->MODER & ~0x000000c0) | 0x00000080; /* alt func for PB3 */
    GPIOB->AFR[0] &= ~0x0000f000;           /* set AF0 (==TRACESWO) on PB3 */
    GPIOB->OSPEEDR |= 0x000000c0;          /* set max speed on PB3 */
    GPIOB->PUPDR &= ~0x000000c0;           /* no pull-up or pull-down on PB3 */
    DBGMCU->CR |= DBGMCU_CR_TRACE_IOEN; /* enable IO trace pins */
}
```

SAM D5x series⁴

```
void trace_init_SAMD5x(void)
{
    GCLK->PCHCTRL[47] = GCLK_PCHCTRL_GEN(0) | GCLK_PCHCTRL_CHEN; /* enable peripheral
                                                                     clock on GCLK_CM4_TRACE */
    PORT->Group[1].PMUX[15].bit.PMUXE = PORT_PMUX_PMUXE(7);      /* set PB30 to SWO */
    PORT->Group[1].PINCFG[30].bit.PMUXEN = 1;                      /* enable PMUX for PB30 */
}
```

LPC13xx series

```
void trace_init_LPC13xx(void)
{
    LPC_SYSCTL->TRACECLKDIV = 1;
    LPC_IOCON->PIO0_9 = 0x83;      /* func 3, no pull-up/down */
}
```

LPC15xx series

```
void trace_init_LPC15xx(int pin)
{
    LPC_SYSCTL->TRACECLKDIV = 1;
    LPC_SWM->PINASSIGN15 = (LPC_SWM->PINASSIGN15 & ~(0xff << 8)) | (pin << 8);
}
```

LPC5410x series

```
void trace_init_LPC15xx(void)
{
    LPC_SYSCTL->TRACECLKDIV = 1;
    LPC_SYSCTL->SYSAHBCLKCTRLSET = 1 << 13;
    LPC_IOCON->PIO0_15 = 0x82;    /* func 2, no pull-up/down, digital */
}
```

⁴ Adapted from the GDB scripts of the Orbuculum project.

LPC5411x series

```
void trace_init_LPC15xx(void)
{
    LPC_SYSCTL->TRACECLKDIV = 0;
    LPC_SYSCTL->SYSAHBCLKCTRLSET = 1 << 13;
    LPC_IOCON->PIO0_15 = 0x82; /* func 2, no pull-up/down, digital */
}
```

LPC546xx series

```
void trace_init_LPC15xx(void)
{
    LPC_SYSCTL->TRACECLKDIV = 0;
    LPC_SYSCTL->SYSAHBCLKCTRLSET = 1 << 13;
    LPC_IOCON->PIO0_10 = 0x306; /* func 6, digital, filter off */
}
```

SWO Tracing on the Cortex M0/M0+

The ARM Cortex M0 and M0+ architectures lack support for SWO tracing. While you still have the option for tracing via a UART or semihosting, if you want to use a uniform debugging environment for all ARM Cortex microcontrollers, it may be worthwhile to emulate SWO tracing on Cortex M0/M0+.

- *Emulating asynchronous mode*

When using a ctxLink or another debug probe that supports asynchronous mode, the first step in emulating SWO is to wire the TxD pin of the UART to TRACESWO on the debug connector. The function to transmit the trace messages must be adapted to add a header byte in front of each packet (as explained in section [TRACESWO Protocol on page 8](#)). In a nutshell, an SWO packet can have a payload 1, 2, or 4 bytes, so there is a header byte for every sequence of payload. Obviously, it must also store the data (header bytes plus payload) in the UART FIFO instead of in the ITM FIFO.

The function `ARM_USART_Send` that is used in the snippet below is appropriate for the Keil implementation of CMSIS (and perhaps others); you may need to replace it with an equivalent function when using another UART driver library. In this implementation, the global variable `TRACESWO_TER` takes over the role of the “Trace Enable Register” (TER) of the ITM. It must be declared as a 32-bit integer, and I recommend that it is initialized to zero. This way, when running the firmware outside a debugger, all traces drop out immediately, but when running under GDB (or a trace viewer that uses the `gdbserver`), the debugger can set this variable to a non-zero value and enable the trace channels. The **BMTTrace** trace viewer ([page 69](#)) and **BMDDebug** front-end ([page 47](#)) check for a variable with the name “`TRACESWO_TER`” and configure it automatically when enabling or disabling channels from the user interface.

```

void trace(int channel, const unsigned char *data, unsigned size)
{
    if (TRACESWO_TER & (1 << channel)) { /* if channel is enabled */
        uint8_t header;
        while (size >= 4) {
            header = (channel << 3) 3;
            ARM_USART_Send(&header, 1);
            ARM_USART_Send(data, 4);
            data += 4;
            size -= 4;
        }
        if (size >= 2) {
            header = (channel << 3) 2;
            ARM_USART_Send(&header, 1);
            ARM_USART_Send(data, 2);
            data += 2;
            size -= 2;
        }
        if (size >= 1) {
            header = (channel << 3) 1;
            ARM_USART_Send(&header, 1);
            ARM_USART_Send(data, 1);
        }
    }
}

```

- *Emulating Manchester mode*

At the moment of writing, the native Black Magic Probe only supports Manchester mode for SWO tracing. The obvious recourse is to emulate Manchester mode via bit-banging, but that is slow, and to keep within the timing constraints the bit-banging routine must run with interrupts disabled. The combination of the two: slow code that runs with interrupts disabled, carries a risk that interrupts are not responded to quickly enough, or even that they are missed altogether.

There is yet a way to implement hardware-supported SWO emulation on a Cortex M0/M0+, if you have a spare SPI interface on your microcontroller. The trick is to expand each bit that is transmitted to a two-bit sequence: a 1 to “10” and a 0 to “01”, and then transmit these through over the MOSI line. This expansion can be efficiently done per 4 bits with a 16-byte lookup table. This same lookup table inverts the bit order: the SPI protocol transmits the most-significant bit first, whereas the SWO protocol (with Manchester encoding) transmits the least-significant bit first.

```

static const uint8_t manchester_lookup[16] = {
    0x55, /* 0000 -> 0101 0101 */
    0x95, /* 0001 -> 1001 0101 */
    0x65, /* 0010 -> 0110 0101 */
    0xa5, /* 0011 -> 1010 0101 */
    0x59, /* 0100 -> 0101 1001 */
    0x99, /* 0101 -> 1001 1001 */
    0x69, /* 0110 -> 0110 1001 */
    0xa9, /* 0111 -> 1010 1001 */
    0x56, /* 1000 -> 0101 0110 */
    0x96, /* 1001 -> 1001 0110 */
    0x66, /* 1010 -> 0110 0110 */
    0xa6, /* 1011 -> 1010 0110 */
    0x5a, /* 1100 -> 0101 1010 */
    0x9a, /* 1101 -> 1001 1010 */
    0x6a, /* 1110 -> 0110 1010 */
    0xaa, /* 1111 -> 1010 1010 */
};

#define M_EXPAND(buffer, byte) \
    ( (buffer)[0] = manchester_lookup[(byte) & 0x0f],           \
      (buffer)[1] = manchester_lookup[(uint8_t)(byte) >> 4] )

```

Apart from the bit expansion, the routine to emulate Manchester encoding is similar to the one that emulates asynchronous encoding for SWO tracing, on [page 66](#), so it is not repeated here. A separate implementation (source code file) is provided among the example files with this book.

Monitoring Trace Data

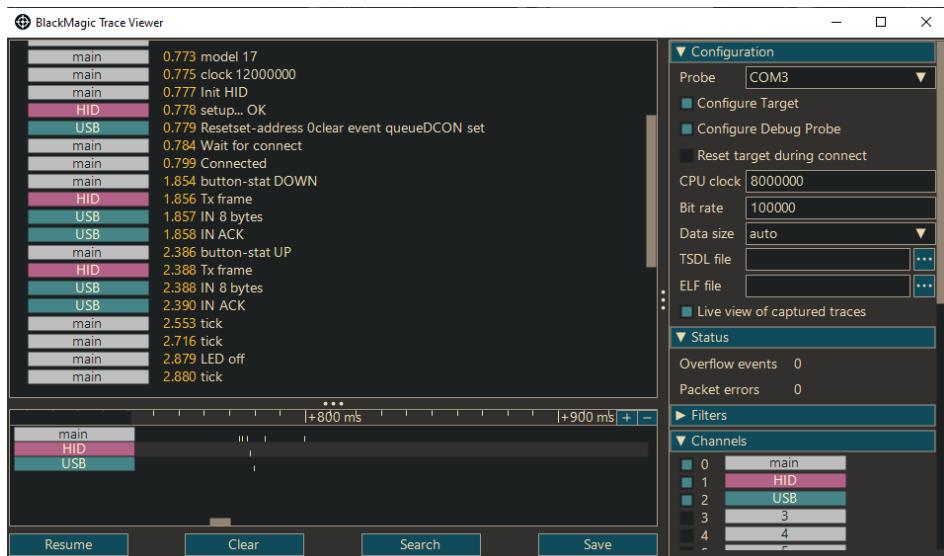
As of version 1.7 of the firmware of the Black Magic Probe, you can redirect the SWO trace data to the virtual UART. The upside is that you only need a serial terminal to view the trace data, and there are many to choose from. However, there are downsides too: channel information is not preserved, and filtering of enabled channels happens in the Black Magic Probe, instead of in the target. Furthermore, and probably a minor point, the Black Magic Probe has only a single UART interface, so you cannot use both the UART and trace redirection at the same time. See the `traceswo` command on [page 44](#) for more information.

To elaborate on the limitations of the transmission of SWO trace data over the UART, there are two practical uses for the channel information. The first is to separate the messages in different lists, or mark them in different colours. The second is to reduce the performance impact of tracing, by disabling the channels that are not relevant in the context of a particular debug session. Trace data that is *not* transmitted implicitly has minimal overhead. When

you let the Black Magic Probe transmit the SWO trace data over the UART, it allows you to enable or disable channels; however, it is the Black Magic Probe that filters out the disabled channels. The target still transmits all trace data for *all* channels to the Black Magic Probe. So you are still subject to the full performance penalty of the SWO transmission —and especially so in the case of Manchester encoding.

The alternative is to capture the SWO trace data directly. This requires a special tool or viewer. An advanced set of tools is the Orbuculum project, which was mentioned earlier. The main program, `orbuculum`, does the hardware capture and provides the data (after some internal processing) onto a TCP/IP port. Other utilities in the project connect to this TCP/IP port for post-processing and visualization. This client-server architecture allows multiple tools or viewers to access the trace data simultaneously. The packet data that the `orbuculum` server makes available on the TCP/IP port has the same format as that of the Segger J-Link probe, thereby allowing you to use the Segger software tools with the Black Magic Probe. Orbuculum runs on Linux, MacOS and a Microsoft Windows.

A stand-alone graphical trace viewer for SWO tracing using the Black Magic Probe is BMTrace: the BlackMagic Trace Viewer. It runs under Microsoft Windows and Linux. The BMTrace utility does not require GDB, because it uses the *Remote Serial Protocol* (RSP) to configure the target and the Black Magic Probe. The BMTrace utility performs the generic configuration for SWO tracing as well as the device-specific configuration for the microcontrollers that it supports. Another distinctive feature of BMTrace is that it supports the [Common Trace Format](#), see [page 72](#).



As described earlier, SWO tracing can use either modes Manchester or Asynchronous, but most variants of the Black Magic Probe support only one of these (and not both). If BMTrace detects that the selected debug probe is a native Black Magic Probe, it sets Manchester mode; likewise, if it detects the ctxLink probe, it sets Asynchronous mode. For other Black Magic Probe variants, you must select the mode in the configuration of BMTrace.

The BMTrace utility optionally configures the target for SWO tracing, and it sets up the Black Magic Probe for SWO tracing as well. For the target configuration, it needs to know the clock that the target microcontroller runs on, as well as the data rate (bit rate) of the transfer.

You can select to skip the target configuration. The target configuration for SWO (both generic and device-specific) then has to be done from GDB, or be performed in the firmware code —like in the code snippets starting on [page 64](#). Setting up the Black Magic Probe for SWO tracing can also be disabled. If both these options are disabled, BMTrace functions as a “passive listener”: it captures SWO trace messages, but does not interact with the target and does not connect to Black Magic Probe’s gdbserver (it connects only to the separate USB endpoint for SWO tracing). The “passive listener” mode allows you to use BMTrace in combination with GDB (which then connects to gdbserver).

Any of the 32 channels can be enabled or disabled. A right-click on the channel selector pops up a window to set a colour and a name for the channel. Note that when running in passive mode, any disabled channels are simply hidden in the trace viewer; they are not disabled in the target (because BMTrace does not communicate with the Black Magic Probe in passive mode). When running in CTF mode ([Common Trace Format](#), see [page 72](#)), the names of the channels are overruled by the “stream” names that are defined in the metadata file for the traces.

Apart from filtering on channels, BMTrace allows filtering incoming messages on keywords in the text. If no filters are set, all messages are shown; if one or more filters are set, only the messages that match any of these filters are shown. If the filter text starts with a “~”, the filter is inverted: the message is *not* shown if it contains the keyword (behind the tilde). Each filter can be enabled or disabled, for quickly toggling them on or off.

The time stamps in the BMTrace utility are relative to the first message that was received. With one exception, these time stamps are of the moment of reception of the trace data. Due to latencies of the USB stack and jitter in the scheduling of the operating system, these time stamps are indicative, but not conclusive. The exception is that *if* the incoming trace data is in the [Common Trace Format](#) and timestamps are present in the CTF stream, BMTrace shows these embedded timestamps instead. These timestamps are generated on the target, and they are generally more accurate.

Tracing with Command List on Breakpoints

Setting and using breakpoints is covered in section [Breakpoints and watchpoints \(page 38\)](#). A feature of GDB is that a list of commands may be attached to a breakpoint, and this list is executed whenever the breakpoint is hit. The trick is: when the final command in this list is “continue”, you have created a breakpoint that “drops through”.

For example, consider a command list with only the `continue` command:

```
(gdb) break 121
Breakpoint 3 at 0x3ce: file blinky.c, line 121.
(gdb) command 3
Type commands for breakpoint(s) 3, one per line.
End with a line saying just "end".
>continue
>end
```

On setting a breakpoint (on hypothetical line 121), GDB responds that breakpoint number 3 was set. When adding a command list, we will therefore have to repeat this identifier.

When running the code, GDB will print lines similar to the following, each time that the breakpoint is hit:

```
Breakpoint 3, main () at blinky.c:121
121      LPC_GPIO->SET[led_iobit] = led_iobit; /* turn LED on */
```

While this only shows that the line was reached, the important difference with the alternative trace methods is that the code does not need to be instrumented with `trace` calls. This implies that no recompilation is necessary if you want to move or add a trace-point. This method of tracing is therefore convenient if you want to check whether a particular line is reached. A limitation of this technique is that there is only a small pool of hardware breakpoints (which are needed when running from Flash).

Any GDB command can be inserted before the `continue` command. For example, a `print` command to show the values of specific variables, or a `backtrace` command to show the call stack that lead to the breakpoint being reached.

The Common Trace Format

As explained in the chapter on [Run-Time Tracing \(page 56\)](#), the intention of run-time tracing is to be a non-intrusive method of debugging. This implies that the trace messages should have negligible overhead, in time and other resources. If the overhead is non-negligible, the software may behave differently when being traced, than when running without tracing: a symptom that is called the *probe effect*.¹

When we focus on the time, the factors that contribute to “overhead” (i.e. latency) are:

- ◊ The need to format the data into a trace message prior to transmitting it.
- ◊ The amount of data to transfer, either to a remote “trace viewer” or internally to a display system.
- ◊ The speed of the data transfer, and any I/O overhead in accessing it.

When it comes to avoiding the probe effect, there is a prevalent fixation on the last point, the speed of the transfer interface. Yet, it is obvious that no matter how well you’ve optimized `sprintf`, skipping it entirely will always be quicker; like it is obvious that transmitting a few bytes is quicker than transmitting many (under equal conditions). Possibly, higher transfer speeds were the easiest goal to achieve in the early days, and perhaps as a corollary to the *Law of the Hammer*,² the reflex is to search for a bigger hammer if the current one won’t do any more.

Both the other two points (avoiding message formatting on the *target* and minimizing the amount of data that needs to be transferred) are addressed by the Common Trace Format (CTF). The Common Trace Format is a specification for a binary data format plus a human-readable “metadata file” to map the binary data to readable text. It thus does away with the formatting and conversion on the microcontroller, and it also skips transferring text strings when it can instead reference these strings in the metadata file. The CTF specification is maintained by the Diagnostic and Monitoring workgroup (DiaMon) of the Linux Foundation.

The metadata file defines the names of trace “events,” the streams that these events belong to, and the names and types of any parameters of each event. This is all recorded in a declarative language with a C-like syntax: the Trace Stream Description Language (TSDL). The metadata is shared (directly or indirectly) between the target that produces the trace messages and the trace

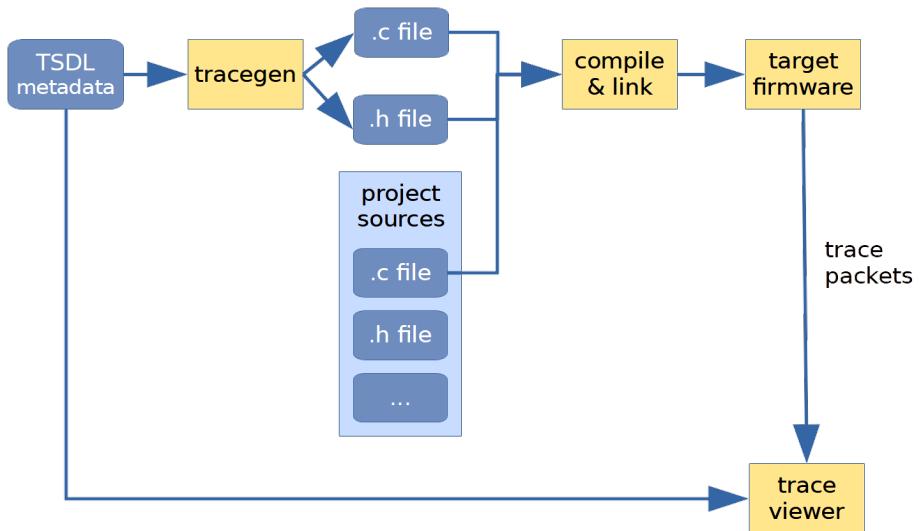
¹ J. Gait; *A probe effect in concurrent programs*; Software: Practice and Experience; March 1986.

² “I suppose it is tempting, if the only tool you have is a hammer, to treat everything as if it were a nail.” [Abraham Maslow; *The Psychology of Science*; 1966]

viewer. The Common Trace Format achieves its compactness because the data in this metadata file is never transmitted.

The Common Trace Format is the cornerstone of LTTng (Linux Trace Toolkit next generation); however, a call into LTTng is not exactly low-overhead in execution time —the rationale for LTTng’s use of CTF is to minimize storage requirements. Besides, it is not an option for embedded systems that run on something other than the full Linux kernel.

Two tools exist that generate OS-independent C code for CTF: `barectf` by the same authors as CTF, and `tracegen` (which is a companion tool to this book). Both tools use the metadata to generate individual C functions to build a binary CTF “packet” for each particular trace event. The generated file is then included in the build for the target firmware, and the source code can call the generated functions to transmit a trace packet in the compact CTF format. The `barectf` tool replaced TSDL with YAML as the metadata language (and it generates a TSDL file for the trace viewer), while the `tracegen` tool sticks with TSDL, but adds some extensions to make it more convenient.



In the above flow chart, the reference to “tracegen” could also be `barectf`. In fact, when using `barectf`, the flow is slightly different: the input to `barectf` is a YAML file, and it creates a TSDL file (along C source and header files) for the trace viewer. The “trace viewer” in the diagram may be either [BMTrace](#) (the BlackMagic Trace Viewer, see [page 69](#)) or another CTF compatible viewer, like Trace Compass.

Binary Packet Format

The Common Trace Format sends trace messages in packets. A packet holds one or more events. An event is basically a single trace message. In practice, packing multiple events in a packet is only useful if the transport protocol imposes a fixed or minimum size on packets. For stream-based protocols like RS232 or SWO (which this book focusses on), a packet holds a single event.



The packet header is optional; it contains a magic value to flag the binary data as the start of a CTF packet and the stream identifier. More information about the packet, such as its size and encoding, may follow in the (equally optional) packet context block.

For each event, an event header is required, because it contains the event identifier (plus possibly a timestamp for the event). The “event fields” block, at the tail of the event, holds any additional parameters that the event has. For example, if you trace a temperature sensor, the event name could be “temperature” and the single field the value in degrees Celsius or Fahrenheit (or Kelvin, for that matter). The “stream context” and “event context” blocks, are usually not relevant for embedded systems. The stream context holds data that applies to all events in the stream, whereas the event context has data that is specific to the event.

Which of the optional headers you should include in the packet depends in part on the transfer protocol. If it is packet-based, like USB or Ethernet, you may choose to omit the packet header, but instead include a packet context with the size of that packet. If, on the other hand, it is a byte stream, like RS232 or SWO, the packet header is as good as mandatory, while the package context is of little use.

A Synopsis of TSDL

The *Trace Stream Description Language* uses a syntax inspired by the C typing system. It will therefore be familiar to most embedded systems' developers. The full specification of this declaration language is on the DiaMon site, see the chapter [Further Information](#) on page 119 for a link.

A minimal example for a specification file is below. It defines a single event called “peltier-plate,” with a field called “voltage” of type “unsigned char.”

```
event {
    name = "peltier-plate";
    fields := struct {
        unsigned char voltage;
    };
};
```

Neither a packet header nor an event header are defined; therefore, these will not be present in the byte stream. Since the size of the single field is a byte, when the byte stream is:

18 1A 1B

it will be translated by the trace viewer to the following three events:

```
peltier-plate: voltage = 24
peltier-plate: voltage = 26
peltier-plate: voltage = 27
```

Merely a single byte needs to be transmitted for a descriptive parametrized event, it does not get much more compact than that. However, this is an exceptional case. When there is more than one event, an event header is needed so that the various events can be distinguished. This leads to the need for a packet header as well: to determine the function of each byte in a byte stream, one must know its position in the packet definition, and therefore one must know where the packet starts in the byte stream.

The following snippet addresses those issues. It defines a packet header in the `trace` section and an event header in the `stream` section.

```
trace {
    major = 1;
    minor = 8;
    packet.header := struct {
        uint16_t magic;
    };
};

stream {
    event.header := struct {
        uint16_t event.id;
    };
};

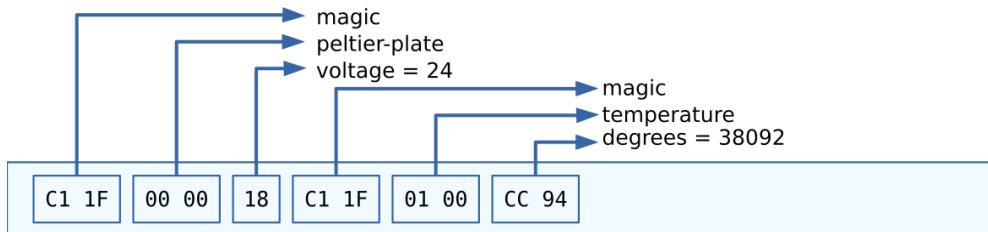
event {
    id = 0;
    name = "peltier-plate";
    fields := struct {
        unsigned char voltage;
    };
};
```

```

event {
    id = 1;
    name = "temperature";
    fields := struct {
        uint16_t degrees;
    };
};

```

An example of a byte stream that matches the above trace description is:



The trace viewer would display the two trace messages:

```

peltier-plate: voltage = 24
temperature: degrees = 38092

```

In tracegen, types like `uint16_t` (as used in the above example) are predefined. When using Babeltrace or another system, you may need to define these types yourself. This can be done with `typedef`, in the same way as in C, or with the more comprehensive `typealias`. The `typealias` construct allows you to set the size of the variable unambiguously, as well as any scaling ("fixed-point" representation), and in which base the number must be displayed (decimal, hexadecimal, binary). The snippet below shows the changes to the "temperature" event.

```

typealias integer {
    size = 16;
    scale = 1024;
    signed = false;
} := fixed_point;

event {
    id = 1;
    name = "temperature";
    fields := struct {
        fixed_point degrees;
    };
};

```

When the event for the temperature sensor is changed to a scaled integer (6 bits integer part, 10 bits fractional part, or a scaling factor of 2^{10}), the byte stream C1 1F 01 00 CC 94 would be displayed as:

```
temperature: degrees = 37.199
```

TSDL knows the standard C types `char`, `int`, `short`, `long`, `float` & `double`, along with their unsigned variants. It also inherits the annoyance of C that the size and byte-order (“endianness”) of these types are implementation-dependent. The tracegen utility uses the following defaults:

<code>char</code>	8-bits, signed
<code>int</code>	32-bits, little-endian
<code>short int</code>	16-bits, little-endian
<code>long int</code>	32-bits, little-endian
<code>float</code>	single precision IEEE-754, 32-bits
<code>double</code>	double precision IEEE-754, 64-bits
<code>enum</code>	“int”-size (so 32-bits by default).

The tracegen utility also predefines the most common “fixed width” integer types typically found in `stdint.h`, like `int16_t` and `uint32_t`. All predefined types can be overruled (with a `typealias`), if so desired.

For zero-terminated text strings, you use the special type “`string`” in TSDL. The tracegen utility converts it to “`const char*`” in the trace support files that it generates, see page 80. The default encoding for strings is UTF-8; though you can set the encoding to ASCII, this is rarely useful —UTF-8 is fully compatible with plain ASCII, and for extended ASCII you would need to know which codepage to use for the upper 128 characters (you cannot set the codepage in TSDL).

For general-purpose trace message output, you could therefore create a definition like the one below:

```
event debug_message {
    fields := struct {
        string msg;
    };
};
```

This creates a function that just transmits a message in a zero-terminated string, and that message can be about anything. If you are used to “`printf`”-style tracing (or methods derived from that), it might be tempting to implement that as the only event and call it from everywhere with pre-formatted strings. However, this does away with the principal advantage of CTF: efficiency, due to a compact encoding. In other words, if you are using CTF in this way, it might not be worth the trouble to use CTF at all.

General structure

At the global level are the declarations of the trace context, the streams, the events, plus any type definitions. The trace context includes the packet header; each stream declaration includes an event header; and each event may contain a declaration of one or more fields. In the generated C code, each event declaration is translated to a function, and the fields are translated to parameters.

Every declaration starts with a keyword, like `trace`, `stream` or `event`. The definition of the attributes of the declaration then follow between curly braces: the declaration *block*. For the `stream` and `event` keywords, a name may optionally be present; this is the name of the stream or the event (this is a tracegen extension).

Trace context

The trace context is a block at the global level, that is introduced with the `trace` keyword. It may contain the following fields inside its declaration block:

<code>major</code>	Major version of the CTF specification.
<code>minor</code>	Minor version of the CTF specification.
<code>version</code>	The version in “major.minor” format; the current version is 1.8. This is a tracegen extension, and it can be used as an alternative to the <code>major</code> and <code>minor</code> fields.
<code>uuid</code>	The UUID (in string format) may be used to ensure that a received packet matches the metadata file (the UUID should then also be included in the packet header).
<code>byte-order</code>	Either <code>le</code> or <code>be</code> for Little-Endian and Big-Endian respectively. Little-Endian is the default.
<code>packet.header</code>	The declaration of the packet header, see the next section.

Packet header

The packet header is nested inside the trace context (the `packet.header` structure). It contains the *type definitions* for one or more of the following fields (in any order):

<code>magic</code>	Must be declared as a 1-byte, 2-byte, or 4-byte unsigned integer. The purpose of this field is to mark the start of a packet in a stream of bytes. A longer magic value gives a more reliable detection of the start of a packet, at the cost of more bytes being transmitted. A 2-byte integer is a common compromise.
--------------------	--

<code>uuid</code>	A user-supplied identifier, used to make sure that the byte stream of the traces matches the definitions in the metadata (the TSDL file). Due to its heavy cost in overhead (16 bytes added to every packet), its use is not recommended for embedded systems.
<code>stream.id</code>	A 1-, 2-, or 4-byte integer with the stream number. Redundant if the trace information uses only a single stream; also redundant for SWO tracing when less than 32 streams are used (because the stream ID is mapped to the SWO channel). This field may also be called “ <code>stream_id</code> ” for compatibility with other CTF implementations.

Stream Declaration

A trace specification may contain one or more streams. Since we need to include a unique ID for each event, and this ID is declared in the *event header*, which in turn is declared in a stream... there is typically at least one stream declaration.

The stream declaration is a block at the global level, that is introduced with the `stream` keyword. It may contain the following fields:

<code>id</code>	The unique numeric identifier for the stream. Both tracegen and barectf can auto-number streams, so this field is optional.
<code>name</code>	A unique name for the stream. This is a tracegen extension.
<code>event.header</code>	The declaration of the event header, see below.

If there is only a single stream, it needs neither an ID nor a name. However, in practice, the `event.header` field should always be present.

When using tracegen, the name of the stream may also appear between the keyword `stream` and the opening brace “{.” For example, the definition:

```
stream PIDcontroller {
    /* other fields */
};
```

is equivalent to the regular form:

```
stream {
    name = PIDcontroller;
    /* other fields */
};
```

Event header

The event header is nested inside the `stream` declaration (the `event.header` structure). It contains the *type definitions* for one or more of the following fields:

<code>event.id</code>	A 1-, 2-, or 4-byte integer with the event ID. This field may also be called “ <code>id</code> ” for compatibility with other CTF implementations.
<code>timestamp</code>	A 4-byte or 8-byte timestamp for the event. The timestamp is linked to the definition of a clock in the TSDL file (see Timestamps further down on this page).

Event Declaration

Event declarations are converted to C functions that send trace packets by the utilities `tracegen` and `barectf`. An event declaration includes the event name and parameters, plus optionally the stream it is part of and attributes.

<code>id</code>	The unique numeric identifier for the event. Both <code>tracegen</code> and <code>barectf</code> can auto-number streams, so this field is optional.
<code>name</code>	A unique name for the event.
<code>stream.id</code>	The numeric ID of the stream that the event belongs to. Redundant if the trace information uses only a single stream. When using <code>tracegen</code> , you may also specify the stream <code>name</code> rather than the numeric ID. This field may also be called “ <code>stream.id</code> ” for compatibility with other CTF implementations.
<code>fields</code>	The declaration of a structure for the parameter names and types (if any).
<code>attribute</code>	Optional GCC attributes that will be included on the C function that <code>tracegen</code> generates.

The event name and (if relevant) the stream name, may be put between the `event` keyword and the opening brace “{” of the declaration. If both a stream name and an event name are present, the two should be separated with a double colon (“::”).

Timestamps

Timestamps must be linked to a clock. This takes two parts: the definition of a clock and the definition of a type that references this clock. The timestamp is then defined as that type.

```

clock {
    name = cycle_counter;
    freq = 1000000000;           /* frequency, in Hz */
};

typealias integer {
    size = 64;
    signed = false;
    map = clock.cycle_counter;
} := tickcount_t;

stream {
    event.header := struct {
        uint16_t event.id;
        tickcount_t timestamp;
    };
};

```

There are more (optional) fields in the `clock` specification, specifically for synchronizing various clocks in a heterogeneous tracing environment, but these are skipped here. The new type `tickcount_t` maps to this clock, and the `timestamp` field in the event header is defined as a `tickcount_t` type. Following the chain backward, the `timestamp` field is now linked to the clock “`cycle_counter`.”

Instead of having the target transmit the timestamps of every event, we recommend that a trace viewer displays the timestamp of when the trace packets are received (and that the timestamp is omitted from the event header). The timestamp of the reception is less accurate (due to latencies and jitter in the transmission protocol), but accuracy in the timestamps is usually only required for specific events: in those events, the timestamp can be transmitted as a parameter (an “event field”).

Scaling up: multiple streams, many events

When there are many trace events or multiple streams involved, a few short-hand notations exist to make maintenance of the metadata easier. When there are multiple streams, each stream should have a unique ID and each event (which should also have a unique ID) must indicate which stream it belongs to.

The `tracegen` utility extends TSDL by allowing a stream to have a name, so that an event can identify its stream by its name rather than a numeric constant. It also supports automatic numbering of streams and events (barectf also supports auto-numbering). For brevity in the TSDL file, the names of a stream and of an event can be placed immediately following the `stream` or `event` keywords, see the snippet below for examples. In the case of an event,

the name of the stream may be prefixed to the event name, with a double colon between the two names.

Below is the example from [page 75](#), extended with a stream name and using the shorthand notations.

```
trace {
    version = 1.8;
    packet.header := struct {
        uint16_t magic;
        uint8_t stream.id;      /* redundant with SWO */
    };
};

typealias integer {
    size = 16;
    scale = 1024;
    signed = false;
} := fixed_point;

stream cooler {
    event.header := struct {
        uint16_t event.id;
    };
};

event cooler::"peltier-plate" {
    fields := struct {
        unsigned char voltage;
    };
};

event cooler::temperature {
    fields := struct {
        fixed_point degrees;
    };
};
```

This snippet defines a stream “cooler” and the events “peltier-plate” and “temperature,” both linked to stream “cooler.” The name “peltier-plate” is between quotation marks, because it contains a “-” character. You may enclose all identifiers in quotation marks, but it is not needed if a name only contains letters, digits and “_” characters (like C identifiers).

Since there is only a single stream in this example, giving the stream a name and referencing its name explicitly in the events is actually redundant. The stream could equally well be anonymous, and the “cooler::” prefix could then be omitted from the event specifications.

When there is a single stream, the `stream.id` field in the `packet.header` structure is usually redundant. With SWO tracing, it is also redundant in the case

of multiple streams, because the stream ID is mapped to the SWO channel. The ID therefore does not have to be repeated in the packet header. Note that you are limited to 32 streams in this case.

Note that these shorthand notations are specific to the tracegen and [BM-Trace/BMDebug](#) utilities. When using a different trace viewer, the basic TSDL syntax (as specified on the site of the DiaMon workgroup) should be used.

Generating Trace Support Files

When running the tracegen utility on the metadata file, it generates a C source and a C header file. These files contain the definitions (prototypes) and the implementations of functions, and each of these functions creates and transmits a packet for an event.

For example, when the snippet on [page 82](#) is saved in a file with the name “peltier.tsdl,” you can run:

```
tracegen -s peltier.tsdl
```

The two files that are created, are named `trace_peltier.c` and `trace_peltier.h`. These contain the implementation and declaration of two functions (because there are two events defined in the TSDL snippet):

```
void trace_cooler_peltier_plate(unsigned char voltage);
void trace_cooler_temperature(fixed_point degrees);
```

The function names contain both the name of the stream and the names of the events. If the stream were anonymous, that part would not be present in the function names either. Any characters that are not valid for use in C identifiers are replaced by an underscore. This happened with the event name “peltier-plate” for example: in the C function name, the “-” is replaced by a “_.”

The “-s” option to tracegen makes it generate code for SWO tracing. When you would use the Common Trace Format for tracing over an RS232 line (or TTL-level UART), this option is not needed.

Also note how the types of the function arguments are copied from the metadata file into the C functions. Your source code should define a `fixed_point` type that matches the definition in the metadata. The alternative is to use the “-t” option on tracegen, in which case it will always attempt to translate the type in the metadata file to a basic C type.

```
tracegen -s -t peltier.tsdl
```

The above call would generate the following function prototype for the temperature event:

```
void trace_cooler_temperature(unsigned short degrees);
```

The function prototypes and implementations in the source and header files are wrapped in conditional compiled sections that test for the NTRACE macro. If the NTRACE macro is defined, the functions are disabled. Thus, if you need to build a release version of the firmware without any tracing functions, rebuild all code with a definition of NTRACE on the compiler command line.

The tracegen utility has a few more options. To see a summary, type:

```
tracegen -?
```

When integrating tracegen with Make, note that the output files have the base name of the input file, but with “trace_” prefixed to it. That is, if the input file is peltier.tsdl, the output are the files trace_peltier.c and trace_peltier.h. An inference rule to match this could look like:

```
trace_% .c : %.tsdl  
    tracegen -s -i:stdint.h $<
```

Integrating Tracing in your Source Code

The tracegen utility generates prototypes and implementations for transmitting trace events, as was shown in the previous section. When integrating this code in your project, one or two additional functions need to be provided by your code.

```
void trace_xmit(int stream_id, const unsigned char *data, unsigned size);  
unsigned long long trace_timestamp(void);
```

The task of the trace_xmit function is to truly transmit the data over a kind of port or interface. For SWO tracing, this would be an adaption of the trace function on [page 61](#), see below:

```
void trace_xmit(int stream_id, const unsigned char *data, unsigned size)  
{  
    if (((ITM->TCR & ITM_TCR_ITMENA) != 0UL && /* ITM tracing enabled */  
        (ITM->TER & (1 << stream_id)) != 0UL) /* ITM channel enabled */  
    {  
        /* collect and transmit characters in packets of 4 bytes */  
        uint32_t value = 0, shift = 0;  
        while (size-- > 0) {  
            value |= (uint32_t)*data++ << shift;  
            shift += 8;  
            if (shift >= 32) {  
                while (ITM->PORT[channel].u32 == 0UL)  
                    {} /* null statement */  
                ITM->PORT[channel].u32 = value;  
            }  
        }  
    }  
}
```

```

        value = shift = 0;
    }
}
/* transmit last collected characters */
if (shift > 0) {
    while (ITM->PORT[channel].u32 == 0UL)
        {}
    ITM->PORT[channel].u32 = value;
}
}
}

```

The above example assumes that you have run tracegen with the “-s” option on the TSDL file. Without the “-s” option, the definition of `trace_xmit` lacks the `stream_id` parameter (the stream ID would instead be present in the packet header).

The `trace_timestamp` function returns a timestamp, which is then transmitted as part of the event header. The return type of this function depends on the declaration of the clock in the TSDL file, see [page 80](#). If the event header does not include a timestamp, there is no need to implement this function (as it will not be called).

Mixing Common Trace Format with Plain Tracing

While the benefit of compactness of Common Trace Format is clear, it adds overhead in the programming effort. Instead of just calling `trace()` with a quick message as a parameter, the programmer now has to spell out the details of the trace message, including any parameters, in a separate TSDL file, and run another tool to create a C file that must be linked with your code. It is more work, and this extra work is worth it for the trace messages that you plan to keep in the code, for regression testing and quality control. For a quick throw-away test, however, this overhead stands in the way.

Fortunately, the two approaches can be mixed when using SWO tracing. A CTF trace message belongs to a stream, which maps to a channel (or *stimulus port*) of the ITM (*Instrumentation Trace Macrocell*), see [SWO Tracing on page 61](#). The trace viewer [BMTrace](#) (and the trace view in the [BMDebug front-end](#)) use the criterion that if a packet is received on a channel that is present in the TSDL file as a stream, that packet is decoded as CTF. Otherwise, the packet is assumed to contain plain text.

Hence, it suffices to reserve a channel for non-CTF (plain text) trace packets. A channel which you never use in TSDL files for stream IDs. Channel 30 is a pragmatic choice, because channel 31 is regularly reserved by an RTOS for tracing and profiling, and auto-numbering of stream IDs by tracegen or barectf starts at 0.

Applications for Run-Time Tracing

When it comes to where and how to use run-time tracing, the application that immediately springs to mind is to print out the program state, or the value of variables, at specific places in the code. This is the embedded equivalent of “printf-style” debugging. Run-time tracing has a wider scope than this, however.

Code Assertions

The *function* of an assertion is to display an error message and abort the program when its parameter evaluates to false. The *goal* of an assertion, however, is to always sit silent, because if it fails (and prints the error message), there is a bug in your code.

Without going into details (see the book *Writing Solid Code* by Steve Maguire¹ for that), note that assertions should therefore not replace error checking. You put assertions in your code to test things that you *know* must be true... provided that the code was called with the correct input parameters, but of which you *know* that these were checked by the caller. The answer to the question of why on earth you would test what you already know to be true, is that you may not know what you *think* you know.

In desktop software, the use of assertions is mainstream, because their use is straightforward, their presence declares pre-conditions, post-conditions and invariants in the code (as an informal expression of the formal specification), and it combines well with unit testing. In embedded development, assertions are less commonplace, and the reason (or at least one of the reasons) is that embedded systems lack a universal console (display) to print the “assertion failed” messages to.

Run-time tracing offers an alternative to the console. Of the methods described in chapter [Run-Time Tracing on page 56](#), semihosting has the advantages that it is always available when running under a debugger, and it requires no additional set-up in the debugger or debug probe. The relative low performance of semihosting is not an issue: an assertion only sends output when it fails —when there is a bug.

There are a few pitfalls in the use of assertions. The most important one is that an assertion should not have a side effect. Changing a variable inside the expression of an assertion is right out of the question, but C functions with side effects, like `strtok()` should be avoided inside an assertion as well. Apart

¹ Maguire, Steve; *Writing Solid Code*; Microsoft Press, 1993; ISBN 978-1556155512; or the second edition by Greyden Press, LLC, 2013; ISBN 978-1570740558.

from that, lengthy operations carry a risk as well, especially in time-sensitive or performance-critical code. Ideally, an assertion should take negligible time (and resources) for testing its condition.

Assertions grow the code size; especially the default implementation of the `assert` macro grows the code because it adds the expression and the filename that the assertion occurs in as strings to the code. For desktop programs, this is a minor issue, because desktop workstations and laptops have ample memory, but embedded systems are regularly quite constrained. In embedded software, it is common to re-implement the `assert` macro so that it is more economical with code space.

A first step is to eliminate the expression as a string. The filename and line number are sufficient to locate the expression that caused the “assertion failed” notification; duplicating the expression that failed, within that notification is redundant. Speaking of filenames, each time you add another `assert()` in a source file, the filename is stored as a string literal. You will want to merge these duplicate strings, so that only a single copy is stored and all `assert` macros reference that single copy. The GCC option to do this is `-fmerge-all-constants`.

The filenames can also be eliminated altogether, by printing the *address* of the error location, instead of the filename and line number. This approach reduces overhead to a minimum. Implementations of assertions typically have two parts, a conditionally defined macro and a function that is called when the assertion fails.

```
#define assert(condition) \
    if (condition) \
        {} \
    else \
        assert_fail()
```

This macro implements `assert` as a statement, as opposed to the standard C library that implements it as a conditional expression. The rationale is that this allows the GCC compiler to catch unintentional assignments in the condition; the standard implementation of `assert` stays silent when you write “`assert(var = 1)`,” even though an assignment inside an `assert` is always wrong. The “`if`” statement has both *then* and *else* parts (with the *then* part as an empty statement) in order to avoid a dangling-else problem.

The core functionality of the `assert` is implemented in the `assert_fail()` function. There is only a single implementation of this function, whilst there are potentially many invocations of the `assert` macro sprinkled throughout your code. Therefore, it saves code space to let the `assert_fail()` function determine the address of the assertion failure, rather than passing it as a parameter to `assert_fail()`.

```

__attribute__((always_inline)) static inline uint32_t __get_LR(void)
{
    register uint32_t result;
    __asm__ volatile ("mov %0, lr\n" : "=r" (result));
    return result;
}

static void addr_to_string(uint32_t addr, char* str)
{
    int i = sizeof(addr) * 2; /* always do 8 digits for a 32-bit value */
    str[i] = '\0';
    while (i > 0) {
        int digit = addr & 0x0f;
        str[--i] = (digit > 9) ? digit + ('a' - 10) : digit + '0';
        addr >>= 4;
    }
}

__attribute__((weak)) void assert_abort(void)
{
    __BKPT(0);
}

void assert_fail(void)
{
    register uint32_t addr = (__get_LR() & ~1) - 4;
    char buffer[] = "Assertion failed at *0x00000000\n";
    addr_to_string(addr, buffer + 23);
    trace(buffer);
    assert_abort();
}

```

The above snippet implements `assert_fail()`, plus three support functions. The first thing `assert_fail()` does is to get the value of the *Link Register*, which holds the address that `assert_fail()` returns to (or that it would return to). That address points behind the call to the function, which is why the size of one instruction is subtracted from it. The lowest bit is also cleared, because that bit is a flag for the ARM Cortex *Thumb mode*. This address is then converted to ASCII and sent out as a trace message.

The last action of `assert_fail()` is to call `assert_abort()`. The default implementation is a software breakpoint, but it can be overridden. Because of the “weak” linkage attribute on the default implementation of `assert_abort()`, it is overruled by a user-defined function with the same name. The intended purpose of `assert_abort()` is to reset all peripherals to a safe state. If the assertion is inside embedded code for a 3D printer, for example, `assert_abort()` would stop all fans and motors and shut heating off.

While on the subject, `__BKPT()` is a CMSIS macro. Other microcontroller support libraries will likely have a similar function for software breakpoints. Otherwise, a simple implementation for GCC is:

```
#define __BKPT(value)    __asm__ volatile ("bkpt "#value)
```

The `trace()` function in `assert_abort()` is a placeholder; it should be replaced by a function that does the actual output of the strings, by the method of your choosing.

The last step is to look up the filename and line number for an address, with help of symbolic information. When the code is loaded in GDB, this information can be obtained with the `info` command. Note that the asterisk is necessary.

```
(gdb) info line *0x08000505
```

On the command line, you can use the utility `addr2line` to get the filename and line number from an address (on a typical toolchain for the ARM Cortex, the full name may be `arm-none-eabi-addr2line`).

```
arm-none-eabi-addr2line -e blinky.elf 0x08000505  
d:\Tools\blinky\blinky.c:168
```

The [BMDebug front-end \(page 47\)](#) automatically looks up file and line information for messages that are printed through semihosting, when those messages contain an address as a hexadecimal number and with an asterisk in front. In particular, when the embedded host writes the following via semihosting:

```
Assertion failed at *0x08000505
```

The [BMDebug front-end](#) will display it in its semihosting view as:

```
Assertion failed at blinky, c:168
```

Tracing Function Entry and Exit

When your code runs in a debugger and halts at a breakpoint, quite often one of the things you want to find out is how you got there: the `backtrace` command is for that purpose. However, when a trace message pops up, the code doesn't halt, and you don't have the context of the message.

The solution is to trace all entries to all functions, as well as exits from them. The GCC compiler has a command-line option to instrument function entries and exits with a call to functions that you must implement:

```
-finstrument-functions
```

This option inserts a call to an “entry” function at each start of a function, and another call to an “exit” function just before the return. A template for these functions is:

```

__attribute__((no_instrument_function))
void __cyg_profile_func_enter(void *this_fn, void *call_site)
{
    /* ... */
}

__attribute__((no_instrument_function))
void __cyg_profile_func_exit(void *this_fn, void *call_site)
{
    /* ... */
}

```

The first parameter of both functions is the address of the function; the second the address of the function that made the call. Both these addresses can be looked up with the symbolic information, as was covered in the previous section on the `assert` macro.

The entry and exit functions themselves should *not* be instrumented, to avoid unbounded recursion. The same applies to *any* function called from the entry and exit functions —including *inline* functions. To avoid a function from being instrumented, you set the attribute “`no_instrument_function`” on it, as was done in the preceding snippet.

In addition to the attribute specifications in the source code, GCC also has command line options to block instrumentation for specific functions or for all functions in specific files, see `-finstrument-functions-exclude-file-list` and `-finstrument-functions-exclude-function-list` in the GCC documentation.

The implementation of the entry and exit functions will typically be a call to a function that outputs a trace message. In this particular case, low overhead is of the essence, and therefore it is particularly suited for the [Common Trace Format](#) (see page 72). If we ignore the `call_site` parameter (which is technically redundant, because you will have received an “entry” message for that caller too), an example implementation for the metadata for the entry and exit functions is in the snippet below. Note that this is not a complete TSDL file (e.g. it lacks the definitions of the packet and event headers), but just the part that declares the events —see the appendix [Code Profiling](#), specifically the section on [Calltree Analysis](#) (page 94) for a full TSDL file for function tracing.

```

typealias integer {
    size = 32;
    signed = false;
    base = symaddress;
} := code_address;

```

```

event enter {
    attribute = "no_instrument_function";
    fields := struct {
        code_address symbol;
    };
};

event exit {
    attribute = "no_instrument_function";
    fields := struct {
        code_address symbol;
    };
};

```

The main feature of the above code is the definition of the `code_address` type, and especially the declaration “`base = symaddress`” (`symaddress` may be abbreviated to `symaddr`). This declaration signals that any parameter with this type is a *symbol address*. This signals the trace viewer to look the address up in the symbolic information.

Currently, the **BMD**Debug**** and **BMT**Trace**** utilities print the function name instead of an address, when the “`base`” for the respective parameter is set to `symaddress`.

The functions generated by `tracegen` for these TSDL events can now be called from the `_cyg_profile_func_enter` and `_cyg_profile_func_exit` functions. The events in the above TSDL snippet have a declaration for an `attribute`, which is set to “`no_instrument_function`.” This attribute is copied to the generated C functions. As an alternative, you can add the option `-no-instr` on the `tracegen` command line, so that it adds the `no_instrument_function` attribute to all generated functions.

```
tracegen -s -t -no-instr blinky.tsdl
```

As covered in section [Integrating Tracing in your Source Code](#) (page 84), the C functions generated from the TSDL file call `trace_xmit`, a function that must be implemented by you. That section also gives an example implementation. When tracing function entry and exit, the `trace_xmit` function must also have the `no_instrument_function` attribute.

Code Profiling

Code profiling is a technique that you use to analyze where the program spends its time; it identifies the “hot spots” or “cycle-eaters” in the code. With this knowledge, you can decide where to optimize the code —and whether to optimize it at all.

Performance optimization is not just about making the code run faster. Increasing the efficiency of the code may allow you to reduce the clock frequency of the microcontroller. Since the power consumption of the microcontroller correlates with the clock frequency, you indirectly reduce power consumption. This may be especially relevant for battery-powered devices.

The two most common methods for profiling are by instrumentating the code and by sampling. Section [Tracing Function Entry and Exit \(page 89\)](#) briefly touched on instrumenting entry and exit of functions. This allows you to profile at a *function* level: it allows you to measure the time spent in each function. It does not tell you which loop or computation inside the function is the hot spot. There are tools to instrument source code such that you get timings on a *source line* level. The drawback is that due to the overhead of instrumentation, the code runs significantly slower, which can affect the sequence in which tasks and interrupts run —and this may then give rise to the probe effect.¹

Sampling-based profiling works by sampling the *Program Counter* (PC) at a regular interval, and then look up which line in the source code corresponds with the address. No instrumentation is needed, and the code runs at its original speed (no slowdown, and therefore no probe effect). On the other hand, the sampling frequency is typically in the range of a few kHz, whereas a microcontroller runs at several MHz. A lot of code can run between two samples. However, if a profiling session runs long enough, it results in a statistical distribution of where time is spent.

Sampling on ARM Cortex

The ARM CoreSight architecture implements PC sampling in its ITM (*Instrumentation Trace Macrocell*) and DWT (*Data Watchpoint & Trace*) units. This means that the sampling is separate from the microcontroller’s arithmetic and logic unit. Thus, sampling is truly non-intrusive; it does not even incur the overhead of interrupt processing.

The sampled values are transferred to the debug probe via the TRACESWO pin and interface. The configuration for profiling therefore has a lot in common with that for [SWO Tracing](#), see [page 61](#). In particular, the device-specific

¹ See also [page 72](#).

configuration for SWO tracing must be performed for profiling as well —see the various snippets starting on [page 64](#) for details.

Note that, as is the case with SWO tracing, the ARM Cortex M0 and M0+ architectures lack support for profiling.

The generic initialization for profiling is below. This snippet may be compared with the one on [page 64](#); there are only few differences in the set-up for profiling versus tracing.

```
void trace_init(int protocol, uint32_t bitrate, uint32_t samplerate)
{
    uint32_t clockfreq = (protocol == 1) ? 2 * bitrate : bitrate;
    uint32_t divider = CPU_CLOCK_FREQ / (1024 * samplerate);
    divider = min((divider > 0 ? divider - 1 : 0), 15); /* clamp in range 0..15 */

    CoreDebug->DEMCR = CoreDebug_DEMCR_TRCENA_Msk;

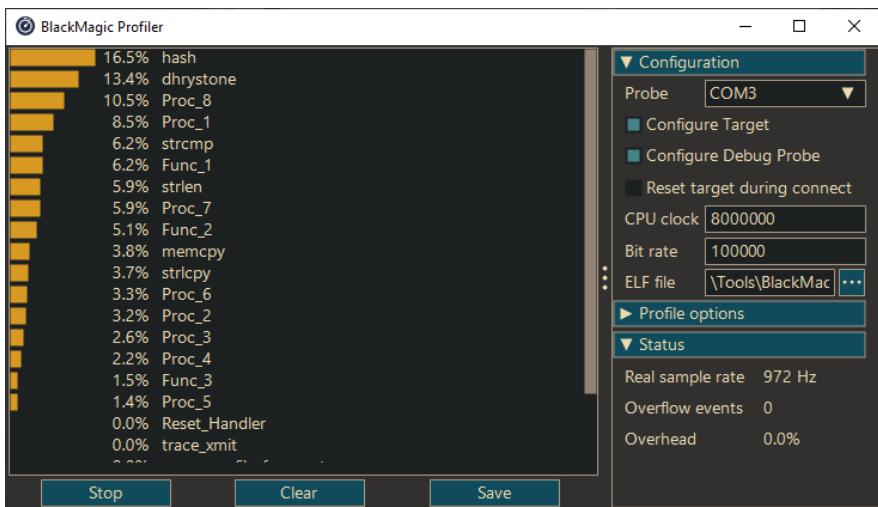
    TPI->CSPSR = 1;           /* protocol width = 1 bit */
    TPI->SPPR = protocol;    /* 1 = Manchester, 2 = Asynchronous */
    TPI->ACPR = CPU_CLOCK_FREQ / clockfreq - 1;
    TPI->FFCR = 0;           /* turn off formatter, discard ETM output */

    ITM->LAR = 0xC5ACCE55;   /* unlock access to ITM registers */
    ITM->TCR = (1 << 16) | ITM_TCR_DWTENA_Msk | ITM_TCR_ITMENA_Msk;
    ITM->TPR = 0;             /* privileged access is off */
    DWT->CTRL = DWT_CTRL_PCSAMPLENA_Msk | DWT_CTRL_CYCTAP_Msk |
                  DWT_CTRL_CYCCNTENA_Msk | ((divider & 0xf) << 1);
}
```

The above snippet limits the sampling frequency to the microcontroller clock divided by 1024. You can increase the maximum sampling frequency by a factor 16 when clearing the CYCTAP bit in the DWT_CTRL register. However, the maximum SWO bitrate that the Black Magic Probe supports, is the limiting factor.

Once the set-up is done, the microcontroller starts streaming the sampled PC addresses to the debug probe, over the TRACESWO pin. Each packet is five bytes long: a header byte (with a fixed value of 0x17) followed by a 32-bit address, transmitted low-byte first (see also section [TRACESWO Protocol](#) on [page 8](#)).

The BlackMagic Profiler (BMProfile) is a small utility that shows the profiling results in a bar-graph. On start-up, it initially shows a list of functions, sorted on sample counts. When you click on a particular function, it shows you the source code for that function, with the bar graph (and percentages of samples) for the source lines. Thus, the utility allows for both *function* level and *source line* level profiling. Clicking in the source view returns to the function list.



Like the [BMTTrace](#) utility, BMProfile can configure the target microcontroller for profiling through the debug interface. It does not use GDB, but relies on the *Remote Serial Protocol* (RSP) to communicate with the target microcontroller (similarly as to how GDB communicates with the target). For several families of microcontrollers, BMProfile can do the device-specific initialization for tracing & profiling as well. These steps are configurable: if the setting “Configure Target” is de-activated in the right column of the user-interface, no device-specific initialization will be done.

Sampled addresses that fall outside the range of the ELF file, or that cannot be attributed to a function, are collected as “overhead” in the bottom panel of the right column. Some microcontrollers have “drivers” or support routines in ROM that firmware code can use, and time spent in these would thus be collected as “overhead”. Similarly, functions linked in from the run-time library sometimes are neither included the DWARF information, nor in the ELF symbol table. When the PC is sampled inside these functions, it is also collected as overhead.

For further analysis, you can save the collected sample data in CSV format (comma-separated values).

Calltree Analysis

The screen capture of the BMProfile utility earlier on this page, shows that the function on top is hash with 16.6% of the MCU time. The total time spent in a routine, like function hash, is the time that a single run takes, multiplied by the number of times the routine runs. Thus, the question is, when a function comes out high in the profiling graph, is it because it runs slowly, or because it

is called so very often. And this is a question that a sampling profiler doesn't answer, on its own.

A first step in deeper understanding of the code and where it is spending its cycles, is to make a calltree of the program. If you are using C, you can use the GNU cflow program to create a calltree by static analysis. A "reverse calltree" generates for each function a list of call sequences that lead to running that function. This quickly tells you by who each function is called. The cflow program creates a reverse calltree with the --reverse command line option:

```
cflow --reverse dhystone.c hash.c
```

When using C++, the situation is less straightforward: cflow does not support C++ and commercial alternatives, such as Understand by SciTools and CppDepend by CoderGears, are fairly expensive. Another caveat is that the static call tree of a program is not necessarily the same as the run-time call tree. Functions passed as parameters, call-back functions and invocations of virtual functions are absent from the static calltree.

The alternative is to create a run-time calltree, by running the code while tracing the function entries and exits. This is the topic of section [Tracing Function Entry and Exit \(page 89\)](#). That section covered the concept and included a snippet of a TSDL file that might be used for function tracing. For completeness, a full (and slightly adapted) TSDL file for tracing functions is below.

```
trace {
    major = 1;
    minor = 8;
    packet.header := struct {
        uint16_t magic;
    };
};

stream function_profile {
    id = 31;
    event.header := struct {
        uint16_t id;
    };
};

typealias integer {
    size = 32;
    signed = false;
    base = symaddress;
} := code_address;

event function_profile::enter {
    attribute = "no_instrument_function";
```

```
    fields := struct {
        code_address symbol;
    };
};

event function_profile::exit {
    attribute = "no_instrument_function";
    fields := struct {
        code_address symbol;
    };
};
```

This TSDL file creates a stream “function_profile” at channel 31 —the channel reserved for profiling and system tracing. The goal now is to run the same code while capturing traces with [BMTrace \(page 69\)](#). As you will observe, the code will run significantly slower, but since the goal is to generate a calltree, accurate timings are not as relevant. What is relevant, though, is that all code paths that were executed during the profiling session, are also executed during the function tracing phase.

After saving the captured trace messages to a CSV file, you can run the `calltree` utility on it to generate a familiar calltree. The `calltree` utility is another software utility that comes with this book. Like the `cflow` utility, `calltree` can create both normal calltrees and reverse trees.

```
calltree -r testrun.csv
```

There are trade-offs between a static calltree and a run-time calltree. As stated, a static calltree misses function calls that are computed on run-time: notably functions called via function pointers and virtual class members. A run-time calltree includes those functions, but on the other hand it lacks calls to library functions —as those are not instrumented.

Firmware Programming

As shown in chapter [Debugging Code on page 29](#), GDB downloads the code in the microcontroller as part of the debugging process. This opens the way for using the Black Magic Probe for small-scale production programming as well.

Using GDB

You can use GDB for uploading code to Flash memory by setting commands on the command line. The following snippet is a single command broken over multiple lines, for the Microsoft Windows command prompt (on Linux, replace the “^” symbol at the end of each line by a “\,” see the second snippet below). In practice, you would put it in a batch file or a bash script.

```
arm-none-eabi-gdb -nx --batch ^
-ex 'target extended-remote COM9' ^
-ex 'monitor swdp_scan' ^
-ex 'attach 1' ^
-ex 'load' ^
-ex 'compare-sections' ^
-ex 'kill' ^
blinky.elf
```

You need to change COM9 to the serial device that is appropriate for your system, and `blinky.elf` to the appropriate filename. On Linux, you may use the `BMScan` utility to automatically fill in the device name for the `gdbserver` virtual serial port:

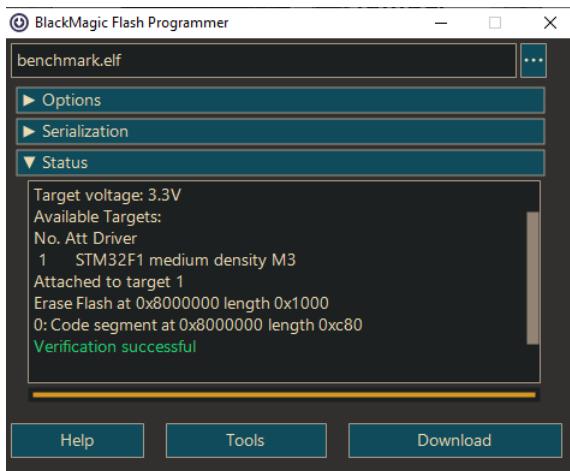
```
arm-none-eabi-gdb -nx --batch \
-ex "target extended-remote `bmscan gdbserver`" \
-ex "monitor swdp_scan" \
-ex "attach 1" \
-ex "load" \
-ex "compare-sections" \
-ex "kill" \
blinky.elf
```

Also see the note on the LPC microcontroller series from NXP regarding the `compare-sections` command on [page 34](#).

Using the BlackMagic Flash Programmer

The `BMFlash` utility is a GUI utility that offers a few additional features over GDB for firmware programming. The `BMFlash` utility uses the *Remote Serial Protocol* (RSP) of GDB to directly communicate with the Black Magic Probe.

GDB is therefore not required to be installed on the workstation on which you perform production programming.



The BMFlash utility automatically scans for the Black Magic Probe on start-up, and connects to it. It also has built-in handling of the idiosyncrasies of the LPC microcontroller series from NXP (see [page 34](#)).

Serialization

BMFlash supports serialization: patching a serial number into the code that is downloaded to the target, and incrementing that serial number for each successful download.

The modes that are available for serialization are:

No serialization	No serialization is performed.
Address	The options for this mode are the name of a section in the ELF file, and the offset in bytes from that section. The offset is a hexadecimal value. The section name is typically ".text" or ".rodata." If the section name is empty, the offset is from the beginning of the ELF file.
Match	In this mode, the BMFlash utility searches for a signature or byte pattern in the original ELF file, and stores another byte pattern ("prefix") plus a serial number at that spot.
	The "match" string can be an ASCII string, like "\$serial\$." It can also contain binary values, which you specify with \ddd or \xhh where ddd is a decimal number of up to three digits and hh is a hexadecimal number of up to two digits (thus, the codes \27 and \x1b are the same). When the code \U* appears in the string, a zero byte is

	<p>added to the match pattern after each byte. The purpose is to make matching Unicode strings easier. The code \A* reverts to single-byte characters.</p> <p>If a backslash must be matched, it must be doubled in the match field.</p> <p>The “prefix” string follows the same syntax as the “match” string. It is optional; if not present, the serial number is written from the start of the signature found in the ELF file. If you want to store the serial number behind the signature in the ELF file (without modifying the signature), the prefix string should be set equal to the match string.</p>
--	--

The starting serial number itself and its width in characters or bytes are decimal values. The serial number can be stored in one of three formats:

Binary	The serial number is stored as an integer, in Little Endian byte order. The width of the serial number will typically be 1, 2, or 4, for 8-bit, 16-bit and 32-bit integers respectively, but other field sizes are valid.
ASCII	The serial number is stored as text, using ASCII characters. The number is stored right-aligned in the field size of the serial number, and padded with zero digits on the left. For example, if the serial number is 321 and the width is 6, the serial number is stored as the ASCII string “000321.”
Unicode	The serial number is stored as text, using 16-bit wide Unicode characters. The width for the serial number should be an even number.

Settings for serialization and other configurations are stored in a file that has the same name as the target (ELF) file, but with the extension “.bmcfg” to it.

Log file

The BMFlash utility can optionally add a row to a log file for each successful download. To activate it, set a check-mark in the “Keep log of downloads” option in the “Options” tab. The log file is in CSV format (comma-separated values). The filename is the same as that of the ELF file (the file that is downloaded to the target), but with the extension “.log” appended.

Each row starts with the date and time of the download. It is followed by three fields identifying the ELF file: the file date & time, the size in bytes, and a POSIX checksum. This checksum is actually a CRC32 of the contents of the file plus the file size. After that, there is an RCS identification string read from the ELF file (if one was present), and finally the serial number patched into the code during the download (if serialization is enabled).

The POSIX checksum enables you to distinguish which version of the ELF file was downloaded. It is calculated over the original ELF file, before patching

a serial number in the code. You can verify whether an ELF file matches the number in the log file, by running `cksum` on the ELF file. The utility `cksum` is a core utility of Unix and Linux distributions; it has also been ported to Windows as part of the GnuWin32 project. A self-contained re-implementation of the `cksum` utility (with minimal features) is also provided with this book.

The RCS identification string is easier to use as a unique identifier of the code that was downloaded. It works in combination with a version-control system and a placeholder for the identification string in the source code. On each commit, the version-control creates a unique stamp and patches that into the placeholder in the source code. With the stamp, you can then look up the matching commit in version-control history.

RCS (*Revision Control System*) is legacy version-control software, but the format for the identification strings lives on. A typical string that you would add to the main source file of your embedded application is:

```
const char __id[] = "$Revision$";
```

Whereas RCS used a handful of keywords, BMFlash only supports `Id` and `$Revision$` (which may be abbreviated to `Rev`). You must furthermore enable keyword expansion in your version-control software, on all source files. In Apache Subversion, add the “`svn:keywords`” property to the source files and add at least the “`Id`” and/or “`Revision`” keywords to the list. For git, you can add the following lines to the `.gitattributes` file:

```
*.c ident  
*.h ident
```

Note that git only supports the `Id` keyword (not `$Revision$`).

These are not the only solutions (in fact, the above solutions have their shortcomings). When using Subversion, a more reliable scheme is to use the `SvnRev` utility as part of the build. The code to add to the main source file changes to the snippet below:

```
#include "svnrev.h"  
const char __id[] = SVNREV_RCS;
```

Enabling keyword expansion is not required when using `SvnRev`. For git, Mercurial and Bazaar, an alternative is Autorevision, but which runs only in Linux. See [Further Information](#) on page 119 for links to the various utilities.

As an aside, if you want to check whether a file contains RCS identification strings, you can use the `ident` utility. This utility is part of the RCS package (and of the GnuWin32 project); a self-contained re-implementation is also provided with this book.

Post-processing

Especially when serialization is active, it may be needed to run a script or program after each successful download —for example, to print a label with the serial number. The BMFlash utility supports this via the “Post-process” option (“Options” tab). If the post-process field holds the name (and path) of an executable program, this program is invoked after each successful download with the ELF file and (optionally) the serial number as arguments.

The post-process runs after updating the log file (if the option to log successful downloads is enabled). A post-process program can therefore extract relevant fields from the log file —for example, to store the information in a database that holds more complete information on the device or firmware.

Miscellaneous tools

The Tools button on the bottom row of the utility provides a few additional commands:

- ◊ Re-scan for Black Magic Probes on the USB bus (e.g. for the case that you launched the utility without first connecting a Black Magic Probe).
- ◊ Erase the full Flash memory of the target, see also the notes below.
- ◊ Erase the option bytes (on microcontrollers that use option bytes), see the notes below.
- ◊ Activate “code read protection” in the option bytes (on microcontrollers that support CRP via option bytes).
- ◊ Verify the code in the microcontroller against the ELF file (without downloading it).

On LPC microcontrollers (from NXP), erasing full Flash memory also clears “code read protection” (CRP). However, on these microcontrollers, CRP has also disabled the SW-DP interface on reset or power-up, with the implication that the Black Magic Probe can no longer access the microcontroller. Therefore, if you accidentally download code with CRP into your development device, you must erase the Flash memory immediately, without leaving the BMFlash utility and without power-cycling (or resetting) the target device. The BMFlash utility notifies you when an ELF file has CRP set, upon opening the ELF file.

The STM32Fxx microcontrollers use option bytes for code protection. Erasing these clears the protection, and by clearing protection, it also erases all Flash memory. The STM32Fxx microcontrollers need a power-cycle, before the change in option bytes is picked up. If the target is powered from the Black Magic Probe, this is handled automatically by the BMFlash utility. When

the target device is self-powered, you should power-cycle it after clearing the option bytes.

Setting code protection with the BMFlash utility currently only works on microcontrollers from the STM32Fxx series, and only for RDP Level 1. After setting it, the target needs to be power-cycled for the new values of the option bytes to be picked up.

Also see the section [Reset Code Protection](#) at page 32 for more information.

Updating Black Magic Probe Firmware

At the time of writing, a new version of the Black Magic Probe hardware has started shipping: version 2.3. The predecessor, “hardware version 2.1,” is now sold out. The firmware is in continuous development, however, for both versions 2.3 and 2.1 (and for other platforms as well). Support for more microcontrollers, as well as new features, are regularly being committed to the GitHub project. There is therefore good reason to update the firmware of the Black Magic Probe to either the latest “stable” firmware release (version 1.8.2 at the time of writing), or an up-to-date “development version.”

You can build the latest firmware yourself, but you often do not need to. The stable releases are available on the GitHub project, and pre-compiled builds of the development release are also updated each day. See chapter [Further Information on page 119](#).

An essential step for Microsoft Windows is to complete the set-up for DFU. See the instructions in Setting up the Black Magic Probe on [page 18](#). As noted in that section, both the DFU interfaces for *normal mode* and *DFU mode* must be installed.

The next step is to install `dfu-util` for your operating system. For Microsoft Windows, download a “binaries” release (see [Further Information on page 119](#) for the download location) and unpack it in a directory of your choice. For Linux, it is more convenient to use the package manager of your distribution to get the latest version; for example:

```
$ sudo apt-get install dfu-util
```

The options on `dfu-util` for updating the firmware are (native Black Magic Probe):

```
dfu-util -d 1d50:6018,:6017 -s 0x08002000:leave -D blackmagic-native.bin
```

For `ctxLink`, use the command below (note the address set with the `-s` option):

```
dfu-util -a 0 -s 0x08000000 -D blackmagic-ctxlink.bin
```

On Linux, you may need to run the command with `sudo` (this depends on whether a `udev` rules file has been installed for the Black Magic Probe, see [page 20](#)).

You can check which firmware version you have with the GDB monitor command (after connecting it as an “extended-remote” target, see also [page 25](#)).

```
(gdb) monitor version
Black Magic Probe v1.8.2, Hardware Version 3
Copyright (C) 2022 Black Magic Debug Project
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
```

The development release of the firmware uses a GitHub hash instead of a version number. In the snippet below, it is the hexadecimal value “0077d31.”

```
(gdb) monitor version
Black Magic Probe 0077d31, Hardware Version 3
Copyright (C) 2022 Black Magic Debug Project
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
```

A drawback of a hash is that they are not monotonically incrementing: a more recent firmware may have a hash value that is lower than the previous version. To find out at what position on the commit timeline a particular hash sits, you have to go to the GitHub project for the Black Magic Probe, and search that repository for the hash number.

The main page for the search results will then tell you that it couldn't find any *code* matching the hash (0077d31 in the above example).

0077d31

Code	0
Commits	1
Issues	1
Discussions	0
Packages	0
Wikis	0

We couldn't find any code matching '0077d31' in [blackmagic-debug/blackmagic](#)
You could [search all of GitHub](#) or try an [advanced search](#).

Advanced search Cheat sheet

Terms Privacy Security Status Docs Contact GitHub Pricing API Training Blog About

© 2022 GitHub, Inc.

In the left column, however, is a list with the topics Code, Commits, Issues, and a few others. If you click on “Commits” (see arrow in the picture) you will get a summary of the relevant commit, plus the date of that commit.

Troubleshooting

The first step in troubleshooting is to check whether the Black Magic Probe has power. When the Black Magic Probe is connected to a USB port, the green and orange LEDs (labelled “PWR” and “ACT” respectively) should be on. The ACT LED may be bright or dim, depending on the firmware version. On the ctxLink, only the ACT LED must be on; the green LED indicates the connection mode, not “power.”

After having verified that the Black Magic Probe is powered-on, the next steps depend on the issue that you are having.

Check whether the system detects the probe

The step to check whether the system can find the Black Magic Probe is covered in section [Checking the Set-up, page 24](#). To summarize, the BMScan utility shows the interfaces of all probes that it can find. For troubleshooting, it is of course recommended connecting only a single debug probe at a time; for the ctxLink, we also recommend to first get the device set up with a USB connection.

```
d:\Tools> bmscan

Black Magic Probe found, serial 7BB180B4:
gdbserver port: COM9
TTL UART port: COM10
SWO interface: {9A83C3B4-0B99-499E-B010-901D6C2826B8}
```

The BMScan utility also does a quick test on whether it can open the ports that it detects. If it indicates “[no access]” after a port (like in the example below), it has detected the port, but at the same time failed to open it.

```
$ bmscan

Black Magic Probe found, serial 7BB180B4:
gdbserver port: /dev/ttyACM0 [no access]
TTL UART port: /dev/ttyACM1 [no access]
SWO interface: 1-2:1.5
```

There are multiple reasons why a serial port cannot be opened. An obvious one is that the port is already open. On Linux, another common cause is access rights: a non-root user must be included in the dialout group to access the ports, see [page 21](#) for more information.

Check whether the probe detects the target

As is common for open-source project, several derivatives of the Black Magic Probe have emerged. The ctxLink is one of these, and one that offers additional features. Most derivatives, however, are optimized on cost —by cutting down on features or quality. Some of these probes do not have sufficient Flash memory to contain the current Black Magic Probe firmware, others are simple (FTDI MPSSE-based) protocol interfaces that don't contain any embedded firmware at all. To make these low-cost probes work, the firmware of the Black Magic Debug project can also be built as a desktop application. Thus, the “firmware” runs on a workstation or laptop, and it communicates with the debug probe via one of several low-level serial protocols. It was originally called the “hosted” set-up, but the project now prefers to refer to it as BMA or BMDA, which stands for the Black Magic (Debug) App.

The relevance, in regard to this chapter, is that the native Black Magic Probe also supports a low-level serial protocol, and (more importantly), the desktop-build of the Black Magic Probe firmware offers additional diagnostics. For the purpose of troubleshooting, the advice is therefore to run the Black Magic code on the desktop, while it is connected to the Black Magic Probe hardware (and while the Black Magic Probe is also connected to a target device).

```
$ blackmagic -t
```

The “-t” option displays the information that the utility gathers about the debug probe and the target. When the target is powered from the Black Magic Probe, you should also add the “-p” option. The utility has more options that may be relevant. Use the following command to list them all:

```
$ blackmagic -h
```

The output of the “blackmagic -t” command, for the case that no problems are detected, is similar to the snippet below:

```
$ blackmagic -t

BMP hosted
for ST-Link V2/3, CMSIS_DAP, JLINK and LIBFTDI/MPSSE
Running in Test Mode
Target voltage: 3.3V Volt
Speed set to 3.2727 MHz for SWD
DPIDR 0x0bb11477 (v1 MINDP rev0)
RESET_SEQ failed
AP 0: IDR=04770021 CFG=00000000 BASE=e00ff003 CSW=03000040 (AHB-AP var2 rev0
Halt via DHCRR: success 01030003 after 1ms
ROM: Table BASE=0xe00ff000 SYSMEM=0x00000001, designer 43b Partno 471
0 0xe000e000: Generic IP component - Cortex-M0 SCS (System Control Space)
(PIDR = 0x04000bb008 DEVTYPE = 0x00 ARCHID = 0x0000)-> cortexm_probe
```

```
CPUID 0x410cc200 (M0 var 0 rev 0)
1 0xe0001000: Generic IP component - Cortex-M0 DWT (Data Watchpoint and Trace)
  (PIDR = 0x04000bb00a DEVTYPE = 0x00 ARCHID = 0x0000)
2 0xe0002000: Generic IP component - Cortex-M0 BPU (Breakpoint Unit) (PIDR =
  0x04000bb00b DEVTYPE = 0x00 ARCHID = 0x0000)
ROM: Table END
*** 1      LPC11xx M0
RAM   Start: 0x10000000 length = 0x2000
Flash Start: 0x00000000 length = 0x20000 blocksize 0x1000
```

One of the first things to look at is the target voltage; it is 3.3V in this example. The Black Magic Probe uses the voltage on the VREF pin for its voltage level shifters on the debug pins. When the voltage on the VREF pin is zero, for example because you did not wire the VREF pin to the target, then the Black Magic Probe won't work.

If the target runs on 3.3V, you can power it from the Black Magic Probe by adding “-p” option on the command line of the `blackmagic` program (in addition to “-t”).

The following snippet shows a case where the debug probe cannot find a target microcontroller:

```
$ blackmagic -t

BMP hosted
for ST-Link V2/3, CMSIS_DAP, JLINK and LIBFTDI/MPSSE
Running in Test Mode
Target voltage: 3.3V Volt
Speed set to 3.2727 MHz for SWD
Exception: SWDP invalid ACK
Trying old JTAG to SWD sequence
Exception: SWDP invalid ACK
No usable DP found
Can not attach to target 1
```

You will also get this response when the target voltage is 0V but that is not the case here. You should double-check the wiring between the Black Magic Probe and the target device. You may want to try to repeat the command again with the “-C” option (for connecting under reset). Other explanations are:

- ◊ The target microcontroller has redefined the SWCLK and/or SWDIO pins, and hence you need to reset to bootloader mode (see [page 27](#)).
- ◊ The SWD interface has been disabled altogether —e.g. because code-read protection is active on an NXP LPC-series microcontroller. See section [Reset Code Protection](#) on [page 32](#) for details.

Another case that may occur is that the target microcontroller does not yet appear in the tables of the Black Magic Probe. New microcontrollers are introduced on the market at a quick pace, and software support for them is often a bit lagging behind. The output for an unsupported microcontroller is similar to the snippet below (this is a contrived example, the particular microcontroller with these designer & part numbers and ID code *is*, in fact, supported by the Black Magic Probe):

```
$ blackmagic -t

BMP hosted
for ST-Link V2/3, CMSIS_DAP, JLINK and LIBFTDI/MPSS
Running in Test Mode
Target voltage: 3.3V Volt
Speed set to 3.2727 MHz for SWD
DPIDR 0x0bb11477 (v1 MINDP rev0)
RESET_SEQ failed
AP 0: IDR=04770021 CFG=00000000 BASE=e00ff003 CSW=03000040 (AHB-AP var2 rev0
Halt via DHCSP: success 00030003 after 2ms
ROM: Table BASE=0xe00ff000 SYSMEM=0x00000001, designer 43b Partno 471
0 0xe000e000: Generic IP component - Cortex-M0 SCS (System Control Space)
(PIDR = 0x04000bb008 DEVTYPE = 0x00 ARCHID = 0x0000)-> cortexm_probe
CPUID 0x410cc200 (M0 var 0 rev 0)
LPC11xx: Unknown IDCODE 0x2998802b
LPC8xx: Unknown IDCODE 0xffffdff88
1 0xe0001000: Generic IP component - Cortex-M0 DWT (Data Watchpoint and Trace)
(PIDR = 0x04000bb00a DEVTYPE = 0x00 ARCHID = 0x0000)
2 0xe0002000: Generic IP component - Cortex-M0 BPU (Breakpoint Unit) (PIDR =
0x04000bb00b DEVTYPE = 0x00 ARCHID = 0x0000)
ROM: Table END
*** 1 Unknown ARM Cortex-M Designer 43b Partno 471 M0
```

For the team maintaining the Black Magic Probe firmware (on GitHub), there are several important values in this dump. The microcontroller is detected as a Cortex-M0 architecture. The numeric ID for the designer is 43b (hexadecimal) and the ID for the part is 471. These values are not conclusive: the value 43b is the code for ARM Ltd., for example—but the part is from NXP.

The information on the architecture, and the designer and part IDs is sufficient, however, to probe deeper. As can be seen from the output, the Black Magic Probe tries to match microcontrollers in the LPC1100 and LPC800 series. Both attempts fail, but the IDCODE values are important. In this particular case, the microcontroller being tested was an LPC11U14, and many microcontrollers from that series are already supported. It may be sufficient to add the number 0x2998802b to a table or list of known codes that are matched against.¹

¹ As stated earlier: this is a contrived example. The given code (0x2998802b) is already in the

When running the Windows build of the `blackmagic` utility, you may need to set the “`-d`” option with the COM port of the Black Magic Probe, in addition to the other options.

```
$ blackmagic -t -d com9
```

Whether or not this option is needed depends on the build options for the utility —more specifically, it depends on which debug probes the utility is built to support. When the `blackmagic` utility is configured (during compilation) to only support the Black Magic Probe (or 100% compatibles like the `ctxLink`), this option is not needed; when the utility is configured to support additional debug probes (e.g. ST-Link V2 or V3, or CMSIS-DAP), you will need to set the COM port for the Black Magic Probe.

How to get the hosted `blackmagic` utility?

When you come to the point that you want to run the `blackmagic` program for troubleshooting, the first hurdle in doing so is... that you don't have it. The hosted build is not part of the official releases.

If you are running Linux, you can get a daily build from GitHub. An executable build of the “hosted” firmware is part of the daily builds.

On Microsoft Windows, you have a few options. One is to set up Linux in a virtual machine, like VirtualBox. You are then running a true Linux and you can run the Linux version of the `blackmagic` utility. Alternatively, you can set up Windows' versions of GCC (we recommend Mingw-w64) together with a Unix-like environment. Sid Price has documented the steps for Cygwin, see [Further Information](#) on [page 119](#) for a link. An alternative is msys2; the choice between Cygwin and msys2 is a matter of personal preference.

Another option is to download a pre-build version of the `blackmagic` utility for Microsoft Windows. For your convenience, an executable `blackmagic` utility for Microsoft Windows is included in the software package that is provided with this book (which is also on GitHub) —again, see [Further Information](#) on [page 119](#) for a link. The drawback is that the software for this book is updated when there is a new edition ready. In other words, the pre-compiled `blackmagic` utility that is distributed with this book may not be the latest version.

Target scan hangs

When the “`monitor swdp_scan`” command appears to hang, this may indicate a problem in the communication between the Black Magic Probe and the target. A typical situation is:

list for the LPC11** series. I used a sabotaged build of the firmware to get this output.

```
(gdb) monitor swdp_scan  
Target voltage: 3.3V
```

Neither does the list with available targets appear, nor is there an error message. The “(gdb)” prompt does not appear either —until you break the connection between your workstation and the Black Magic Probe.

The target voltage is measured by the Black Magic Probe itself; this first part of a scan does not involve communication with the target. The second phase of a scan, however, is to query the target using the SW-DP protocol. More pointedly, the target scan is typically the first command that involves communication with the target.

The *hang-up* is the result of the target responding, but responding in a way that confuses the debug probe or GDB. This can be caused by the wiring between the Black Magic Probe and the target: a flaky connection, noise picked up by long wires, reflections (transmission line effects), ... So the first step is to check connections and whether there is a noise source (like, for example, a noisy power supply that powers the target).

If the above step does not (reliably) fix the problem, an alternative is to reduce the speed of the SW-DP protocol, using the “monitor frequency” command. On start-up, the SWCLK clock runs at approximately 3.5 MHz. You can lower this with the following command:

```
(gdb) monitor frequency 2M  
Max SWJ freq 00225510
```

This sets the SWCLK frequency to roughly 2 MHz —roughly, because the slowdown is done by inserting “wait states” in the toggling of the SWCLK pin, and those wait states are in discrete amounts. The value that the command returns, 00225510 in the above example, is in hexadecimal in firmware 1.8.

You need to run the “monitor frequency” command *before* an “swdp_scan”, of course.

GDB crashes on “attach”

If you get either of the following errors on attaching to the target, this relates to a bug in GDB versions 11.1, 11.2 and 12.1.

```
(gdb) attach 1
.../gdb/remote.c:7979: internal-error: ptid_t remote_target::select_thread_for_ambiguous_stop_reply(const target_waitstatus*): Assertion `first_resumed_thread != nullptr' failed.
A problem internal to GDB has been detected,
further debugging may prove unreliable.
Quit this debugging session? (y or n) [answered Y; input not from terminal]

This is a bug, please report it. For instructions, see:
<https://www.gnu.org/software/gdb/bugs/>.
```

or

```
(gdb) attach 1
Attaching to program: blinky, Remote target
.../gdb/thread.c:72: internal-error: thread_info* inferior_thread(): Assertion `current_thread_ != nullptr' failed.
A problem internal to GDB has been detected,
further debugging may prove unreliable.
```

The GDB bug affects communication with gdbserver stubs in general, so it is not specific to the Black Magic Probe or any particular microcontroller. A patch has been available for some time, but at the time of writing, there is no confirmation that this patch will be included in the 12.2 release.

While waiting for a fix, there are two workarounds:

- ◊ revert to GDB version 10;
- ◊ update the Black Magic Probe firmware to 1.8.2 (or later), which contains a bypass for the GDB issue.

Failure to erase Flash memory

When the target microcontroller is from the STM32F series, the following error on a load command is likely due to **readout protection**” (RDP).

```
(gdb) load
Error erasing flash with vFlashErase packet
```

For more information, see [Reset Code Protection](#) on page 32.

Spying on the communication

GDB communicates with the Black Magic Probe via the *Remote Serial Protocol* (RSP). This is largely a text-based protocol. GDB lets you view the commands and responses of this protocol with the following command:

```
(gdb) set debug remote 1
```

The strings of RSP are intermixed with the other console output of GDB. They are easy to recognize, though: each string is prefixed with “Sending packet:” or “Packet received:.” In RSP, binary data, and monitor commands and their replies, are transmitted in encoded form. The strings that GDB prints in its console are the data as it transmits it to the Black Magic Probe; that is, in encoded form.

For example, the snippet below shows the output of a `monitor` command. The `monitor` command itself is translated to the `$qRcmd` command; its parameter (the string “`tpwr enable`”) is encoded as two hexadecimal digits per character. The reply starts with the letter “0” and then a string that is encoded in the same way. The decoded message (“Enabling target power”) is printed next.

```
(gdb) set debug remote 1
(gdb) monitor tpwr enable
Sending packet: $qRcmd,7470777220656e61626c65#07...Ack
Packet received: 0456e61626c696e672074617267657420706f7765720a
Enabling target power
Packet received: OK
```

The Remote Serial Protocol is described in detail in the GDB manual, “Debugging with GDB.” See [Further Information](#) on page 119 for a reference.

TRACESWO Capture

If a trace monitor, such as the [BlackMagic Trace Viewer](#) (page 69), stays fully silent —no trace messages & no error messages, the first things to check is whether there is a good voltage on the VREF pin of the Black Magic Probe. The reason is that the TRACESWO pin goes through a voltage level shifter on the Black Magic Probe (like the other debug pins), and that the “target side” of this level shifter is powered through the VREF pin. If you power the target through the Black Magic Probe (see the `monitor tpwr` command, [page 44](#)), this step is moot. If a trace monitor successfully attaches to the Black Magic Probe, this is also an indication that there is an adequate voltage at the VREF pin.

Once you have established that the VREF pin is powered, you can subsequently check, with a logic analyser or an oscilloscope, whether trace data is received at all on the TRACESWO line. Absence of data means that on the TRACESWO line means that SWO tracing has not been configured (or not configured correctly) in the target. Depending on the trace monitor, you may need to configure tracing from inside your source code, or you may need to run a particular script from GDB or some other tool.

If there is data on the TRACESWO line, check whether it is in the correct format and with a bit-rate that is in range with the capabilities of the debug probe. For instance, the Black Magic Probe hardware supports Manchester mode, whereas ctxLink only supports asynchronous mode. For Manchester encoding on the Black Magic Probe, the bit-rate is limited to approximately 200 kb/s; asynchronous encoding typically supports higher bit-rates.

If the trace monitor shows an error message along the lines of “access denied” or “failure opening device,” this may indicate either a missing driver (on Microsoft Windows), or a missing udev rule (on Linux). See chapter [Setting up the Black Magic Probe](#), and specifically the relevant section on either [Microsoft Windows](#) (page 18) or [Linux](#) (page 20), for details.

Another test that you can do is to redirect the TRACESWO data to the virtual UART interface of the Black Magic Probe. You can view it using a serial terminal. Use the following command —see also [page 44](#):

```
(gdb) monitor traceswo decode
```

If data now shows up, it means that the target is configured correctly, and that the Black Magic Probe correctly receives the TRACESWO data. You can then focus on receiving the data on the dedicated interface.

TTL-Level UART

The Black Magic Probe has a TTL-level UART for general purpose communication with a device. This UART can be used together with the SWD interface or on its own.

A feature of the interface is that the RxD and TxD lines run through level shifters (just like the pins on the Cortex Debug Header). Thus, you can use the UART to interface with microcontrollers running at 5V as well as at 3.3V, 2.5V... down to 1.2V. The implication is, however, that there needs to be a voltage on the secondary side of the level shifters. This is why the UART connector (4-pin 1.25 mm pitch “PicoBlade”) has a VCC pin —VREF would be a more accurate name. The logic voltage of the device should be connected to this pin.

The VCC pin on the UART connector is the same as the VREF pin on the debug header. When the VREF pin is powered from the Black Magic Probe (the `monitor tpwr` command), so is the VCC pin. If the attached target is running on 3.3V, the wire to VCC is optional if you instead power the secondary side of the level shifters through the Black Magic Probe.

GDB on Microsoft Windows

GDB may optionally use an “index cache” to increase performance on debugging large executables. It stores this cache in the “home” directory of the workstation. To find the home directory, it uses the `HOME` environment variable. In Microsoft Windows, this variable is not set by default. Therefore, on launching GDB, you may be greeted with the warning:

```
warning: Couldn't determine a path for the index cache directory.
```

This warning may be safely ignored; for executable files of the scale that fit in a microcontroller, you are unlikely to notice any reduced performance.

Alternatively, you can set the `HOME` environment variable before launching GDB:

```
SET HOME=%USERPROFILE%
```

or in PowerShell:

```
$Env:HOME = $Env:USERPROFILE
```

To keep the variable permanently set (instead of having to re-type it each time that you open a console to run GDB), you can add the variable to the list of static environment variables, in the System Properties dialog, TAB Advanced. After clicking on the button Environment Variables (near the bottom of the dialog), you will be presented with a new dialog with two lists of environment variables: one for the variables for the current user and one for the system variables (valid for all users). If you are the only user of the workstation, it makes no difference which one you take.

Microcontroller Driver Support

Microcontrollers frequently need some configuration to set up specific GDB functions or SWO tracing. For example, the section [Flash Memory Remap](#) on [page 31](#) addressed a step that is needed before you can download code into the microcontrollers of the LPC families from NXP. In that section, we also recommended defining a command for that step in the .gdbinit file.

The utilities [BMFlash](#) and [BMDebug](#) run MCU-specific scripts to remap memory, and the utilities [BMTrace](#) and [BMDebug](#) also run MCU-specific scripts to configure SWO tracing. These utilities contain the scripts embedded in the executable, and they establish which script to run by evaluating the name of the MCU driver that the Black Magic Probe returns on attaching to it. However, the Black Magic Probe is continuously enhanced and extended, and microcontroller support is growing. To that end, the predefined hard-coded scripts can be extended or overruled.

Script definitions for new (or modified) scripts must be stored in a file with the name “bmscript” (no file extension). On Microsoft Windows, this script must be stored in the “BlackMagic” directory on the (roaming) “Application Data” folder. The “INI” files for the diverse utilities are stored here as well. On Linux, the bmscript file must be stored in the “.local/share/BlackMagic” directory below the home directory of the current user.

The syntax of the definitions in the bmscript file is similar to that of .gdbinit, but it is not compatible with it. Only “define” statements can occur in bmscript, and these define statements must conform to either a register definition, or a script definition.

```
define SYSCON_SYSMEMREMAP [ lpc8xx, lpc11xx, lpc12xx, lpc13xx ] = {int}0x40048000
define SYSCON_SYSMEMREMAP [ lpc15xx ] = {int}0x40074000
define SCB_MEMMAP [ lpc17xx ] = {int}0x400FC040
define SCB_MEMMAP [ lpc21xx, lpc22xx, lpc23xx, lpc24xx ] = {int}0xE01FC040

define memremap [ lpc8xx, lpc11xx, lpc12xx, lpc13xx ]
    set SYSCON_SYSMEMREMAP = 2
end

define memremap [ lpc15xx ]
    set SYSCON_SYSMEMREMAP = 2
end

define memremap [ lpc17xx ]
    set SCB_MEMMAP = 1
end

define memremap [ lpc21xx, lpc22xx, lpc23xx, lpc24xx ]
    set SCB_MEMMAP = 1
end
```

As is apparent in the above example, each register and each script has a list of microcontroller driver names between square brackets after the name. These driver names are the names that the Black Magic Probe reports when it scans the attached target. The name may end with an asterisk, for a wildcard. For example, if “STM32F1*” appears in this list, it matches STM32F101T8 as well as STM32F103C8.

The list of MCU drivers is a filter for the definition. Because of this filter, there is no conflict to define the same register name or script name twice, provided that there is no overlap in MCU driver names.

The names of the registers may be freely chosen, but the names of the scripts are predefined by the [BMFlash](#), [BMTTrace](#) and [BMDebug](#) utilities. The scripts that are currently defined are:

<code>memremap</code>	To make sure that the microcontroller’s Flash memory map conforms to the ELF file layout —see also page 31 .
<code>partid</code>	To read the microcontroller’s “device id” or “id-code”.
<code>swo_device</code>	For the MCU-specific configuration for SWO tracing.
<code>swo_trace</code>	For the configuration for SWO tracing that is common to all ARM Cortex microcontrollers.
<code>swo_channels</code>	To set the mask for enabled channels (for SWO tracing); this script is common to all ARM Cortex microcontrollers.
<code>swo_profile</code>	To configure statistical profiling through the <i>Instrumentation Trace Macrocell</i> (ITM) and <i>Data Watchpoint & Trace</i> (DWT) units; this script is also common to all ARM Cortex microcontrollers.
<code>swo_close</code>	To disable SWO tracing and profiling (reset ITM and DWT); this script is common to all ARM Cortex microcontrollers.

You will typically only add (or replace) the first three of this list, but you can override the generic scripts for a particular microcontroller as well.

The only operations allowed on the registers (within a script) are assignment with “=,” “|=” and “&=,” which function in the same way as in GDB (and the C language). Numbers can be in decimal or hexadecimal notation, a “~” may prefix a value (or register) to denote the bitwise inversion of the value.

The operand on the left of the assignment operator is always considered an address. A literal number or a parameter on the right of the operator is considered a value. Thus, the literal or the value of the parameter is stored at the address at the left. The C *dereference operator* “*” can be used to interpret the value as a pointer, and read the value that the literal points to. A register does not need a dereference operator; it is always considered to be an address that needs to be dereferenced.

A parameter is specified as “\$0” to “\$9”. The number of parameters and their values depends on the script. For the swo_channels script, for example, the single parameter \$0 is the bit mask of enabled channels. See the “**bmscript**” file that is provided with the utilities for details on the parameters for each script. When a parameter appears at the right side of an assignment, you may add both a shift value and a literal value behind it. The value of the parameter is then shifted left by the shift value and the literal value is merged with a binary “or” operation.

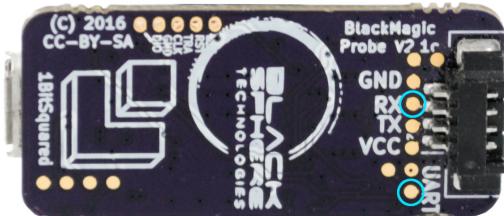
For a script that has a result (to be used by **BMD**Debug**, **BMP**Profile**, **BMT**Trace** or **BM**Flash**********

```
define DBGMCU_IDCODE [ stm32f0* ] = {int}0x40015800
define partid [ stm32f0* ]
    set $ = DBGMCU_IDCODE
end
```

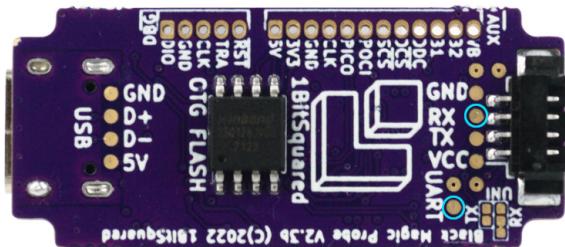
The assignment to the “\$” parameter should also be the last statement in the script. At the moment, this is relevant only to the “partid” script (the other scripts do not return a value).

Linking TRACESWO to UART-RxD

Part of why this book exists is for our own reference in how to configure and use the Black Magic Probe and its surrounding utilities. That may already have been conspicuous in the coverage of the ARM Cortex M0/M0+ architectures and the peculiarities of the NXP LPC series of microcontrollers —we happen to use these low-end microcontrollers a lot. Presented here, in this chapter, is a small modification that we make to the Black Magic Probes that we own —even though it is doubtful that many others will find it as useful as we do.



The modification is that we add a miniature slide switch (specifically: TE Connectivity product MLL1200S) that optionally connects the TRACESWO pin on the Cortex Debug connector to the RxD pin on the TTL-level UART connector. The switch is glued to the bottom of the PCB and two of its three pins are soldered to the test pads indicated with the light blue circles. The image above is for hardware version 2.1; the one below is for hardware version 2.3.



Linking TRACESWO to RxD enables the Black Magic Probe to receive the asynchronous SWO tracing protocol, although it now receives it via the UART interface instead of through the dedicated raw-data interface. Alternatively, it allows you to use UART tracing ([page 57](#)) over the debug connector, so that you only need a single cable between the debug probe and a Cortex M0 microcontroller (which lacks support for SWO tracing). The criterion for a “single cable” is especially relevant for our use of the tag-connect cable, see [page 15](#).

Further Information

Hardware

Black Magic Probe: The official Black Magic Probe hardware is available from:

1BitSquared	https://1bitsquared.de/products/black-magic-probe
adafruit	https://www.adafruit.com/product/3839
elektor	https://www.elektor.com/

ctxLink: can be obtained from:

Sid Price	http://www.sidprice.com/ctxlink/
Crowd Supply	https://www.crowdsupply.com/sid-price/ctxlink
Mouser	https://www.mouser.com/

3D Printed Enclosures for the Black Magic Probe can be found on Thingiverse. A simple clip that offers some protection for the 10-pin Cortex Debug header (see page 21) is “thing” 2387688 (by Michael McAvoy); a full enclosure with openings for the connectors, LEDs and button is “thing” 2836934 (by Emil Fresk).

<https://www.thingiverse.com/>

Designs for ctxLink enclosures can be downloaded from Sid Price’s GitHub page:

https://github.com/sidprice/ctxLink_cases

tag-connect: cables with a pogo-pin plug, specifically suited for firmware programming and debugging. The cable suitable for the ARM Cortex SWD interface are TC2030-CTX and TC2030-CTX-NL. See also [page 15](#).

<https://www.tag-connect.com/>

freeconnect: open-source a pogo-pin connectors by Rafael Silva, as an alternative to tag-connect cables.

<https://perigoso.github.io/freeconnect/>

PCBite: PCB holders and needle probes:

<https://sensepeek.com/>

Software

Black Magic Probe: The project website contains links to downloads, documentation and schematics:

<https://black-magic.org/>

Active development of the firmware happens on the GitHub project.

<https://github.com/blackmagic-debug/blackmagic>

Notes on building the firmware are in the wiki of this GitHub project. However, these notes are Linux-centric. For building on Microsoft Windows, see

the additional notes on Sid Price's blog, specifically:

<http://www.sidprice.com/2020/03/24/building-blackmagic-probe-on-windows/>

<http://www.sidprice.com/2018/05/23/cortex-m-debugging-probe/>

dfu-util: A utility to update the firmware of USB devices that support the DFU protocol.

<http://dfu-util.sourceforge.net/>

Zadig: A utility for installing the drivers for SWO tracing and firmware update, see chapter [Setting up Black Magic Probe](#), specifically [page 18](#).

<https://zadig.akeo.ie/>

libusbK: Drivers, support DLLs and development files for generic USB device access.

<http://libusbk.sourceforge.net/UsbK3/>

Orbuculum: A set of utilities to process the output ARM Cortex Debug interface (SWO tracing, exception trace, performance profiling, ...), see also [page 61](#).

<https://orbcode.org>

gdbgui: Various GDB front-ends were mentioned in chapter [Requirements for Front-ends \(page 11\)](#), but we have singled out gdbgui because it is cross-platform and open-source, and it offers the required features in a simple interface.

<https://www.gdbgui.com/>

Troll: A source-level debugger supporting the Black Magic Probe, that is independent of GDB.

<https://github.com/stoyan-shopov/troll>

turbo: A GDB front-end with specific support for the Black Magic Probe, by the author of the Troll debugger.

<https://github.com/stoyan-shopov/turbo>

Cortex Debug: Visual Studio Code extension that adds debugging functions, with (initial) support for the Black Magic Probe.

<https://github.com/Marus/cortex-debug>

GnuWin32: The GnuWin32 project has native Windows ports of many GNU utilities, like the core utilities and RCS.

<http://gnuwin32.sourceforge.net/>

SvnRev & Autorevision: Utilities to extract revision numbers or hashes from version-control repositories.

<https://www.compuphase.com/svnrev.htm>

<https://autorevision.github.io/>

Articles, Books, Specifications

Debugging with GDB; Richard Stallman, Roland H. Pesch & Stan Shebs; Free Software Foundation, 2011; ISBN 978-0-9831592-3-0.

The book is also available in PDF and HTML formats on:

<https://www.gnu.org/software/gdb/documentation/>

Embedded Debugging with the Black Magic Probe (this book) is available on GitHub in PDF format. The tools mentioned in this book, BMScan, BMDebug, BMProfile, BMTrace, BMFlash, elf-postlink and tracegen, live there as well.

<https://github.com/compuphase/Black-Magic-Probe-Book>

The Art of Debugging with GDB, DDD, and Eclipse; Norman Matloff & Peter Jay Salzman; No Starch Press, 2008; ISBN 978-1593271749.

The Definitive Guide to the ARM Cortex-M3, second edition; Joseph Yiu; Newnes Press, 2009; ISBN 978-1856179638.

Writing Solid Code, second edition; Steve Maguire; Greyden Press, LLC, 2013; ISBN 978-1570740558.

Common Trace Format: The specification of the binary format as well as the Trace Stream Description Language (TSDL), see chapter The Common Trace Format on [page 72](#).

<https://diamon.org/ctf/>

SEGGER RTT: Resources and documentation for SEGGER's Real Time Transfer protocol, including portable source code to include in your firmware, viewers, loggers and other tools.

<https://wiki.segger.com/RTT>

SVD repository: A collection of "System View Description" files for various microcontrollers can be found on GitHub:

<https://github.com/posborne/cmsis-svd>

Index

-
- ! * (asterisk), 89
.bmcfg file, 53, 99
.gdbinit file, **26**, 30, 31, 42, 63, 115
~ (filter character), 70
1BitSquared, 1
-
- A** Access memory, 31
Access point, 22, 23
Adapter board, 24
addr2line utility, 89
Address look-up, 89
Addyi (drug), 2
ADIS, 7
Akeo Consulting, 18
Altering execution flow, 37, 38
Apache Subversion, *see* Subversion
Application Data folder, 115
ARM CoreSight, 56, 63, 92
ARM Cortex, 1, 59
 M0/M0+ architecture, 8, 27, 28, 45, 59,
 61, 66, 93, 118
ASCII, 77
Assembly code, 43, 51
Assertions, 86, 87
Asterisk, 89, 116
Asynchronous encoding, 8, 44, 53, 63, 64,
 69
Attach target, 26
 ~ troubleshooting, 107, 110
Autocompletion, 48, 49
AutoDesk Fusion, 17
Automatic download, 47, 54
Autorevision utility, **100**, 120
-
- B** Babeltrace, 76
backtrace (command), **42**, 89
barectf utility, 73, 81, 85
Base (number), 76
Battery, 17, 22
 ~ connector, 14
 ~ status, 22
Bit rate, 63
Bit-banging, 67
BKPT (instruction), 59
Black Sphere Technologies, 1
blackmagic utility, 106, 108, 109
bmcfg file extension, 53, 99
BMDA, *see* Hosted set-up
BMDebug (front-end), 30, 42, **47**, 47–51, 53,
 54, 59, 85, 89
BMFlash utility, 7, 34, 47, **97**, 98
BMPProfile utility, 93

- BMScan utility, 18, 20, 23, **24**, 97
bmscript file, 115
BMTrace utility, 7, **69**, 70, 73, 85, 112
Boot pin, 28
Bootloader (MCU), 16, **28**, 31, 107
Break on exceptions, 45
Break-out board, 15, 24
Breakpoint, 10, 38, 39, 49, 56, 71
 command list, 71
 conditional ~, 39
 disable ~, 49
 enable / disable, 39
 enable ~, 49
 hardware ~, 10, 35, 49, 71
 ~ ID, 38
 one-time ~, 38
 software ~, 58, 88
Button, *see* Push-button
-
- C** Calculator, 41
Call stack, **42**, 71, 89
Calltree, 95, 96
Case-insensitive (search), 49
CDC class driver, 18–20
cflow (utility), 95
Channel (tracing), 61, 70, 82
Checksum, 99
 vector table, **34**, 47
cksum utility, 99
Clone (debug probe), 57
CMSIS, 50, 59, 61, 88
Code instrumentation, 56, 58, 71
 function entry & exit, 89
Code Read Protection, **32**, 33, 34, 46, 101,
 102, 107
CoderGears, 95
COM port, 25
Command file (GDB), 27, *see also* .gdbinit
 file
Command line
 autocompletion, 48, 49
 history, 48
Command list, 71

Commands (GDB), 26

- attach, 26
- backtrace, **42**, 89
- compare-sections, 34, 97
- connect_srst, 27
- continue, 71
- define, 26
- disassemble, 43, 51
- display, 40, 50
- file, 30
- help, **35**, 37, 54
- info, **35**, 54, 89
- load, 31, 32, 48, 54
- monitor, 11, 25, 32, 33, 40, **43**, 103, 104
- print, 41
- reset, 46, 54, 55
- run, 34
- start, 34
- target, 25
- trace, 52
- user-defined, 26, 31

Common Trace Format, 52, 53, 69, 70, **72**, 74, 75, 85, 90, 121

compare-sections (command), 34, 97

Conditional breakpoint, 39

Conditional compilation, 57

connect_srst (command), 27

Connector pin-out, 24

Console (GDB), 11, 30, 48

continue (command), 71

CoreDebug, 59

CoreSight architecture, 56, 63, 92

Cortex Debug header, 13, 15, 17, **24**, 24, 61, 118

Cortex M0/M0+ architecture, 8, 28, 45, 59, 61, 66, 93, 118

CppDepend (utility), 95

CRC mismatch, 47

Crowd Supply, 119

CSV file, 94, 96, 99

CTF, *see* Common Trace Format

- ~ packet, 74

ctxLink (debug probe), 1, 6, 14, 17, 22, 25, 103, 105

Cygwin, 109

D Dangling-else problem, 87

Data trace, 56

DDD (front-end), 5, 12

Debug Access Port, 1, 27

Debug connector, *see* Cortex Debug header

Debug Port, 8

Debug probe, 1, 5, 6

- ~ effect, 72

Debug symbols, 30, 31, 89

Debugger attached check, 28, 59

Development release (firmware), 103

Device Manager (Microsoft Windows), 18

DFU

- ~ mode, 20
- ~ protocol, 18, 20, 103
- dfu-util, 20, 103, 120
- DHCP, 23
- DHCSR, 28, 59
- dialout group, 21, 105
- DiaMon, 72, 74
- Disable breakpoint, 39, 49
- disassemble (command), 43, 51, *see also* Assembly code
- display (command), 40, 50
- Download to target, 31, 47, 54, 97
- DTR (serial port), 25
- Duplicate strings, 87
- DWARF, 30, 94, *see also* Debug symbols
- DWT, 56, 116

E Eclipse (front-end), 5, 11

Edit-Compile-Debug Cycle, 54

ELF file, 34, 47, 98

elf-postlink utility, 34, 47

Emulating SWO tracing, 66

Enable breakpoint, 39, 49

Enclosure, 17, 119

Endianness, 77, 78

Entry (function), 89

Entry point, 47

Environment variables, 26, 114

- HOME, **114**

Erase Flash memory failure, 111

ESD-protection, 17

Ethernet, 74

ETM, 56

Event

- ~ header, 74, 80
- ~ id, 80

Event (CTF), 80

Exceptions, 45

Execution point, 42, 48

- altering ~, 38

Exit (function), 89

F file (command), 30

Filter, 70

Find text (in source code), 49

Firmware download, *see* Download to target

Firmware update, 18, 19, **103**

Fixed-point numbers, 76

Flash memory, 10, 11, 27

- erase troubleshooting, 111
- ~ programming, 27, **97**, 97, *see also* BM-Flash
- ~ remap, 31

Frame (call stack), 42

freeconnect (plug-of-nails), 16, 119

Fresh, Emil, 17, 119

Front-end, 5, 11, 12
BMDebug, 30, 42, 47, 47–51, 53, 54, 59
gdbgui, 29, 54
FTDI MPSSE, 106
Function entry & exit, 89, 92
Function key, 49, 51

G

Gait, J., 72
Galvanic isolation, 17

GDB
 commands, *see* Commands (GDB)
 ~ console, 11, 30, 48
 versions 11 & 12, 110
gdbgui (front-end), 12, 29, 40, 54, 120
gdbserver, 1, 5–7, 18, 20, 25, 43, 70
git, 100
GitHub, 3, 50, 103, 104, 108, 119, 121
 ~ hash, 104
GnuWin32, 99, 100, 120

H

Halfword, 41
Hammer (Law of the ~), 72
Hard reset, 55
HardFault handler, 59, 60
Hardware breakpoint, 10, 35, 49, 71
 number of ~, 10
Header byte (SWO), 9
heapinfo (semihosting), 45
help (command), 35, 37, 54
History (commands), 48
HOME environment variable, 26, 114
Hosted set-up, 106, 109
HTTP provisioning, 22

IDC header, 24
ident utility, 100
Identifier (format), 82, 83
Index cache directory, 114
Inference rule (Make), 84
info (command), 35, 54, 89
Inline function, 90
Inlined function, 55
Instruction trace, 56
Instrumented profiling, 92
Instrumented trace, 56
Instrumenting code, *see* Code instrumentation
Interrupt Service Routine, 10
Invariants, 86
Isolation, *see* Galvanic isolation
ISP, 34
ITM, 8, 9, 61, 63, 85, 116

J
J-Link (Segger), 5, 69
Jitter, 81
JST PH connector, 14, 17
JTAG, 1, 5, 7, 24, 25
 ~ header, 15, 24

K

KDdbg (front-end), 5, 11
Keil ULINK-ME, 5

L
Latency, 72
Law of the Hammer, 72
LED, 20, 24, 25, 45, 105
Level shifters, 24, 44, 107, 112, 113
Li-Po battery, 15, 17, 22
libopencm3, 50, 59
librdimon, 59
libusbK, 19, 20, 120
License, 4
Line number lookup, 89
Link Register, 88
Linux, 109
Linux Foundation, 72
Little Endian, 99
load (command), 31, 32, 48, 54
Log file, 99, 101
LPC microcontrollers, 28, 31, 33, 34, 47, 65,
 66, 98, 101
 Flash Memory Remap, 31
LTtng, 73

M
MAC address, 23
Machine code, *see* Assembly code
Maguire, Steve, 86
Make (utility), 84
Manchester encoding, 8, 9, 52, 63, 64, 69
 clock derivation, 9
 emulation, 67
Maslow, Abraham, 72
Matloff, Norman, 3
McAvoy, Michael, 17, 119
MCU support scripts, 115
MEMMAP (register), 31, 115
Memory
 ~ access, 31
 display / set ~, 40, 41
 ~ watch, 51
Merge strings, 87
Metadata file (CTF), 52, 53, 72, *see also*
 TSDL
micro-USB, 13
Microsoft Windows, 108, 109
monitor (command), 11, 25, 32, 33, 40, 43,
 103, 104
Morse code, 45
msys2, 109

-
- N** Needle probes, 16
Nemiver (front-end), 11
Network scan, 23, 25
newlib C library, 45, 59
Nightly build, *see* Development release
Non-intrusive debugging, 56, 72
NRZ encoding, *see* Asynchronous encoding
NTRACE macro, 84
Number base, 76
-
- O** On-the-go programming, 13
One-time breakpoint, 38
OpenOCD, 6
Optimized code, 37, 55, 55
Option bytes (STM32), 32, 33, 101, 102
Orbuculum, 63, 69, 120
Overvoltage protection, *see* Galvanic isolation
-
- P** Packet
 - ~ header (CTF), 74, 75, **78**, 82, 85
 - ~ header (ITM), **9**, 61
 - ~ layout (CTF), 74Packet-based protocol, 74
Parity bit, 8
Part ID, 116
Passive listener, 52, 70
PCBite, 16, 119
Performance optimization, 92
Peripheral register, 8, 41
PicoBlade connector, 13, 14, 113
plugdev group, 21
Pogo-pins, 15, 119
POSIX checksum, 99
Post-conditions, 86
Post-mortem analysis, 56
Post-processing, 101
Power saving (microcontroller), 92
Power selection (ctxLink), 15
Power-cycle, 33, 55, 101, 102
Pre-conditions, 86
Price, Sid, 17, 109, 119
print (command), 41
printf, 59
printf-style debugging, 11, 86
Probe, *see* Debug probe
 - ~ effect, 72, 92Production code, 57
Profiling, 92, 116
Protocol
 - packet-based ~, 74
 - remote ~, *see* RSP
 - stream-based ~, 74Push-button (on board), 13, 20, 22
-
- R** RCS identification string, 99, 100
rdimon.specs, 59
RDP (STM32), 32, 33, 101, 102, 111
Read Protection, *see* Code Read Protection
Register
 - debug ~, 8
 - peripheral ~, 8, 41, 50
 - view ~, 43, 50Remote Serial Protocol, *see* RSP
reset (command), 46, 54, 55
Reverse calltree, 95
Ribbon cable, 15
RS232, 5, 25, 74, *see also* UART
RSP, 5, 7, 69, 94, 97, **111**
RTOS, 61, 85
RTT, 121
run (command), 34
Run-time calltree, 95, 96
Run-time tracing, 11, **56**, 56, 86
RZ encoding, *see* Manchester encoding
-
- S** *Salzman, Peter*, 3
SAM microcontrollers, 65
Sampling (profiling), 92
Scan targets, 44
 - ~ troubleshooting, 109*SciTools*, 95
Scope (variables), 42
Script, 115
Scripts (MCU support), *see* MCU support scripts
Section (ELF file), 98
Segger, 121
 - ~ J-Link, 5, 69Self-destruct code, 34
Semihosting, 45, 51, 57–59, 86, 89
Sensepeek, 16, 119
Serial monitor, *see* Serial terminal
Serial number, 98
Serial port, 21, 70, *see also* UART
Serial terminal, 44, 51, 57
Serial Wire Debug, *see* SWD
Serialization, 98, 99
Side-effect, 86
Silva, Rafael, 119
Software breakpoint, 58, 88
Software trace, 56
sprintf, 72
SSID (Wi-Fi), 22, 23
ST-Link clone, 57
Stable release (firmware), 103
Stack
 - ~ frame, 42
 - ~ pointer, 60
 - ~ trace, 39start (command), 34
Static calltree, 95, 96

Statistical profiling, 92, *see also* Profiling
stderr, 58
stdint.h, 77
Stepping through code, 36
 by instruction, 43, 51
 skip functions, 37
Stimulus ports, 61, 85
STM32 microcontrollers, 28, 32, 33, 64,
 101, 102
 option bytes, 32, 33, 101
 RDP, 32, 101
Stop & Stare, 56
Stream (CTF), 79–81
Stream-based protocol, 74
String, 77
Stub (debugger), 5, *see also* gdbserver
Subversion, 100
sudo, 21, 103
SVC (instruction), 59
SVD file, 50, 121, *see also* System View Description
SVDConv utility, 50
SvnRev utility, 100, 120
SW-DP protocol, 26, 110
SWCLK, 7, 24, 28, 45, 55, 110
SWD, 1, 5, 7, 24, 25, 27, 28
SWDIO, 7, 24, 55
SWDP scan, 27
 ~ troubleshooting, 109
SWO Tracing, 51, 61, 74, 79, 82, 83, 85
 emulation, 66
 protocol, 9, 62
 troubleshooting ~, 112
Symbolic information, *see* Debug symbols
SYSMEMREMAP (register), 31
System View Description, 50, 121

T tag-connect (plug-of-nails), 15, 16, 118, 119
Target
 attach ~, 26
 GDB command, 25
 ~ list, 44
 ~ power, 24, 112
 scan ~, 44, 109
Temporary breakpoint, *see* One-time breakpoint
Text search, *see* Find text
Thingiverse, 119
Thumb mode (ARM), 88
Time stamp, 70
Torx (screw head), 3
TPIU, 63
trace (command), 52
Trace capture, 18, 19, 21, 63
Trace context (CTF), 78
Trace Viewer, 69, 72, 73, *see also* BMTrace
tracegen utility, 73, 76, 77, 79, 81, 83, 84,
 85, 91

TRACESWO, 8, *see also* SWO Tracing
 ~ pin, 9, 24, 57, 61, 112, 118
traceswo (command), 63
Tracing, *see* Run-time tracing
Trailing-zero compression, 9, 53, 62
Transfer speed, 61, 72
Troll debugger, 7, 120
Troubleshooting, 105
 connection, 105
 GDB, 110
 SWO Tracing, 112
 target, 106
 UART, 113
TSDL, 72, 73, 74, 85, 95, 121, *see also* Common Trace Format
 types, 77
TTL-level UART, *see* UART
TUI, 5, 12
turbo (front-end), 120
Turnaround, 7
typealias, 76, 77
typedef, 76
Types (TSDL), 77

U UART, 8, 13, 14, 18, 20, 52, 57, 113, 118
 SWO tracing, 44, 68
 troubleshooting, 113
udev rules, 21, 22, 103, 113
ULINK-ME (Keil), 5
Understand (utility), 95
Unicode, 99
Unit testing, 86
USB, 74
USB ID, 20
User-defined command, 26, 31
UTF-8, 77
UUID, 78, 79

V Value history, 41
Variable, 40
 display format, 50
 ~ watch, 40, 50
Vector table checksum, 34, 47
vector_catch (command), 40
Version-Control software, 100
vFlashErase packet, 32, 111
VID:PID, 6, 20
VirtualBox, 109
Visual Studio, 12
Visual Studio Code, 12, 120
VisualGDB (front-end), 12
Voltage level, 24, 107, 112
VS Code, *see* Visual Studio Code

W Watch variable, 40, 50
register, 43
Watchpoint, 11, 38, 39
Weak linkage, 88
Wi-Fi link, 6, 22
Wiki, 119
Wildcard character, 116
WinGDB (front-end), 12

WinUSB device, 19, 20
WPS, 22

Y YAML, 73
Yiu, Joseph, 62

Z Zadig, 18-20, 120
Zero-terminated string, 77